

Operating Systems Concepts

Processes



CS 4375, Fall 2025

Instructor: MD Armanuzzaman (*Arman*)

marmanuzzaman@utep.edu

September 3, 2025

Summary


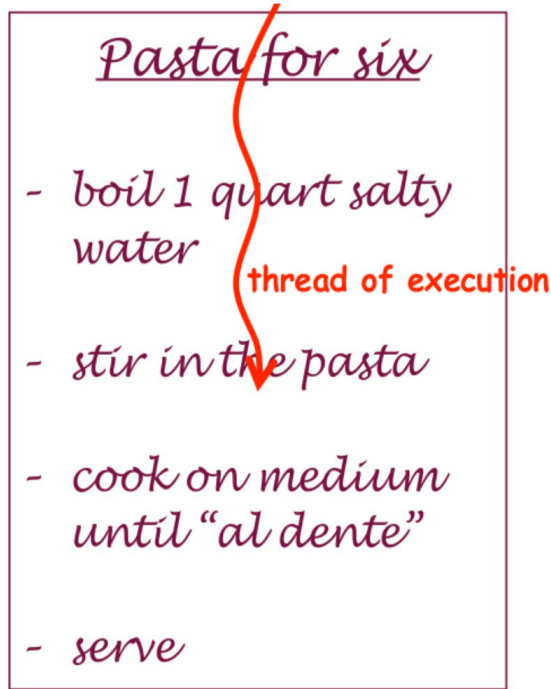
- A very brief history of computer architecture
- High level topics of the course
- Goals or Responsibilities of OS
- Tasks of an OS
- Major OS components
- OS design approaches

Agenda

- Processes
 - Process representation in OS
 - Process creation
 - Process termination
 - Context switching
 - Process queues
 - Process scheduling
 - Interprocess communication



Processes Concept

- A **process** is a program in execution.



Process

Process vs Thread



Processes Concept

- A process image consists of **three**

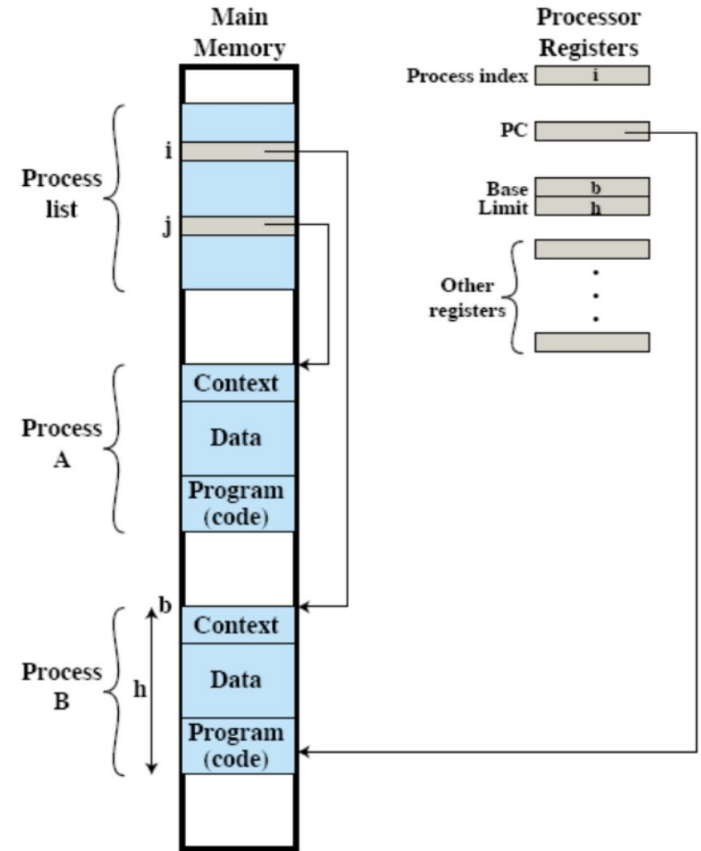
components:

- User Address space {
- An **executable** program
 - The associated **data** needed by the program
 - The execution **context** of the process which contains all information the OS needs to manage the process (ID, state, CPU registers, stack, etc.)

Processes Concept


- A process image consists of **three components**:
 - An **executable** program
 - The associated **data** needed by the program
 - The execution **context** of the process which contains all information the OS needs to manage the process (ID, state, CPU registers, stack, etc.)

User Address space



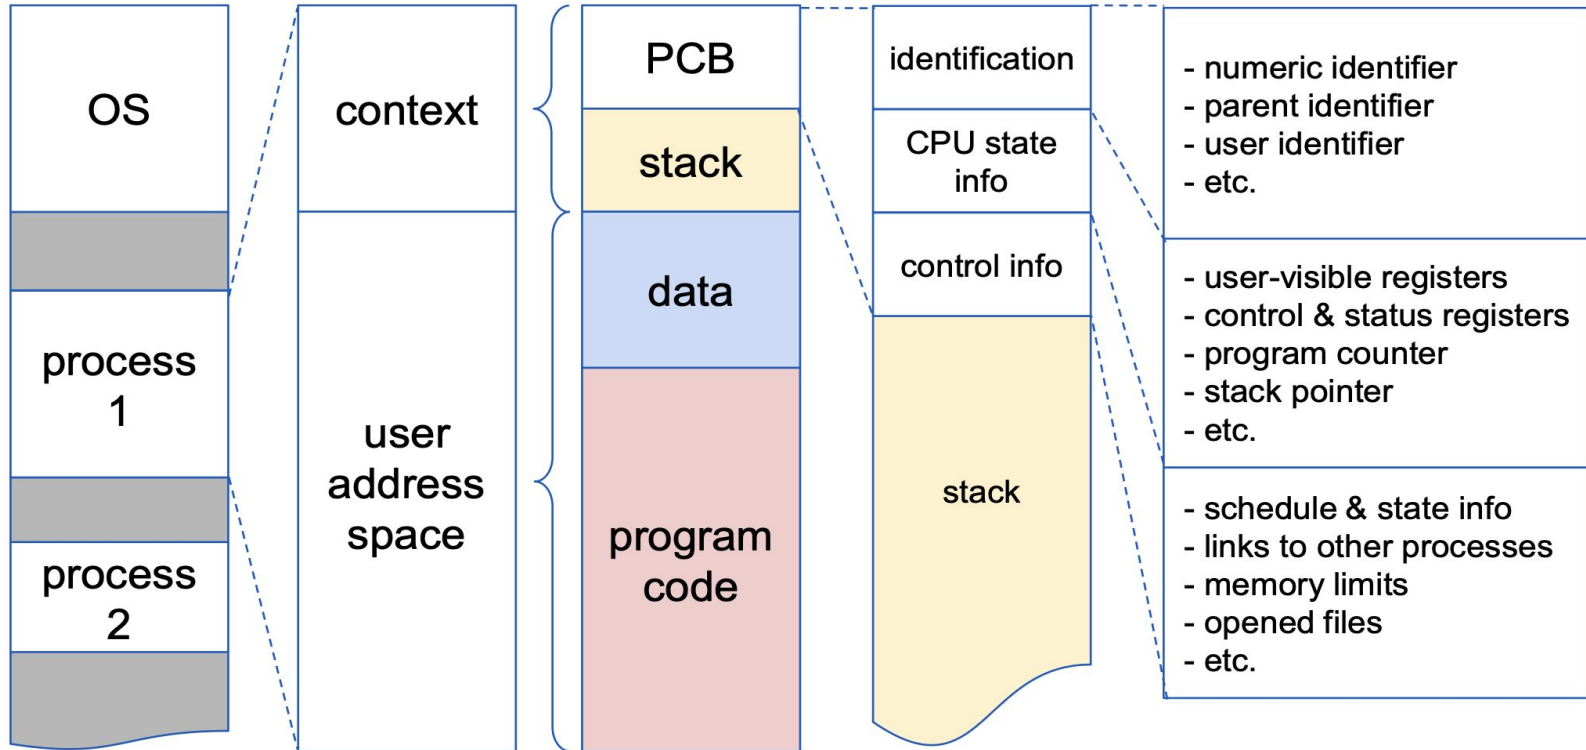
Process Control Block (PCB)

- PCB is included in the **context**, along with the stack
- PCB is a “**snapshot**” that contains all necessary and sufficient data to **restart** a process where it left off (ID, state, CPU registers, etc.)
- PCB is one entry in the OS’s **process table** (array or linked list)




Process Control Block (PCB)

Illustrative contents of a process image in (virtual) memory



Process State

- A process changes its **state** during execution:
 - **New**: The process is being created
 - **Ready**: The process is waiting to be assigned to a processor
 - **Running**: Instructions are being executed
 - **Waiting**: The process is waiting for some event to occur
 - **Terminated**: The process has finished execution



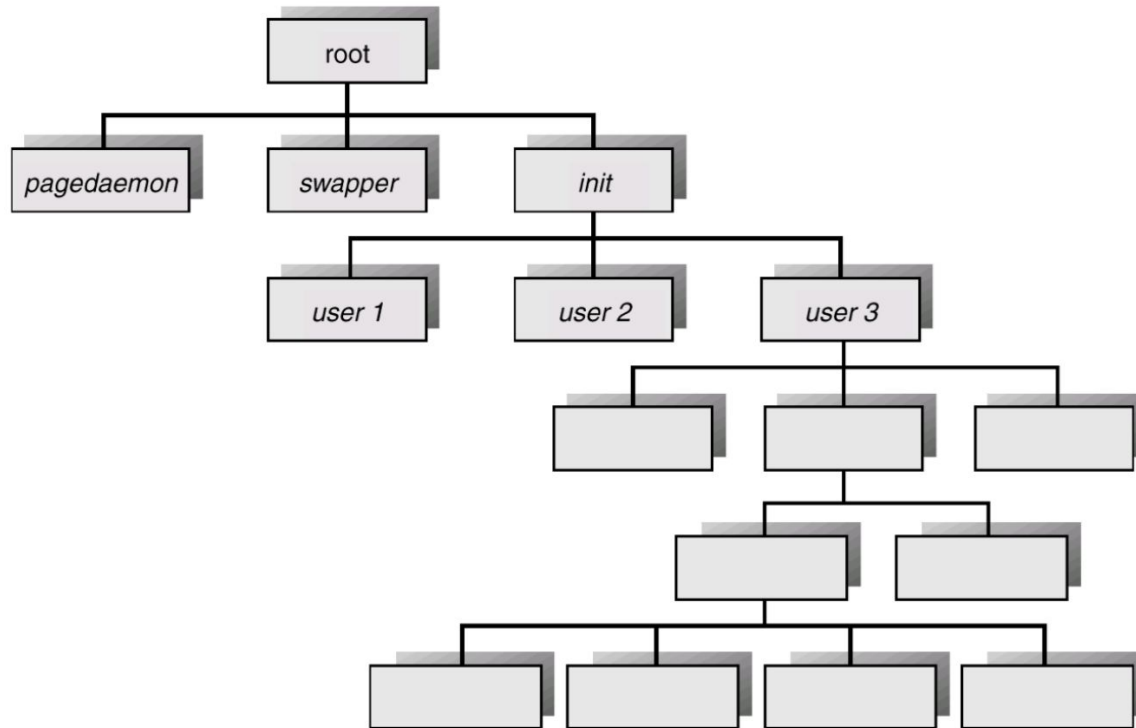
Process Creation

Some events that lead to process creation:

- The system **boots**
 - When a system is **initialized**, several background processes or “**daemons**” are started (email, logon, etc.)
- A user requests to **run** an application
 - By typing a command in the CLI shell or double-clicking in the GUI shell, the user can **launch** a new process
- An existing process spawns a **child** process
 - For example, a **server** process (e.g. web server, file server) may create a new process for each request it handles
 - The **init** daemon waits for user login and **spawns** a shell
- A batch system takes on the **next job** in line

Process Creation

Process creation by spawning. A **tree** of processes on a typical **UNIX** system:



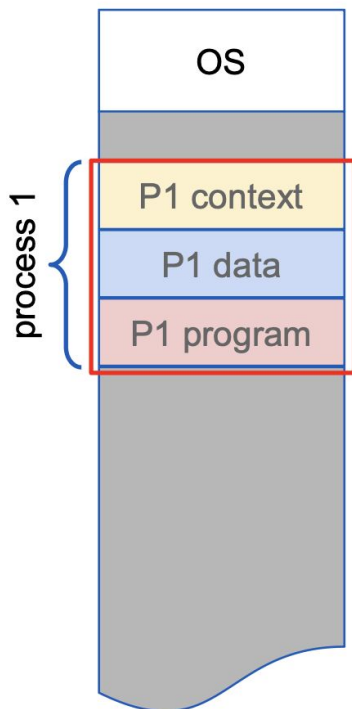
Process Creation: preliminaries

- `fork()` system call
 - `pid_t fork(void);`
 - Creates a new process in Unix systems, which is called the **child process**
 - Returns
 - **0**: On success to the child
 - **Child PID**: On success to the parent
 - **-1**: On failure and type in `errno`
- `exece()` system call
 - Family of functions replaces the current process image with a new process image
 - `int execvp(const char *file, char *const argv[]);`
 - Returns **-1** on failure

Process Creation

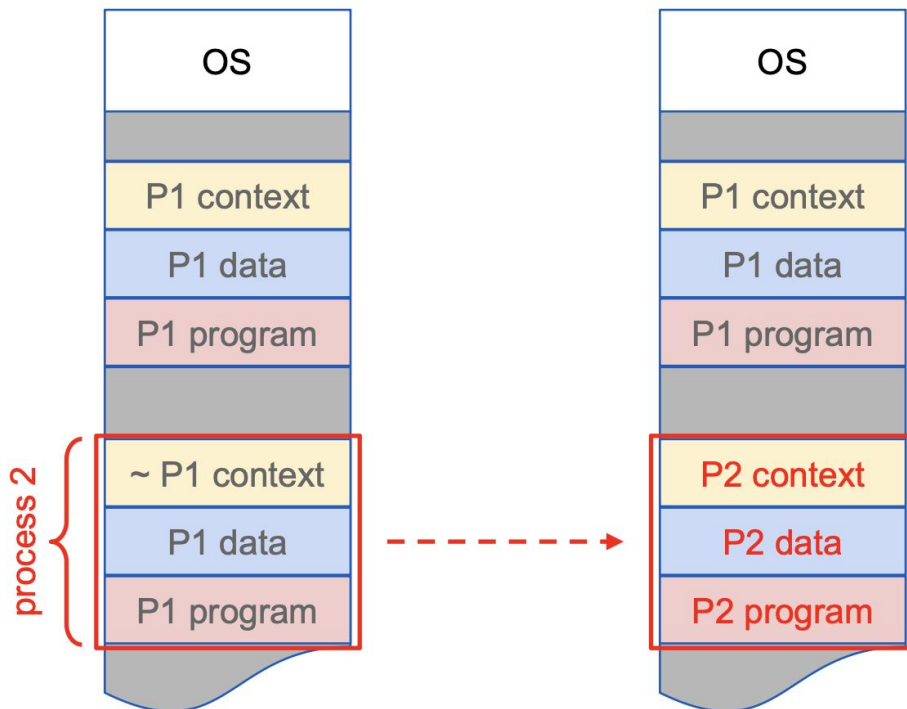
1. Clone child process

`pid = fork()`



2. Replace child's image

`execvp(name, ...)`



Fork Example 1

```
1  #include <stdio.h>
2  main()
3  {
4      int ret_from_fork, mypid;
5      mypid = getpid(); /* who am I? */
6      printf("Before: my pid is %d\n", mypid); /* tell pid */
7
8      ret_from_fork = fork();
9      sleep(1);
10     printf("After: my fork returns pid : %d, said %d\n",
11            |   |   |   |   |   |   |   ret_from_fork, getpid());
12 }
```

Process Creation

```
1  ...
2  int main(...)
3  {
4      ...
5      if ((pid = fork()) == 0) // create a process
6      {
7          fprintf(stdout, "Child pid: %i\n", getpid());
8          err = execvp(command, arguments); // execute child process
9          fprintf(stderr, "Child error: %i\n", errno);
10         exit(err);
11     }
12     else if (pid > 0) // we are in the parent
13     {
14         process
15         |   fprintf(stdout, "Parent pid: %i\n", getpid());
16         |   pid2 = waitpid(pid, &status, 0); // wait for child process
17         |   ...
18     }
19     ...
20     return 0;
21 }
```


Fork Example 2

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main()
5  {
6      fork();
7      fork();
8      fork();
9      printf("Process pid is %d\n", getpid());
10     return 0;
11 }
```

- How many lines of output will this code produce?

Process Termination

Some events that lead to process termination:

- Regular **completion**, with or without error code.
 - The process voluntarily executes an ***exit(err)*** system call to indicate to the OS that it has finished.
- **Fatal** error (uncatchable or uncaught)
 - Service errors: no memory left for allocation, I/O error, etc.
 - Total time limit exceeded
 - Arithmetic error, out-of-bound memory access, etc.
- **Killed** by another process via the **kernel**
 - The process receives a **SIGKILL** signal
 - In some systems the parent takes down its children with it.

Process Pause/Dispatch

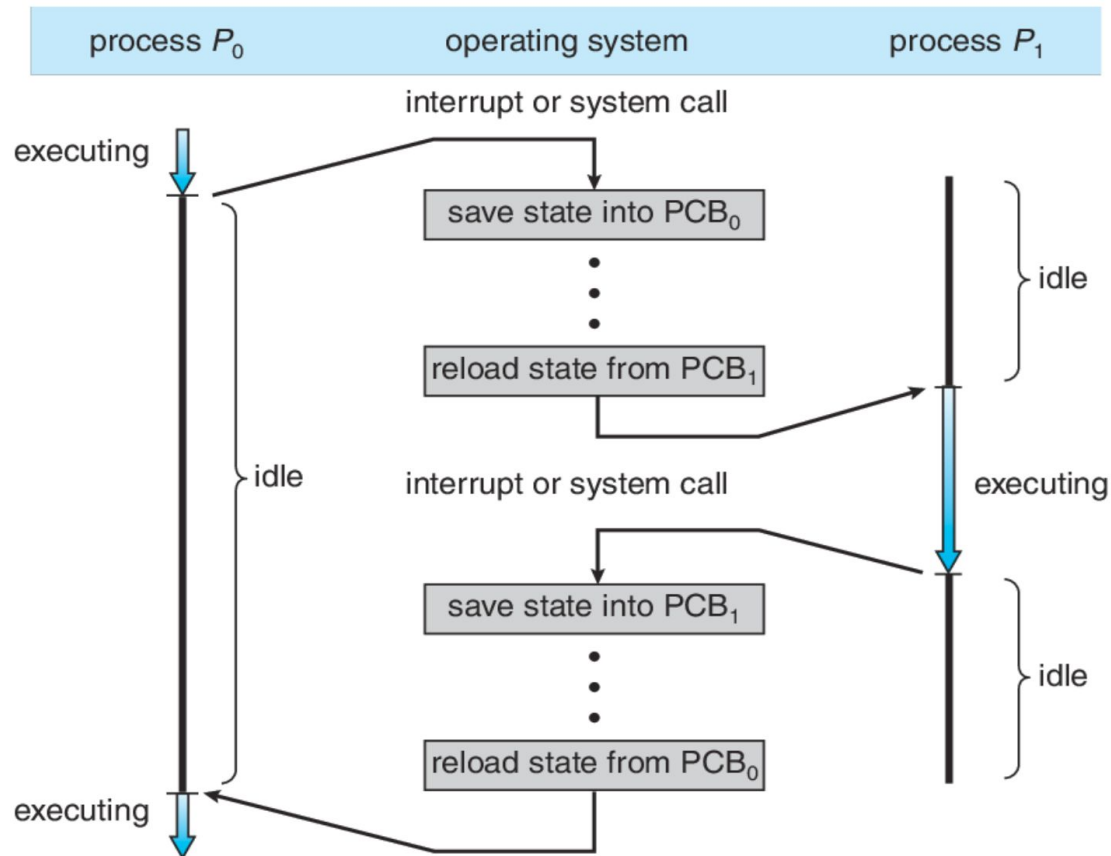
Some events that lead to process pause/dispatch:

- I/O wait - OS triggered (*e.g. system call*)
 - A process invokes an I/O system call that blocks waiting for the I/O device: The OS puts the process in “Waiting” mode and dispatches another process to the CPU.
- Preemptive timeout - Hardware interrupt triggered (*e.g. timer*)
 - The process receives a timer interrupt and relinquishes control back to the OS dispatcher: The OS puts the process in “Ready” mode and dispatches another process to the CPU.
 - Not to be confused with “**Total time limit exceeded**”, which leads to process termination.

Process Context-switching

- When CPU switches to another process, the system must **save the state** of the old process and **load** the saved state for the new process.
- Context-switch time is **overhead**; the system does no useful work while switching.
- Switching time is **dependent** on hardware support.

Process Context-switching



Process Context-switching, Step by Step

1. **Save** CPU context, including PC and registers (the only step needed in a simple mode switch)
2. Update process **state** (to “Ready”, “Blocked”, etc.) and other related fields of the PCB.
3. Move the PCB to the appropriate **queue**
4. Select another process for execution: This decision is made by the CPU **scheduling algorithm** of the OS.
5. Update the **PCB** of the selected process (change state to “Running”).
6. Update memory management **structures**
7. **Restore** CPU context to the values contained in the new PCB

Process Context-switching

What events trigger the OS to switch processes?

- **Interrupts:** External, **asynchronous** events, independent of the currently executing process instructions.
 - **Clock interrupt:** OS checks time and may block process
 - **I/O interrupt:** Data has come, OS may unblock process
 - **Memory fault:** OS may block process that must wait for a missing page in memory to be swapped in.
- **Exceptions (traps):** Internal, **synchronous** (but involuntary) events caused by instructions: OS may terminate or recover process.
- **System calls (traps!):** Voluntary **synchronous** events calling a specific OS service: After service completed, OS may either resume or block the calling process, depending on I/O, priorities, etc.


Process Queues

- **Job queue:** Set of all jobs in the system
- **Ready queue:** Set of all processes residing in main memory, ready and waiting to execute
- **Device queues:** Set of processes waiting for an I/O device

Processes migrate among the various queues

Process Queues

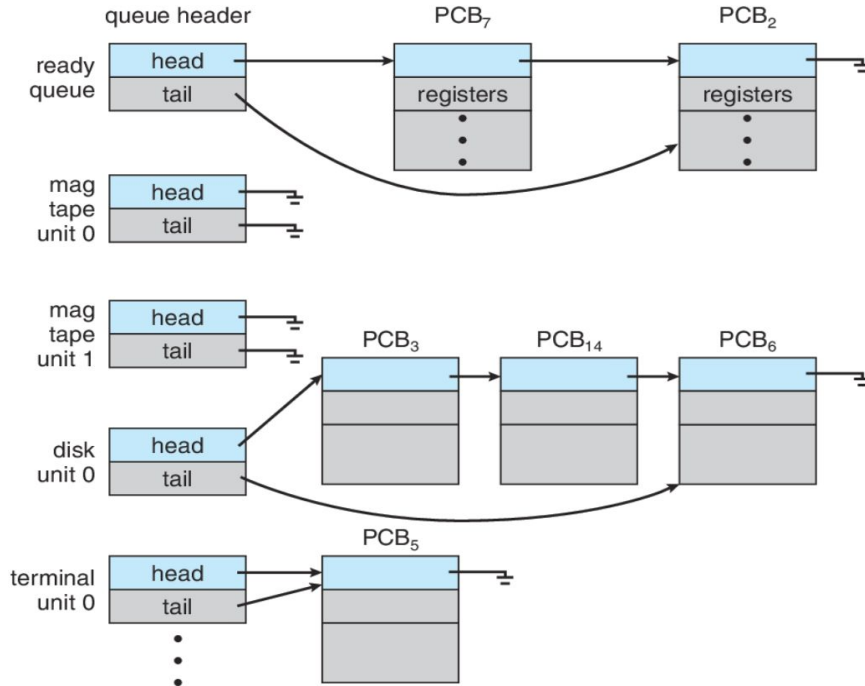
The process table can be **split** into per-state queues: PCBs can be linked together if they contain a pointer field



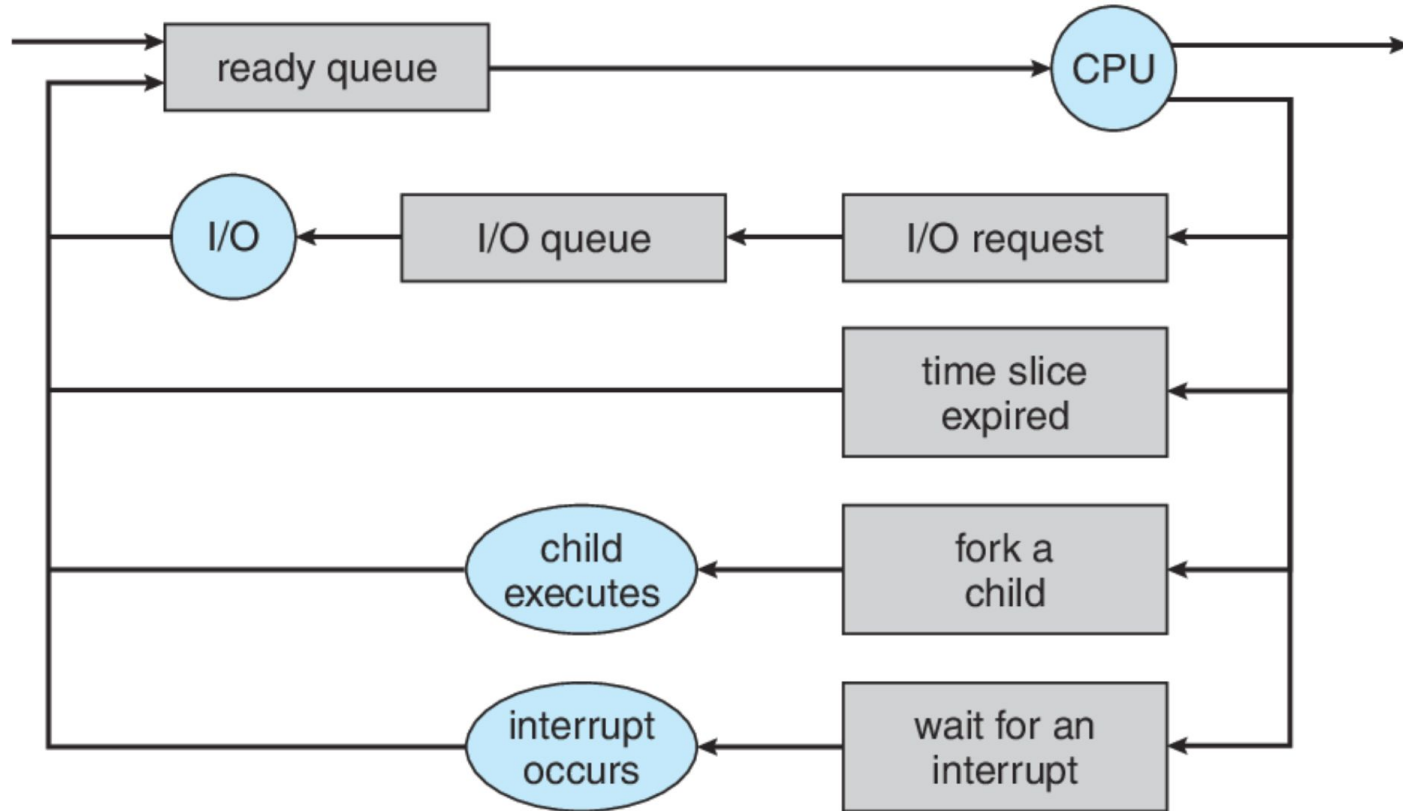
Process Queues

Ready Queue and various I/O device queues.

I/O scheduling: The decision to handle a process's pending I/O request



Representation of Process Scheduling

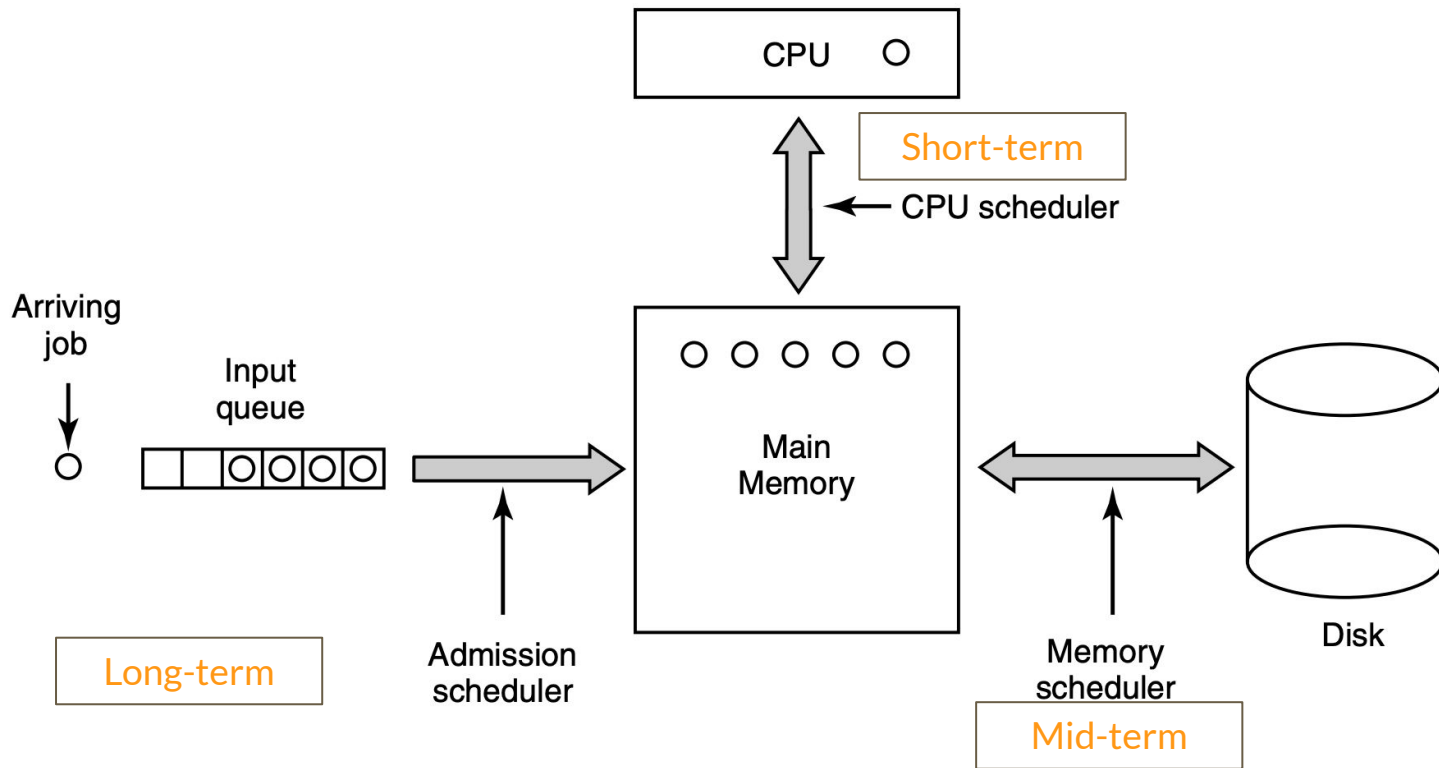


OS Scheduling

- Algorithm depends on the execution environment
 - All systems
 - Batch systems
 - Interactive systems
 - Real-time systems


OS Scheduling

Three Level Process Scheduling



OS Scheduling

- Long-term scheduling
 - The decision to add a program to executed (job scheduling)
- Medium-term scheduling
 - The decision to add to the number of processes that are partially or fully in main memory (“swapping”)
- Short-term scheduling (CPU scheduling)
 - The decision as to which available processes in memory are to be executed by the processor (“dispatching”)

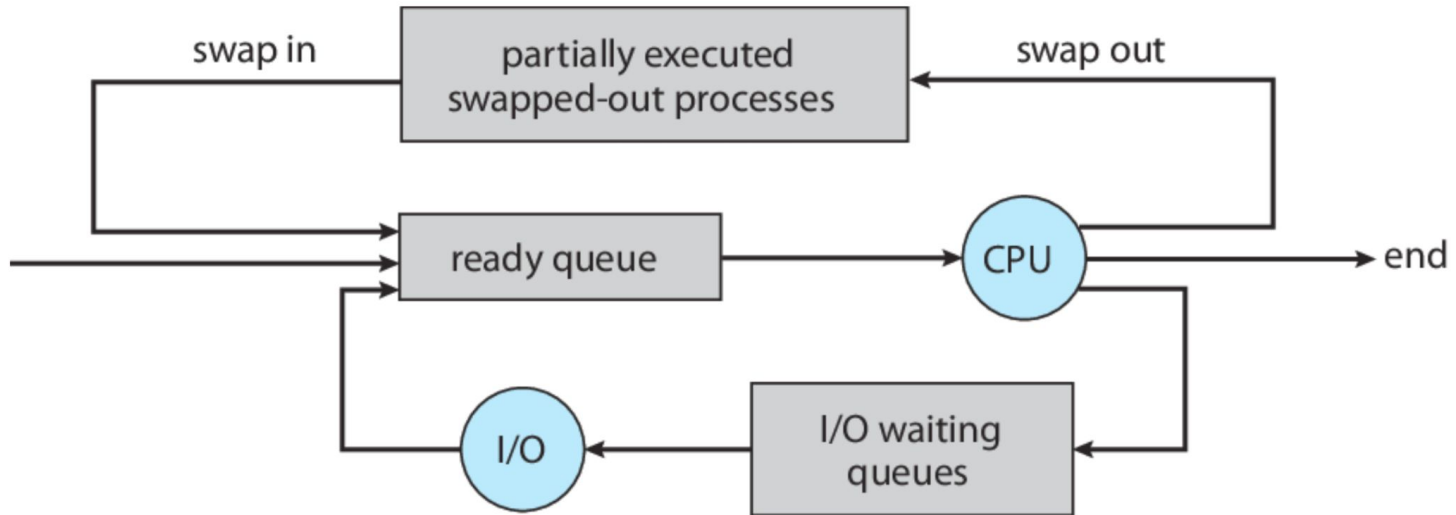


Schedulers

- Short-term scheduler is invoked very frequently (*milliseconds*) \Rightarrow (*must be fast*)
- Long-term scheduler is invoked very infrequently (*seconds, minutes*) \Rightarrow (*may be slow*)
- The long-term scheduler controls the degree of *multiprogramming*
- Processes can be described as either:
 - *I/O-bound process*: Spends more time doing I/O than computations, many short CPU bursts
 - *CPU-bound process*: Spends more time doing computations; few very long CPU bursts
- Long-term schedulers need to make careful decision

Addition of Medium Term Scheduling

- In time-sharing systems: remove processes from memory “temporarily” to reduce degree of multiprogramming.
- Later, these processes are resumed → **Swapping**



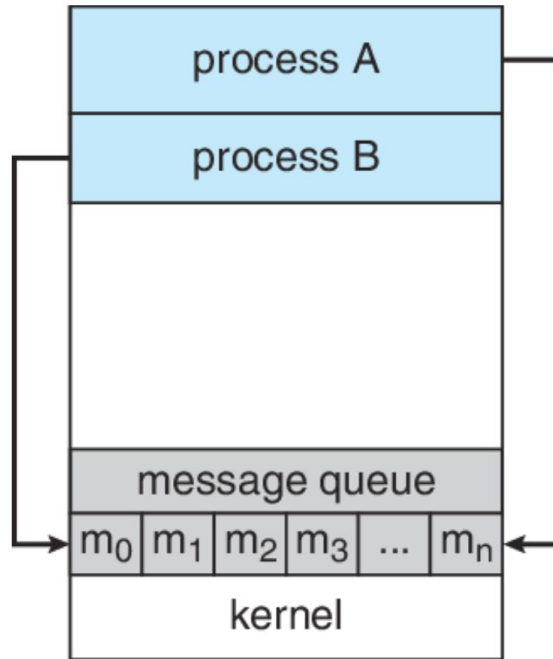
Cooperating Processes

- **Independent** process cannot affect or be affected by the execution of another process.
- **Cooperating** process can affect or be affected by the execution of another process.
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience
- Disadvantage
 - Synchronization issues and race conditions

Interprocess Communication (IPC)


- Mechanism for processes to communicate and to synchronize their actions
- **Shared Memory:** by using the same address space and shared variables
- **Message Passing:** processes communicate with each other without resorting to shared variables

Communications Models



(a)

Message Passing



(b)

Shared Memory

Message Passing

- Message passing facility provides two operations:
 - **send** (*message*) – message size fixed or variable
 - **receive** (*message*)
- If P and Q wish to communicate, they need to:
 - Establish a communication **link** between them
 - Exchange messages via send/receive
- Two types of message passing
 - **Direct** communication
 - **Indirect** communication

Message Passing - Direct Communication

- Processes must name each other explicitly:
 - **send** ($P, message$) – send a message to process P
 - **receive** ($Q, message$) – receive a message from process Q
- Properties of communication link:
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional
- Symmetrical vs Asymmetrical direct communication:
 - **send** ($P, message$) – send a message to process P
 - **receive** ($id, message$) – receive a message from any process
- **Disadvantage** of both: limited modularity, hardcoded

Message Passing - Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Primitives are defined as:
 - **send** (*A, message*) – send a message to mailbox A
 - **receive** (*A, message*) – receive a message from mailbox A

Message Passing - Indirect Communication

- Operations
 - Create a new mailbox
 - Send and receive messages through mailbox
 - Destroy a mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Message Passing - Indirect Communication

- Mailbox sharing
 - $P1$, $P2$, and $P3$ share mailbox A
 - $P1$, *sends*; $P2$ and $P3$ receive
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a “*receive*” operation
 - Allow the system to select the receiver arbitrarily. Sender is notified who the receiver was.

Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - **Blocking send:** The sender blocks until the message is received
 - **Blocking receive:** The receiver blocks until a message is available
- Non-blocking is considered **asynchronous**
 - **Non-blocking send:** The sender sends the message and continues
 - **Non-blocking receive:** The receiver receives a valid message or null

Buffering

- Queue of messages attached to the link; implemented in one of three ways:

- Zero capacity – 0 messages

Sender must wait for receiver

- Bounded capacity – finite length of n messages

Sender must wait if link full

- Unbounded capacity – infinite length

Sender never waits

Recap

- Processes
 - Process representation in OS
 - Process creation
 - Process termination
 - Context switching
 - Process queues
 - Process scheduling
 - Interprocess communication

Acknowledgements

- “Operating Systems Concepts” book and supplementary material by A. Silberschatz, P. Galvin and G. Gagne
- “Operating Systems: Internals and Design Principles” book and supplementary material by W. Stallings
- “Modern Operating Systems” book and supplementary material by A. Tanenbaum
- R. Doursat and M. Yuksel from University of Nevada, Reno
- Farshad Ghanei from Illinois Tech
- T. Kosar and K. Dantu from University at Buffalo