Operating Systems Concepts

System Calls in xv6

CS 4375, Fall 2025
Instructor: MD Armanuzzaman (*Arman*)
marmanuzzaman@utep.edu
September 22, 2025

Summery

- Operating system Services

- Operating system design choices

  - Direct execution protocol

  - Limited direct execution protocol

- System calls

  - Implementation of system calls

- Context switch between processes

Agenda

- Assembly programming

- Privileged CPU features

    - Registers and instructions

- xv6 system call

    - How does the whole cycle of system calls work?

    - Go over code snippets

Assembly programming

- Two main approaches:

  - Inline-assembly: assembly inside C source files

    - Compiler extension to call assembly instructions

    - Compiler still helps with some things automatically (e.g., allocating, saving, and restoring registers)

    - e.g., riscv.h

  - Assembly source: separate assembly file

    - Code written directly in assembly

    - Everything must be done manually by programmer, even calling conventions

    - e.g., trampoline.S

Extended ASM syntax

```
asm("assembly code" : outputs :
inputs : clobbers);
```

- **Outputs:** registers or memory that contain outputs

- **Inputs:** registers or memory that contain inputs

- **Clobbers:** registers or memory that are overwritten, not explicitly marked as

  inputs, and you don't care about their values

- **Volatile:** qualifier that disables certain optimizations

Extended ASM modifiers

- **Outputs:**
  - "=r": specifies a register (compiler decides which one)
  - "=m" specifies a memory location
  - "=rm" specifies either a register or memory location

- **Inputs:**
  - "r": specifies a register (compiler decides which one)
  - "m" specifies a memory location
  - "rm" specifies either a register or memory location

- **Clobbers:**
  - Specific registers
  - "memory" or "cc"

Extended ASM example

```
static inline void w_satp(uint64 x) {
  asm volatile("csrw satp, %0" : : "r" (x));
}
```

csrw <csr>, <src> = Control and Status Register Write.

- What does this code do?
- **%0:** specifies the 0th argument in outputs + inputs

Same example in .S file

```
csrw satp, a1
```

- No compiler help allocating, saving, or restoring registers

- Code manually specifies a1, overwriting its value

- Compiler cannot reorder instruction
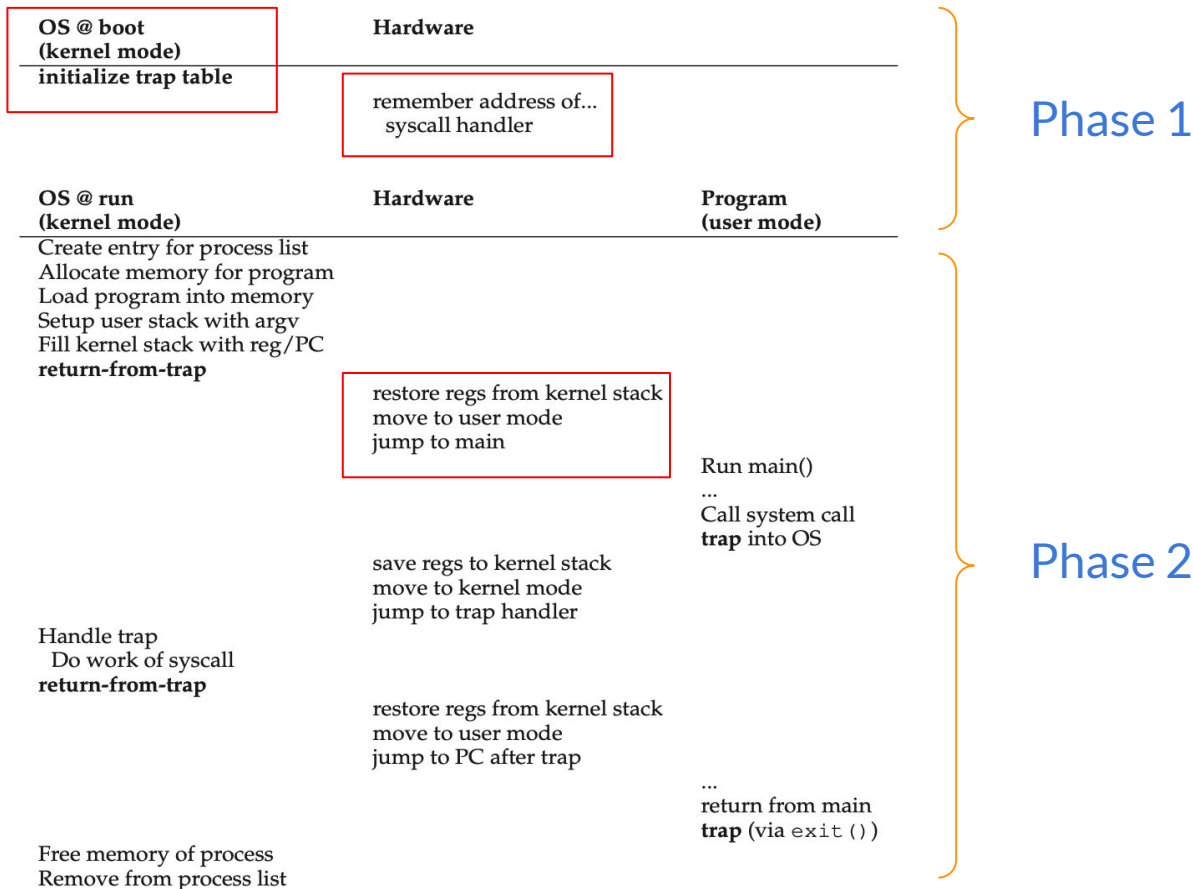
Privileged CPU features

- Two main parts:

  - Registers that can only be accessed by kernel

  - Instructions that can only be executed by kernel

- What do these look like in RISC-V?

If accessed/executed in user mode will generate a fault

Important privilege RISC-V features

- Privileged registers:

  - **satp:** physical address of page table root

  - **stvec:** ecall jumps here, points to trampoline

  - **sepc:** ecall saves the user's PC here

  - **sscratch:** scratch space; used to store temporary data

- Privileged instructions:

  - **Access regs:** csrr (read), csrw (write), csrrw (swap)
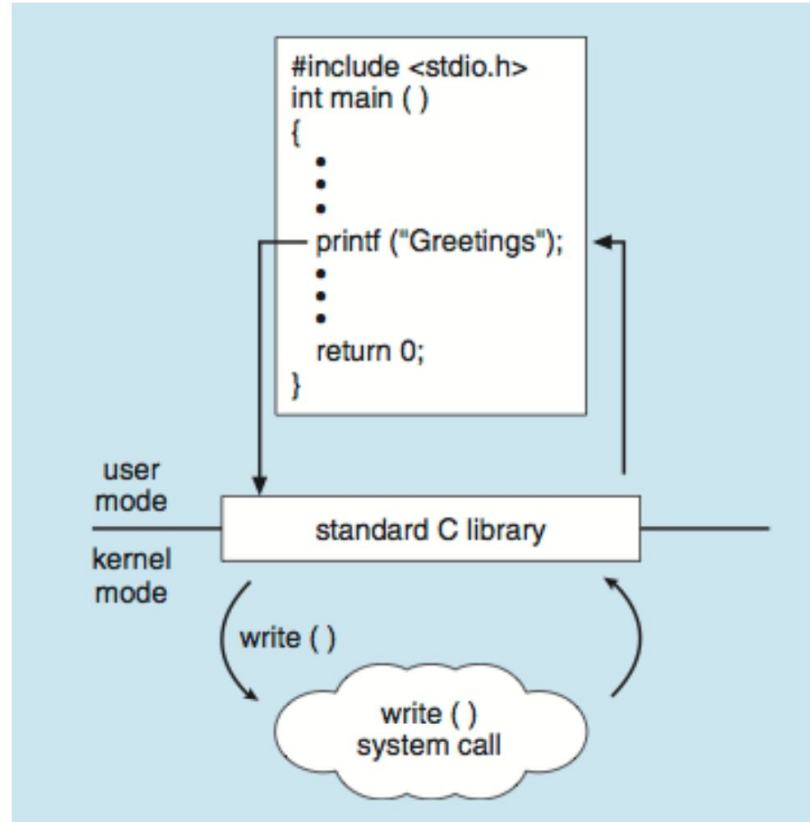
  - **sret:** return to userspace

Recall: LDE Protocol

OS @ boot (kernel mode)	Hardware	
initialize trap table		
	remember address of… syscall handler	

Phase 1

OS @ run (kernel mode)	Hardware	Program (user mode)
Create entry for process list		
Allocate memory for program		
Load program into memory		
Setup user stack with argv		
Fill kernel stack with reg/PC		
return-from-trap		
	restore regs from kernel stack move to user mode jump to main	
		Run main() … Call system call **trap** into OS
	save regs to kernel stack move to kernel mode jump to trap handler	
Handle trap Do work of syscall **return-from-trap**		
	restore regs from kernel stack move to user mode jump to PC after trap	
		… return from main **trap** (via exit())
Free memory of process Remove from process list		

Phase 2

11

RISCV Calling conventions

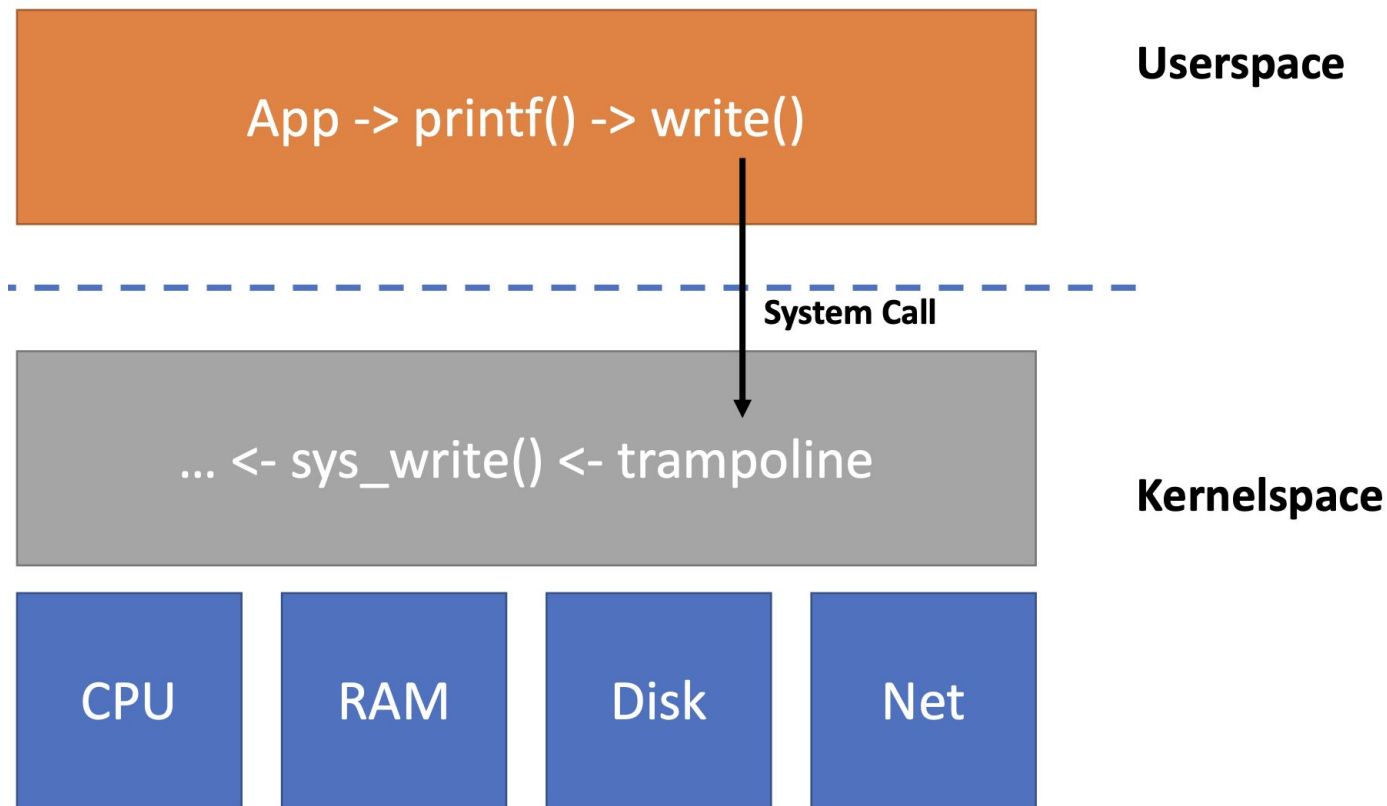| Register	ABI Name	Description	Saver
x0	zero	Hard-wired zero	—
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	—
x4	tp	Thread pointer	—
x5–7	t0–2	Temporaries	Caller
x8	s0/fp	Saved register/frame pointer	Callee
x9	s1	Saved register	Callee
x10–11	a0–1	Function arguments/return values	Caller
x12–17	a2–7	Function arguments	Caller
x18–27	s2–11	Saved registers	Callee
x28–31	t3–6	Temporaries	Caller
f0–7	ft0–7	FP temporaries	Caller
f8–9	fs0–1	FP saved registers	Callee
f10–11	fa0–1	FP arguments/return values	Caller
f12–17	fa2–7	FP arguments	Caller
f18–27	fs2–11	FP saved registers	Callee
f28–31	ft8–11	FP temporaries	Caller

User -> kernel transitions

- System calls, faults, and interrupts enter kernel same way

- Important for isolation and performance

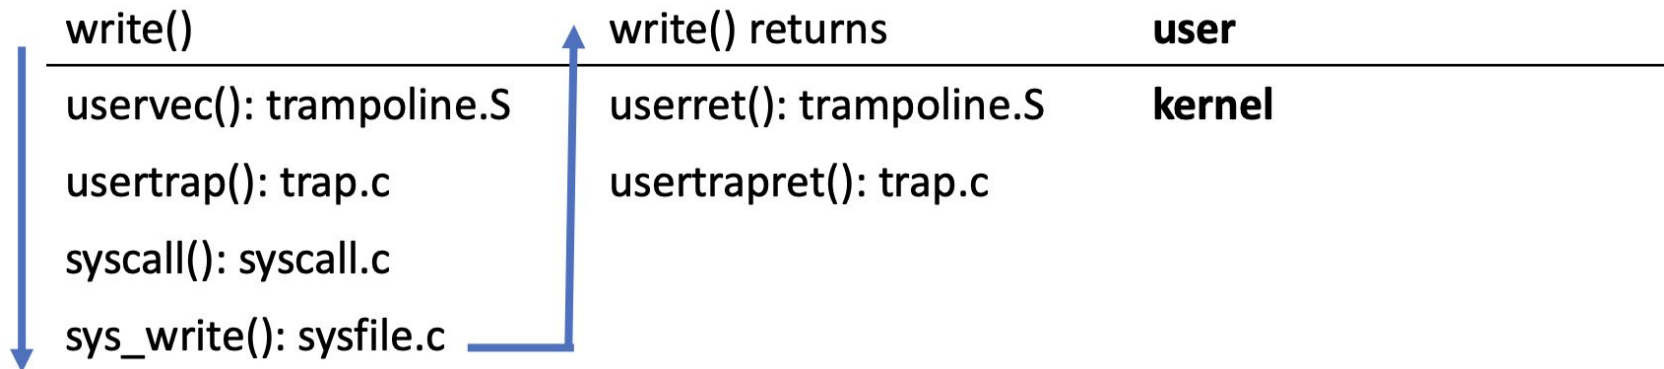- Lots of careful design choices; details matter!

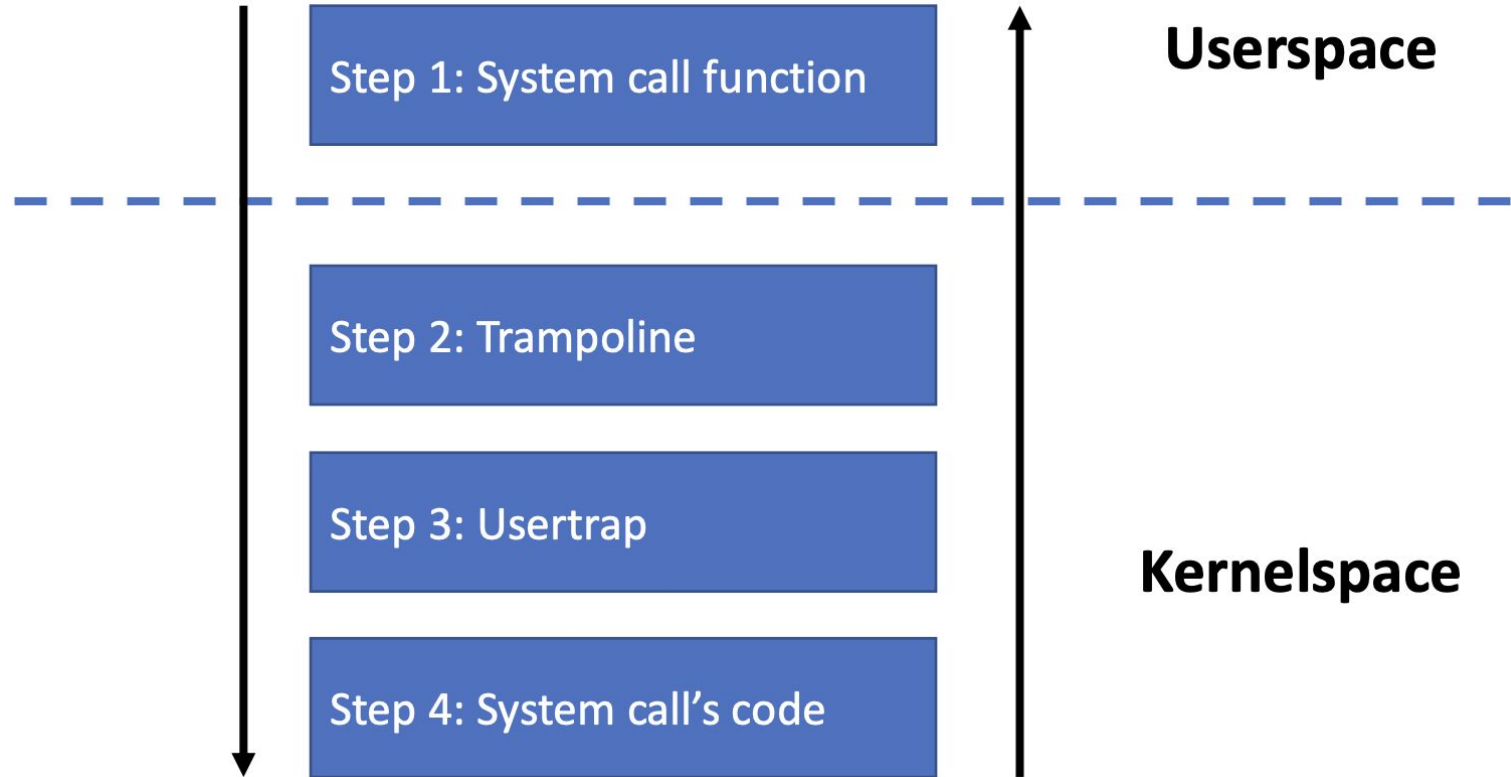# System Call Interface

System calls

App -> printf() -> write()

Userspace

System Call

... <- sys_write() <- trampoline

Kernelspace

CPU   RAM   Disk   Net

15

What needs to happen?

- Issue: CPU registers are set up for user, not kernel

- 32 registers, sp, pc, privilege mode, satp, stvec, sepc, …

- Steps needed:
    - Switch to supervisor mode
    - Save 32 registers + pc
    - Switch to kernel page table and stack
    - Jump to kernel C code

- User code can't interfere with transition

- Must be transparent, resume without disturbing user code (traps and interrupts can't be observable)

Example: write() system call

write()	write() returns	**user**
uservec(): trampoline.S	userret(): trampoline.S	**kernel**
usertrap(): trap.c	usertrapret(): trap.c	
syscall(): syscall.c		
sys_write(): sysfile.c		

Steps in a system call

Step 1: System call function

Step 2: Trampoline

Step 3: Usertrap

Step 4: System call's code

Userspace

Kernelspace

18

Step 1: System call function

- Use **ecall** instruction to enter supervisor mode

- **a7** register stores system call number. Why?

- System call's arguments stored in a0-a6

- Return value stored in **a0**

- No other registers changed after return! Why?

Step 1: System call func (usys.S)

```
…
.global write
write:
 li a7, SYS_write
 ecall
 ret
…
```

ecall does very little

- ecall performs just two actions:

  - Change to supervisor mode

  - Jump to **stvec** location

  - That's it!

- Why so simple?

  - Give designers flexibility to optimize their kernel

- Optimization ideas:

  - No pgtbl switch: Use one table for user and kernel

  - Don't use a stack for some simple system calls

  - Optimize which registers need to be saved

Step 2: Trampoline

- Must save registers to the trapframe. Why?

  - So they can be restored before returning

  - So the kernel can retrieve the syscall # and arguments

- Must restore the kernel's stack. Why?

  - So normal kernel code can execute

- Must restore the kernel's page table. Why?

  - Kernel's code and data aren't mapped in user pgtbl

- Must jump into usertrap(). Why?

  - Figures out how to handle different types of traps

Step 2: Trampoline (trampoline.S)

```
…
  # swap a0 and sscratch
  # so that a0 is TRAPFRAME
  csrrw a0, sscratch, a0
…
```
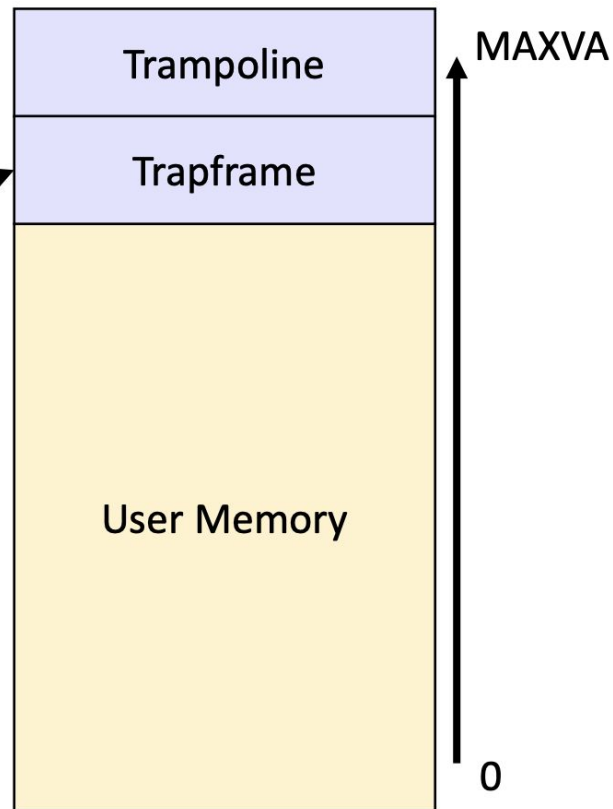
Contents of trapframe

```
struct trapframe {
 /*   0 */ uint64 kernel_satp;   // kernel page table
 /*   8 */ uint64 kernel_sp;     // top of process's kernel stack
 /*  16 */ uint64 kernel_trap;   // usertrap()
 /*  24 */ uint64 epc;           // saved user program counter
 /*  32 */ uint64 kernel_hartid; // saved kernel tp
 /*  40 */ uint64 ra;
 /*  48 */ uint64 sp;
 /*  56 */ uint64 gp;
 /*  64 */ uint64 tp;
 /*  72 */ uint64 t0;

 …
};
```

Step 2: Trampoline (trampoline.S)

```
# save the user registers in TRAPFRAME

        sd ra, 40(a0)

        sd sp, 48(a0)

        sd gp, 56(a0)

        sd tp, 64(a0)

        sd t0, 72(a0)

        sd t1, 80(a0)

        sd t2, 88(a0)

        sd s0, 96(a0)

        sd s1, 104(a0)

        sd a1, 120(a0)

        sd a2, 128(a0)

        sd a3, 136(a0)

        sd a4, 144(a0)

        sd a5, 152(a0)

        sd a6, 160(a0)

        sd a7, 168(a0)

        sd s2, 176(a0)

        sd s3, 184(a0)

        sd s4, 192(a0)

        sd s5, 200(a0)

        sd s6, 208(a0)

        sd s7, 216(a0)

        sd s8, 224(a0)

        sd s9, 232(a0)

        sd s10, 240(a0)

        sd s11, 248(a0)

        sd t3, 256(a0)

        sd t4, 264(a0)

        sd t5, 272(a0)

        sd t6, 280(a0)

        # save the user a0 in p->trapframe->a0

        csrr t0, sscratch

        sd t0, 112(a0)
```

Step 2: Trampoline (trampoline.S)

```
…
# restore kernel stack pointer from p->trapframe->kernel_sp
ld sp, 8(a0)
# make tp hold the current hartid, from p->trapframe->kernel_hartid
ld tp, 32(a0)
# load the address of usertrap(), p->trapframe->kernel_trap
ld t0, 16(a0)
# restore kernel page table from p->trapframe->kernel_satp
ld t1, 0(a0)
csrw satp, t1
sfence.vma zero, zero
# jump to usertrap(), which does not return
jr t0
```

Step 2: Trampoline


```
csrw satp, t1
```

Change page table!
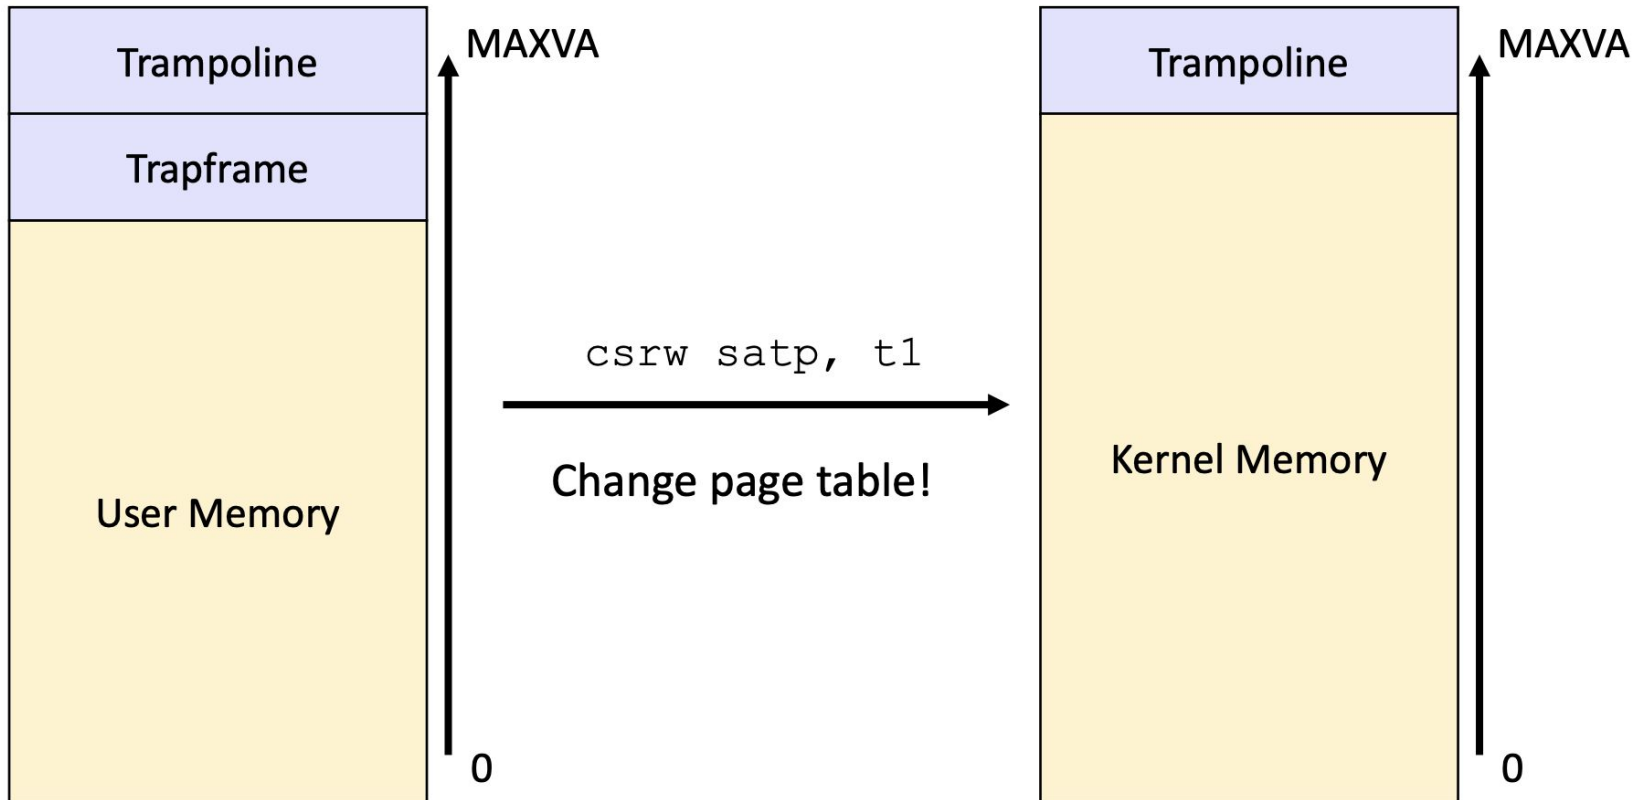
Step 3: Usertrap

- Entry point into kernel C code

  - Handles system calls, traps, and interrupts

  - Must figure which it is to dispatch to right part of kernel

- r_scause() helps usertrap() figure this out

  - A value of 8 indicates system call

Step 4: System call

- Must determine which system call

- Kernel maintains table of function pointers

- **a7** register value in **Trapframe** determines index in table

- e.g., sys_write()

Summary

- System calls are much more complex than func calls

- Major reasons:

  - Requirement of isolation

  - Desire for simple, fast, and flexible hardware

- Questions to consider for entry/exit

  - Can an evil program abuse the entry mechanism?

  - Can you think of ways to make hardware or software simpler?

  - Can you think of ways to make traps faster?

Announcement

- Homework 1

  - Due Monday September 22nd

  - **DON'T FORGET TO GIVE GITHUB REPO ACCESS PERMISSION TO THE TA**

    - danielmarin350@gmail.com

- Quiz 2

  - Next class

- Homework 2

  - Will be released today

  - Due on monday october 1st, 11.59 PM