

# Operating Systems Concepts

CPU Scheduling



CS 4375, Fall 2025

Instructor: MD Armanuzzaman (*Arman*)

[marmanuzzaman@utep.edu](mailto:marmanuzzaman@utep.edu)

September 24, 2025

# Summery

- Assembly programming
- Privileged CPU features
  - Registers and instructions
- xv6 system call
  - How does the whole cycle of system calls work?
  - Go over code snippets

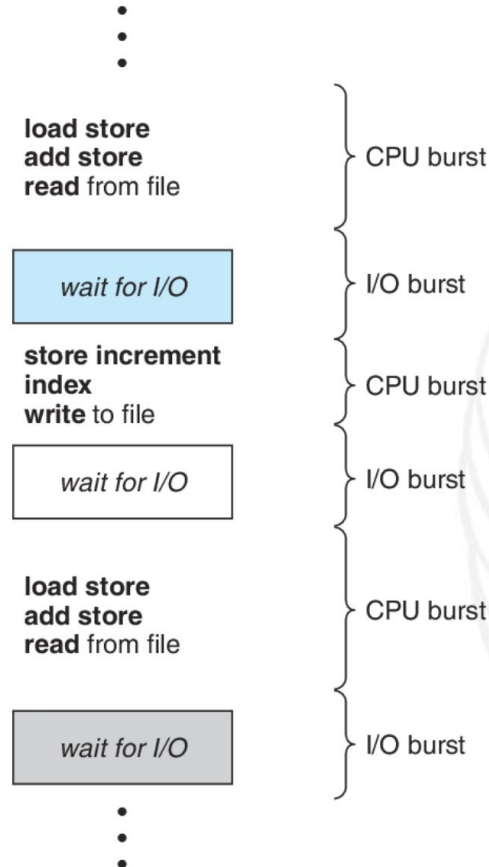
# Agenda

- CPU Scheduling:
  - Basic concepts
  - Scheduling Criteria & Metrics
  - Different Scheduling Algorithms
    - FCFS
    - SJF
    - Priority
    - RR
  - Preemptive vs Non-preemptive Scheduling
  - Gantt Charts & Performance Comparison

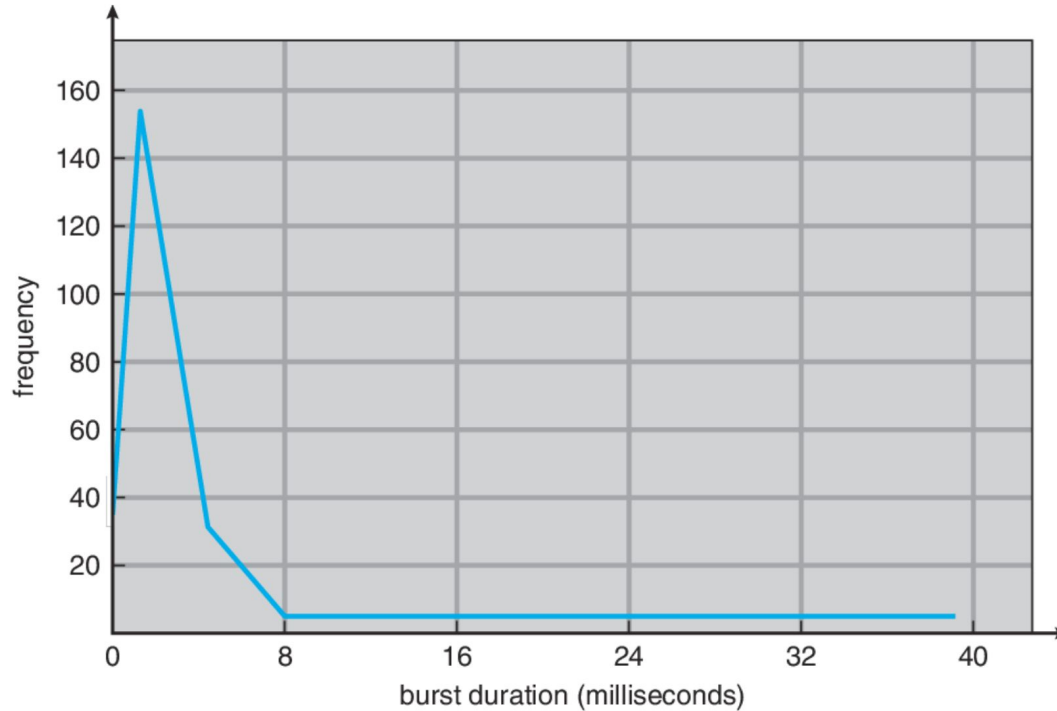
# Basic Concepts

- Multiprogramming is needed for efficient CPU utilization
- **CPU Scheduling**: deciding when to execute which processes
- Process execution begins with a **CPU burst**, followed by an **I/O burst**
- CPU-I/O burst cycle:
  - Process execution consists of a *cycle* of **CPU execution** and **I/O wait**

# Alternating Sequence of CPU And I/O Bursts

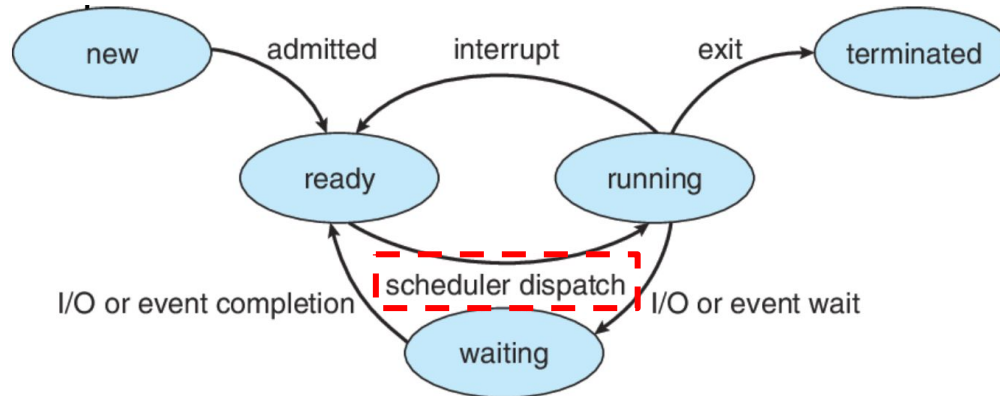


# Histogram of CPU-burst Durations



# Process State

- A process changes its **state** during execution:
  - **New**: The process is being created
  - **Ready**: The process is waiting to be assigned to a processor
  - **Running**: Instructions are being executed
  - **Waiting**: The process is waiting for some event to occur
  - **Terminated**: The process has finished execution



# CPU Scheduler

- Selects from among the processes in memory that are **ready** to execute, and allocates the CPU to one of them  $\Rightarrow$  short-term scheduler
- CPU scheduling decisions may take place when a process:
  1. A new process arrives *preemptive*
  2. Switches from running to ready state *preemptive*
  3. Switches from waiting to ready *preemptive*
  4. Switches from running to waiting state *non-preemptive/cooperative*
  5. Terminates *non-preemptive/cooperative*



# CPU Scheduler

- Scheduling under 4 (wait) and 5 (termination) is *non-preemptive/cooperative*
  - Once a process gets the CPU, keeps it until termination/switching to waiting state/release of the CPU
- All other schedulings are *preemptive*
  - Most OSs use this
  - e.g. time quota expires
  - Cost associated with access to shared data

# Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler.
- It's function involves:
  - Switching context
  - Switching to user mode
  - Jumping to the proper location in the user program to resume that program
- **Dispatch latency:** Time it takes for the dispatcher to stop one process and start another one.

# Scheduling Criteria

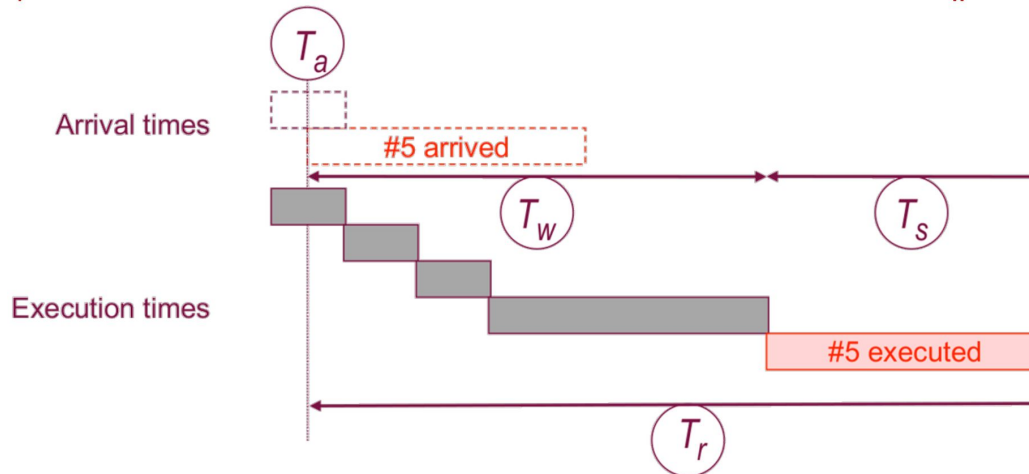
- **CPU Utilization:** Keep the CPU as busy as possible **Maximize**
- **Throughput:** # of processes that complete their execution per time unit **Maximize**
- **Response Time:** Amount of time it takes from when a request was submitted until the first response (not output) is produced (for time-sharing environment) **Minimize**
- **Waiting Time:** Total amount of time a process has been waiting in the ready queue **Minimize**
- **Turnaround Time:** Amount of time passed to finish execution of a particular process i.e. execution(service) time + waiting time **Minimize**

# Optimization Criteria

- Maximize CPU Utilization
- Maximize Throughput
- Minimize Response Time
- Minimize Waiting Time
- Minimize Turnaround Time

# Scheduling Metrics

- $T_a$  - **Arrival time**: Time the process became “READY” [again]
- $T_w$  - **Waiting time**: Time spent waiting for the CPU
- $T_s$  - **Service time**: Time spent executing in the CPU
- $T_r$  - **Turnaround time**: Time spent waiting and executing =  $T_w + T_s$

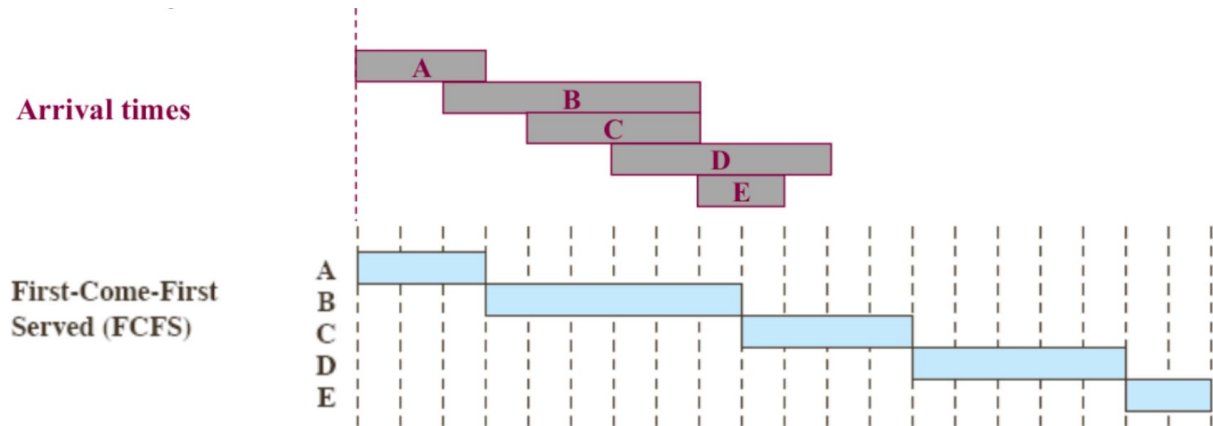


$$T_r / T_s = 2.5$$

# Scheduling: First-Come, First-Served (FCFS)

- Processes are assigned the CPU in the order they request it
- When the running process blocks, the first “READY” process is selected
- When a process gets “READY”, it is put at the end of the queue

FCFS  
Scheduling  
Policy



FCFS	Finish Time	<b>A</b>	3	<b>B</b>	9	<b>C</b>	13	<b>D</b>	18	<b>E</b>	20	<b>Mean</b>
	Turnaround Time ( $T_T$ )		3		7		9		12		12	8.60
	$T_T/T_S$		1.00		1.17		2.25		2.40		6.00	2.56

# Scheduling: FCFS Example

- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$

Process	Burst Time
$P_1$	24
$P_2$	3
$P_3$	3

- The Gantt chart for the schedule is:



- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27) / 3 = 17$

# Scheduling: FCFS Example

- Suppose that the processes arrive in the order:  $P_2, P_3, P_1$
- The **Gantt chart** for the schedule is:




- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3) / 3 = 3$
- Much better than previous case



# Scheduling: FCFS Example

- Suppose that the processes arrive in the order:  $P_2, P_3, P_1$
- The **Gantt chart** for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3) / 3 = 3$
- Much better than previous case

**Convoy effect:** short process behind long process

# Scheduling: FCFS Example



**Convoy effect:** short process behind long process

# Scheduling: Shortest-Job-First (SJF)

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
  - **Non-preemptive** – once CPU given to the process it cannot be preempted until completes its CPU burst
  - **Preemptive** – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt.
    - This scheme is known as the **Shortest-Remaining-Time-First (SRTF)**
- **Assumes the run times are known in advance!**
- SJF(SRTF) is optimal – gives minimum average waiting time for a given set of processes


# Scheduling: Non-Preemptive SJF

- Among several equally important “**READY**” jobs (or CPU bursts), the scheduler picks the one that will finish the earliest

SJF  
Scheduling  
Policy

Arrival times

Shortest Job  
First (SJF)

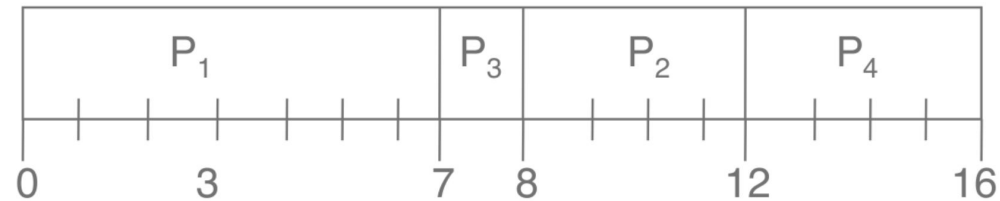


SJF	Finish Time	A	3	B	9	C	15	D	20	E	11	Mean
	Turnaround Time ( $T_T$ )		3		7		11		14		3	7.60
	$T_T/T_S$		1.00		1.17		2.75		2.80		1.50	1.84

# Scheduling: Non-Preemptive SJF Example

Process	Arrival Time	Burst Time
P <sub>1</sub>	0	7
P <sub>2</sub>	2	4
P <sub>3</sub>	4	1
P <sub>4</sub>	5	4


- The Gantt chart for the schedule is:



- Average waiting time:  $(0 + 6 + 3 + 7) / 4 = 4$

# Scheduling: Preemptive SJF

- Choose the process whose **remaining** run time is shortest.
- Allows new short jobs to get good service



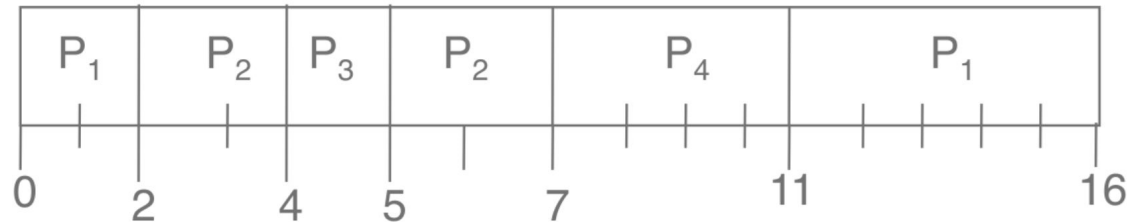
SRT  
Scheduling  
Policy

SRT	Finish Time	A	3	B	15	C	8	D	20	E	10	Mean
	Turnaround Time ( $T_r$ )		3		13		4		14		2	7.20
	$T_r/T_s$		1.00		2.17		1.00		2.80		1.00	1.59

# Scheduling: Preemptive SJF Example

Process	Arrival Time	Burst Time
P <sub>1</sub>	0	7
P <sub>2</sub>	2	4
P <sub>3</sub>	4	1
P <sub>4</sub>	5	4

- The Gantt chart for the schedule is:



- Average waiting time:  $(9 + 1 + 0 + 2) / 4 = 3$

# Scheduling: Priority

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (**smallest/largest** number)
- Two schemes:
  - **Non-preemptive** – Once CPU given to the process it cannot be preempted until completes its CPU burst
  - **Preemptive** – If a new process arrives with higher priority, preempt
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- **Problem: Starvation**– Low priority processes may never execute
- **Solution: Aging** – As time progresses, increase the priority of the process



# Scheduling: Priority Example

Process	Arrival Time	Burst Time	Priority
P <sub>1</sub>	0	7	2
P <sub>2</sub>	2	4	1
P <sub>3</sub>	4	1	4
P <sub>4</sub>	5	4	3

- Non-preemptive:
  - $P_1 \rightarrow P_2 \rightarrow P_4 \rightarrow P_3$
- Preemptive:




# Scheduling: Round-Robin (RR)

- Each process gets a small unit of CPU time (i.e. *time quantum*), usually 10-100 milliseconds
- After this time has elapsed, the process is preempted and added to the end of the ready queue
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. No process waits more than  $(n-1)*q$  time units
- Performance:
  - Large  $q \rightarrow$  FCFS
  - Small  $q \rightarrow q$  must be large enough with respect to context switch time, otherwise overhead would be too high

# Scheduling: Round-Robin (RR)

- Preemptive FCFS, based on a timeout interval, the **time quantum (TIME\_SLICE)**
- The running process is interrupted by the clock and put last in a FIFO **“READY”** queue

RR ( $q = 1$ )  
Scheduling  
Policy

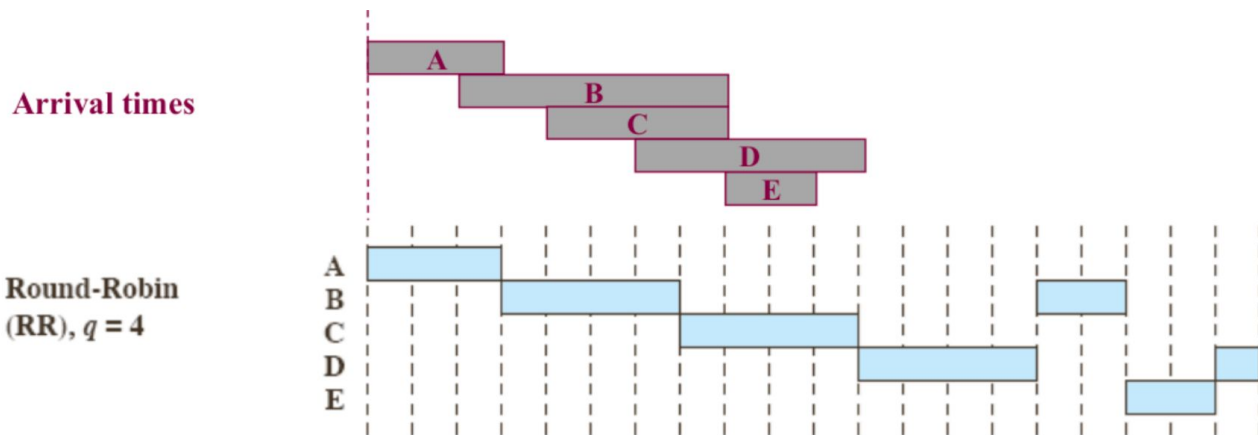


RR $q = 1$	Finish Time	<b>A</b>	4	<b>B</b>	18	<b>C</b>	17	<b>D</b>	20	<b>E</b>	15	<b>Mean</b>	
	Turnaround Time ( $T_r$ )		4		16		13		14		7		10.80
	$T_r/T_s$		1.33		2.67		3.25		2.80		3.50		2.71

# Scheduling: Round-Robin (RR)

- A crucial parameter is the quantum  $q$  ( $\sim 10\text{-}100\text{ms}$ )
  - $q$  should be large compared to context switch latency ( $\sim 10\mu\text{s}$ )
  - $q$  should be less than the longest CPU burst, or RR **degenerates** to FCFS

RR ( $q = 4$ )  
Scheduling  
Policy



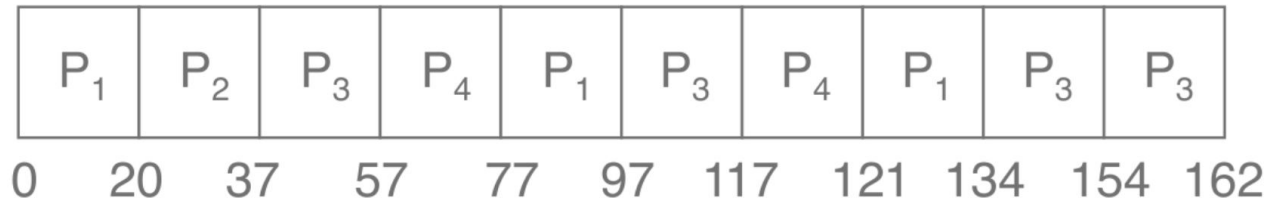
RR $q = 4$	Finish Time	<b>A</b>	3	<b>B</b>	17	<b>C</b>	11	<b>D</b>	20	<b>E</b>	19	<b>Mean</b>
	Turnaround Time ( $T_T$ )		3		15		7		14		11	10.00
	$T_T/T_S$		1.00		2.5		1.75		2.80		5.50	2.71

# Scheduling: RR Example

- Consider  $q = 20$

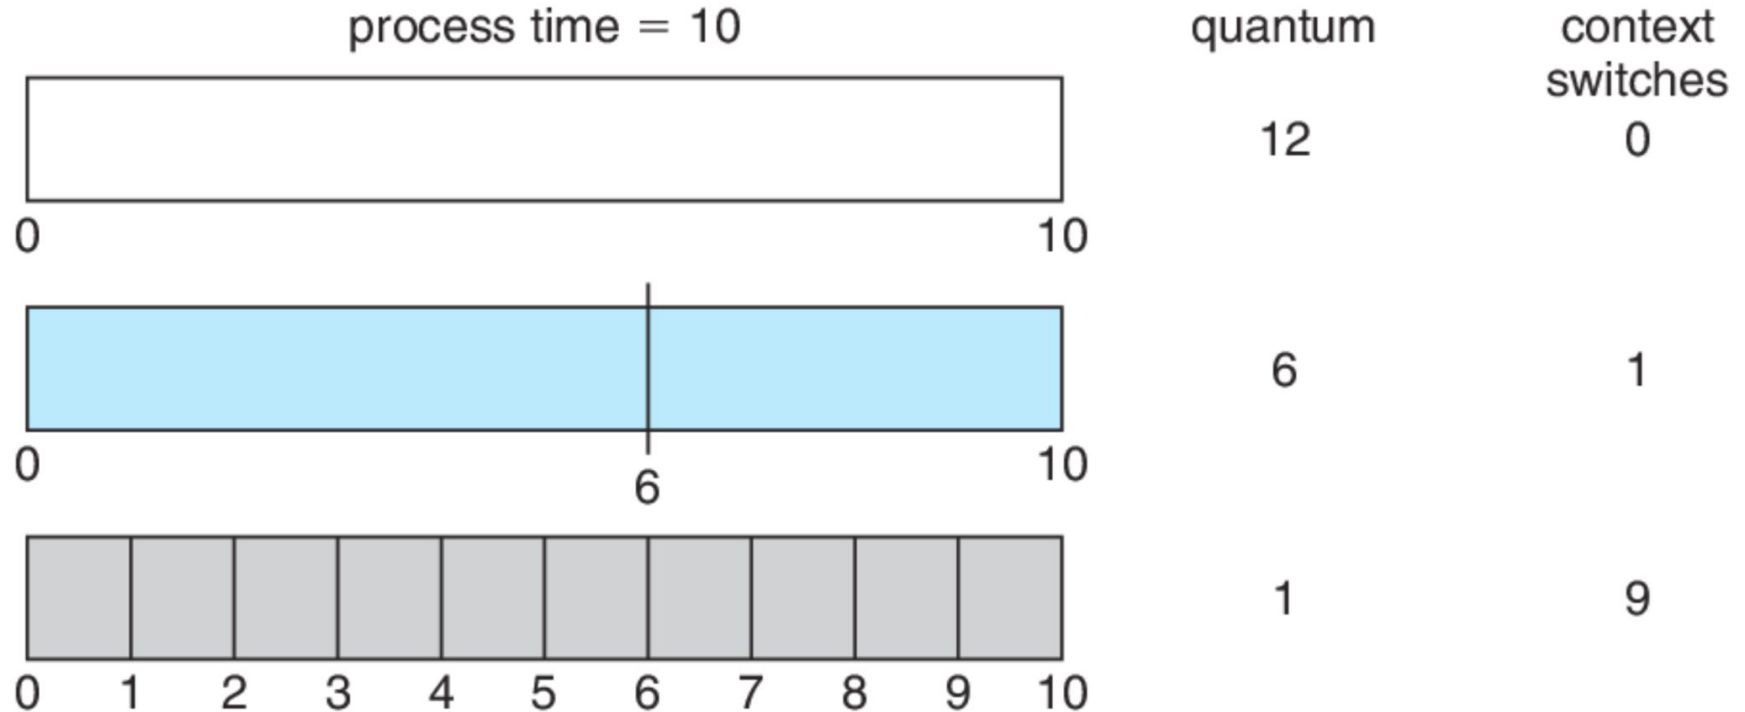
Process	Burst Time
$P_1$	53
$P_2$	17
$P_3$	68
$P_4$	24

- The Gantt chart for the schedule is:

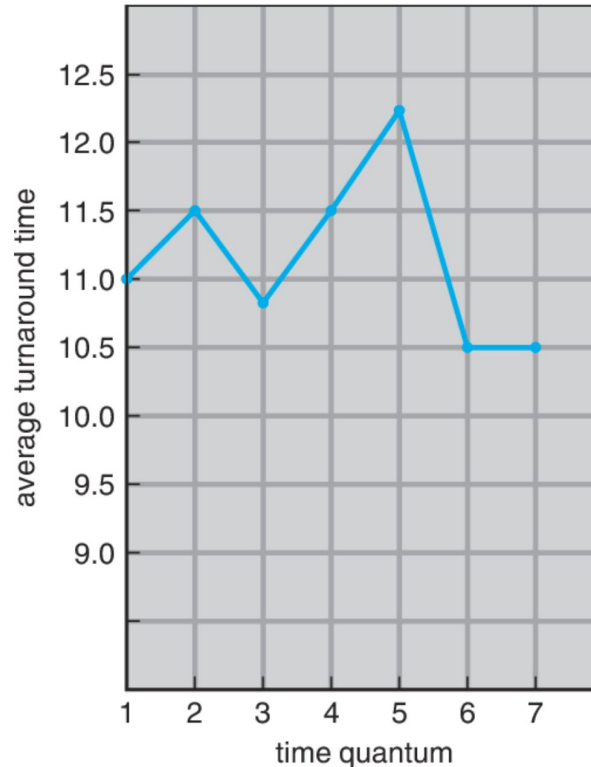


- Typically, higher average turnaround than SJF, but better *response*

# Time Quantum and Context Switch Time



# Turnaround Time Varies With The Time Quantum



process	time
$P_1$	6
$P_2$	3
$P_3$	1
$P_4$	7

# Exercise

- Consider a system using round-robin scheduling with a fixed quantum  $q$ . Every context switch takes  $s$  milliseconds. Any given process runs for an average of  $t$  milliseconds before it blocks or terminates (burst time)
- Determine the **fraction of CPU time** that will be wasted because of context switches for each of the following cases (Your answer should be in terms of  $q$ ,  $s$ , and  $t$ )
  - $t \leq q$
  - $t \gg q$  ( $t$  is much greater than  $q$ )
  - $q \rightarrow 0$

*Context switching time*

---

*Context switching time + useful CPU burst*



# Exercise

- Consider a system using round-robin scheduling with a fixed quantum  $q$ . Every context switch takes  $s$  milliseconds. Any given process runs for an average of  $t$  milliseconds before it blocks or terminates (burst time)
- Determine the **fraction of CPU time** that will be wasted because of context switches for each of the following cases (Your answer should be in terms of  $q$ ,  $s$ , and  $t$ )
  - $t \leq q$   
 $s / (t + s)$
  - $t \gg q$  ( $t$  is much greater than  $q$ )  
 $s / (q + s)$
  - $q \rightarrow 0$   
100%

# Comparison: FCFS

- PROS:

- It is a fair algorithm
  - Schedule in the order that they arrive

- CONS:

- Average response time can be lousy
  - Small requests wait behind big ones (convoy effect)
- May lead to poor utilization of other resources
  - FCFS may result in poor overlap of CPU and I/O activity
    - E.g., a CPU-intensive job prevents an I/O intensive job from doing a small bit of computation, thus preventing it from going back and keep the I/O subsystem busy

# Comparison: SJF

- PROS:

- Provably optimal with respect to average waiting time
  - Prevents convoy effect

- CONS:

- Can cause starvation of long jobs
- Requires advanced knowledge of CPU burst times
  - This is not easy to predict!

# Comparison: Priority

- PROS:

- Guarantees early completion of high priority jobs

- CONS:

- Can cause starvation of low priority jobs
- How to decide/assign priority value?

# Comparison: Round-Robin

- PROS:

- Great for timesharing
  - No starvation
- Does not require prior knowledge of CPU burst times
- Generally reduces average response time

- CONS:

- What if all jobs are almost the same length
  - Increases the turnaround time
- How to set the “best” time quantum?
  - If too small, it increases context-switch overhead
  - If too large, response time degrades

# Exercise

Process	Arrival Time	Burst Time	Priority
P <sub>1</sub>	9	3	5
P <sub>2</sub>	6	4	4
P <sub>3</sub>	8	5	3
P <sub>4</sub>	7	9	2
P <sub>5</sub>	5	7	1

- Consider Shortest-Remaining-Time-First (SRTF) scheduling policy is used.
  - a. Draw the Gantt chart
  - b. Find the response time, waiting time, and turnaround time for each process.

# Exercise

Process	Arrival Time	Burst Time	Priority
$P_1$	0	20	3
$P_2$	5	15	1
$P_3$	10	10	2
$P_4$	15	5	4

- Draw Gantt charts, find average turnaround time, waiting time, response time for each processor, for each of FCFS, SJF (non-preemptive), SRTF (preemptive), RR ( $q=4$ ), Priority (preemptive and non-preemptive)

# Announcement

- Quiz 2 (Released on Blackboard)
  - Due at 11.59 PM tomorrow
  - 30 minutes, single attempt, 15 MCQs
- Homework 2
  - Due on Wednesday october 1st, 11.59 PM