

Operating Systems Concepts

xv6 scheduling



CS 4375, Fall 2025

Instructor: MD Armanuzzaman (*Arman*)

marmanuzzaman@utep.edu

October 8, 2025

Summary

- Main Memory:
 - Fixed and Dynamic Memory Allocation
 - External and Internal Fragmentation
 - Address Binding
 - Hardware Address Protection
 - Paging
 - Segmentation


Agenda

- xv6 scheduling
 - Default round robin scheduling of xv6
 - Understand the code base
- Implement priority scheduling in xv6
 - Tasks in homework 3
- Implement aging policy

Scheduling: Round-Robin (RR)

- A crucial parameter is the quantum **q** (~10-100ms)
 - **q** should be large compared to context switch latency (~10μs)
 - **q** should be less than the longest CPU burst, or RR **degenerates** to FCFS

RR ($q = 4$) Scheduling Policy



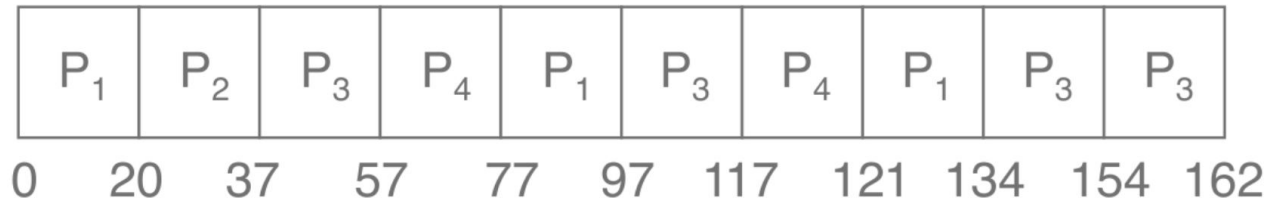
RR $q = 4$	Finish Time	A 3	B 17	C 11	D 20	E 19	Mean
	Turnaround Time (T_r)	3	15	7	14	11	10.00
	T_r/T_s	1.00	2.5	1.75	2.80	5.50	2.71

Scheduling: RR Example

- Consider $q = 20$

Process	Burst Time
P_1	53
P_2	17
P_3	68
P_4	24

- The **Gantt chart** for the schedule is:



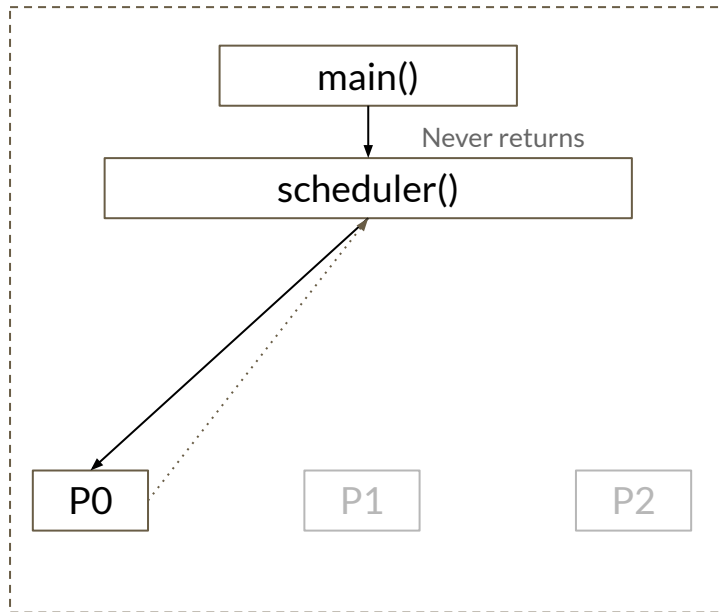
- Typically, higher average turnaround than SJF, but better *response*

xv6 scheduling

CPU 0

CPU 1

CPU 2




xv6 scheduling

CPU 0

CPU 1

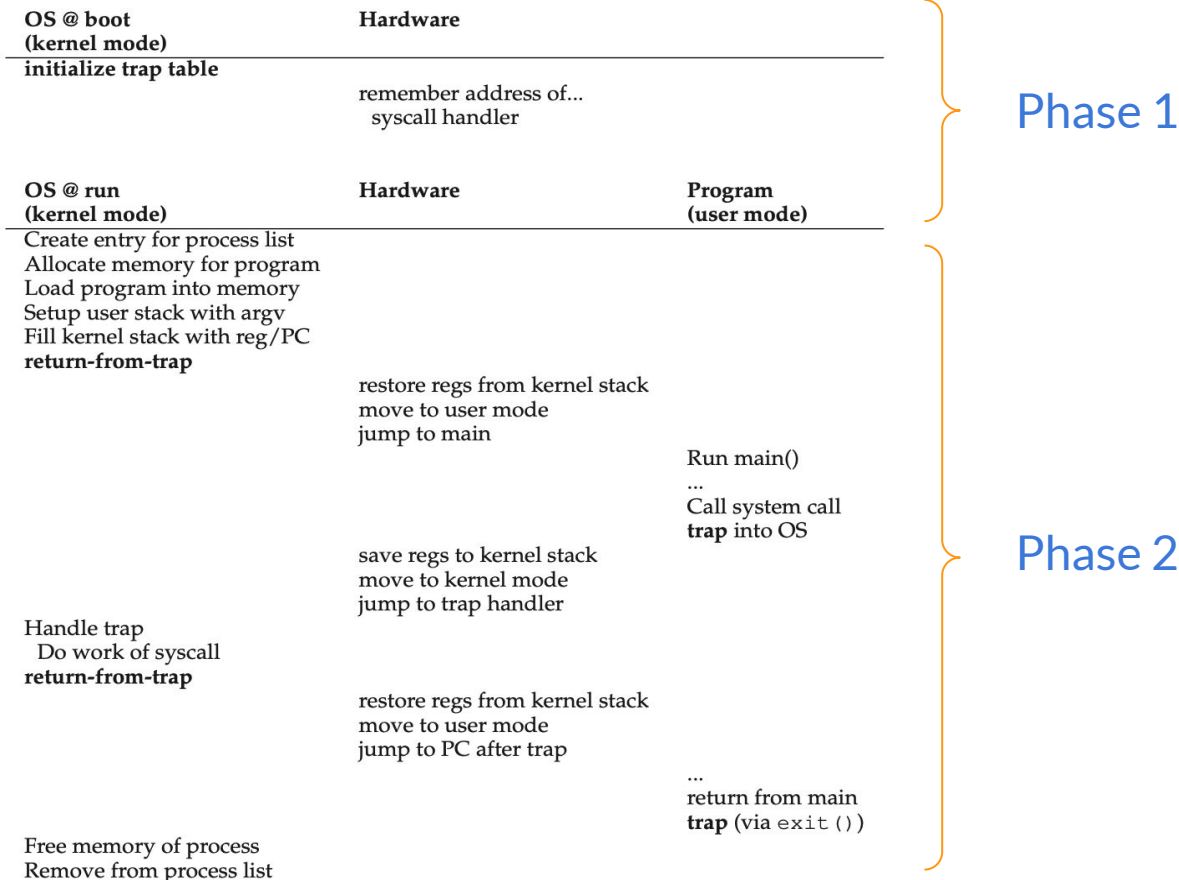
CPU 2



Round Robin Scheduling of xv6

- `scheduler()` in `proc.c`
 - Invoked from `kernel/main.c`
 - **INFINITE LOOP** -> it never returns
 - Control loop that executes all user applications
 - Selects the next `RUNNABLE` state
 - Switches to that process (`swtch.S`)
 - Where it last yielded/slept, or its start
 - NEED TO SAVE THE CONTEXT OF SCHEDULER ITSELF!!!
 - Releases the lock at the end (pay attention to lock acquire releases)

Limited Direct Execution Protocol

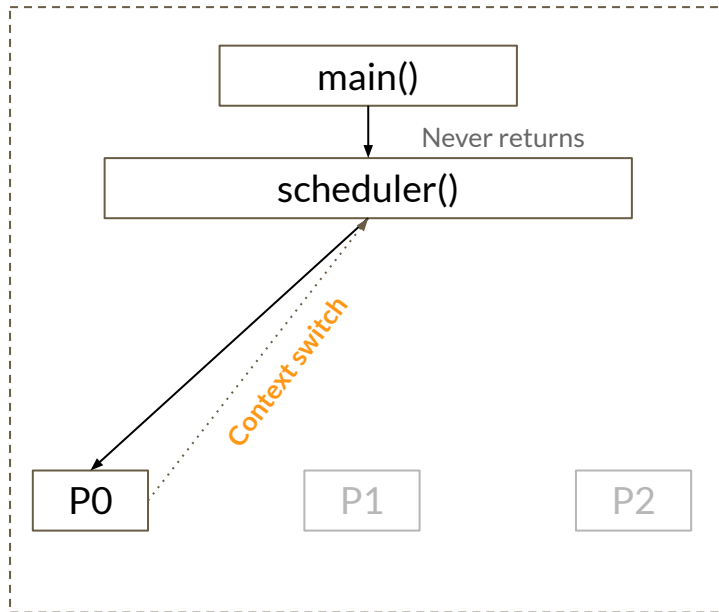


xv6 scheduling

CPU 0

CPU 1

CPU 2



Round Robin Scheduling of xv6

- A user process to give up CPU
 - `exit()`
 - `yield()`
 - `sleep()`
 - Preempted by the timer interrupt
 - 100ms

`sched()` in `proc.c`

Round Robin Scheduling of xv6

- `sched()` in `proc.c`
 - `myproc()`
 - Currency process
 - Hold only `p->lock`
 - Process must not be RUNNING
 - RUNNABLE, SLEEPING, or ZOMBIE
 - Interrupt bookkeeping
 - `mycpu()->context`
 - Holds the context of the `scheduler` process

Implementing Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (**largest**)
- Xv6 Schemes:
 - **Preemptive** – The timer interrupt yields() the running process and should provide a chance to schedule a new process with higher priority
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- **Problem: Starvation**– Low priority processes may never execute
- **Solution: Aging** – As time progresses, increase the priority of the process

Scheduling: Priority Example

Process	Arrival Time	Burst Time	Priority
P ₁	0	7	2
P ₂	2	4	1
P ₃	4	1	4
P ₄	5	4	3

- Non-preemptive:
 - $P_1 \rightarrow P_2 \rightarrow P_4 \rightarrow P_3$
- Preemptive:



Task 1

- Implement `getpriority()` and `setpriority()` system calls
 - Where do you add the priority field?
 - When do you initialize?
- Forked child should inherit the parent priority
 - `fork()`
 - Where do you set the priority when new child is created with `fork()`?
- Modify `ps` program to show priority of each process
- Test program:
 - `task1.c`

Task 2

- Implement priority scheduling
 - Compile time option in `param.h` to select between round robin and priority scheduling
 - `#define Scheduler_ALGO 0 or 1`
 - Which function should host the implementation of the priority scheduling?
- Test program
 - `task2.c`

Task 3

- Keep track of process aging
 - Add a field in `proc` to store `readytime`
 - When a process is changed for other state to `RUNNABLE`
- Modify `ps` program to print age of a process
 - Update given `pstat.h`
- `age`
 - `current time - ready time`
- Test your implementation

Task 4

- Implement and merge process aging with priority scheduling
 - Add a field in `proc` to store `readytime`
 - When a process is changed for other state to `RUNNABLE`
- Modify `ps` program to print age of a process
 - Update given `pstat.h`
- **age**
 - $\text{current time} - \text{ready time}$
- Test your implementation

Task 4


- age
 - current time - ready time

```
#define MAXEFPRIORITY 99
#define AGING_DIV 25

effective_priority = min(MAXEFPRIORITY, priority + (currtime -
    readytime)/AGING_DIV)
```

Scheduling Metrics (Extra credit)

- T_a - **Arrival time**: Time the process became “READY” [again]
- T_w - **Waiting time**: Time spent waiting for the CPU
- T_s - **Service time**: Time spent executing in the CPU
- T_r - **Turnaround time**: Time spent waiting and executing = $T_w + T_s$



$$T_r / T_s = 2.5$$

Announcement

- Quiz 3
 - On blackboard
- Homework 3
 - Due on Monday October 27th, 11.59 PM