

# Operating Systems Concepts

xv6 scheduling



CS 4375, Fall 2025

Instructor: MD Armanuzzaman (*Arman*)

[marmanuzzaman@utep.edu](mailto:marmanuzzaman@utep.edu)

October 8, 2025

# Summary

- Main Memory:
  - Fixed and Dynamic Memory Allocation
  - External and Internal Fragmentation
  - Address Binding
  - Hardware Address Protection
  - Paging
  - Segmentation

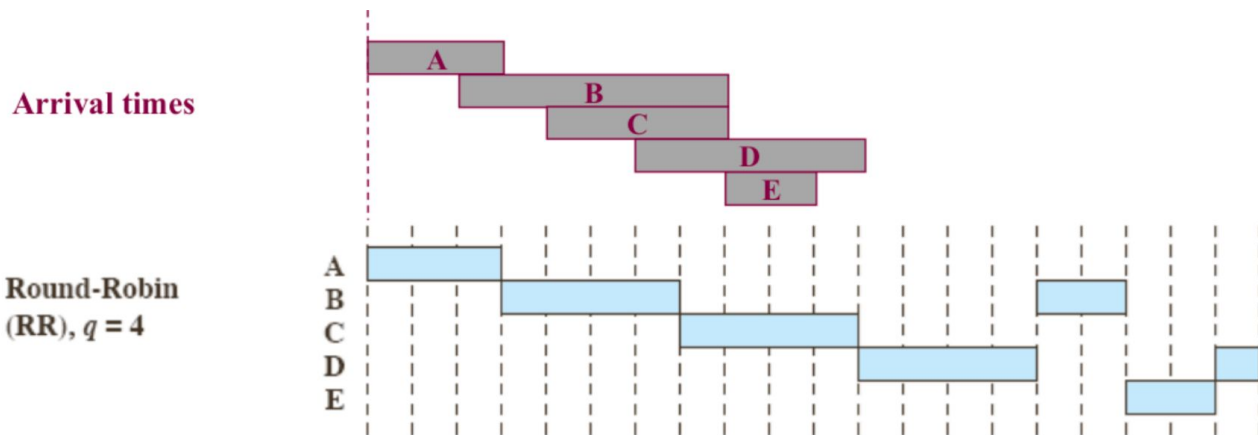
# Agenda

- xv6 scheduling
  - Default round robin scheduling of xv6
  - Understand the code base
- Implement priority scheduling in xv6
  - Tasks in homework 3
- Implement aging policy

# Scheduling: Round-Robin (RR)

- A crucial parameter is the quantum  $q$  (~10-100ms)
  - $q$  should be large compared to context switch latency (~10 $\mu$ s)
  - $q$  should be less than the longest CPU burst, or RR **degenerates** to FCFS

RR ( $q = 4$ )  
Scheduling  
Policy



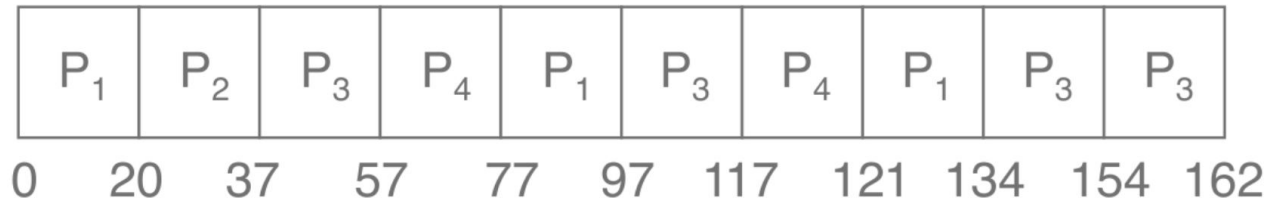
|            |                           |            |             |             |             |             |             |
|------------|---------------------------|------------|-------------|-------------|-------------|-------------|-------------|
| RR $q = 4$ | Finish Time               | <b>A</b> 3 | <b>B</b> 17 | <b>C</b> 11 | <b>D</b> 20 | <b>E</b> 19 | <b>Mean</b> |
|            | Turnaround Time ( $T_T$ ) | 3          | 15          | 7           | 14          | 11          | 10.00       |
|            | $T_T/T_S$                 | 1.00       | 2.5         | 1.75        | 2.80        | 5.50        | 2.71        |

# Scheduling: RR Example

- Consider  $q = 20$

| Process | Burst Time |
|---------|------------|
| $P_1$   | 53         |
| $P_2$   | 17         |
| $P_3$   | 68         |
| $P_4$   | 24         |

- The Gantt chart for the schedule is:



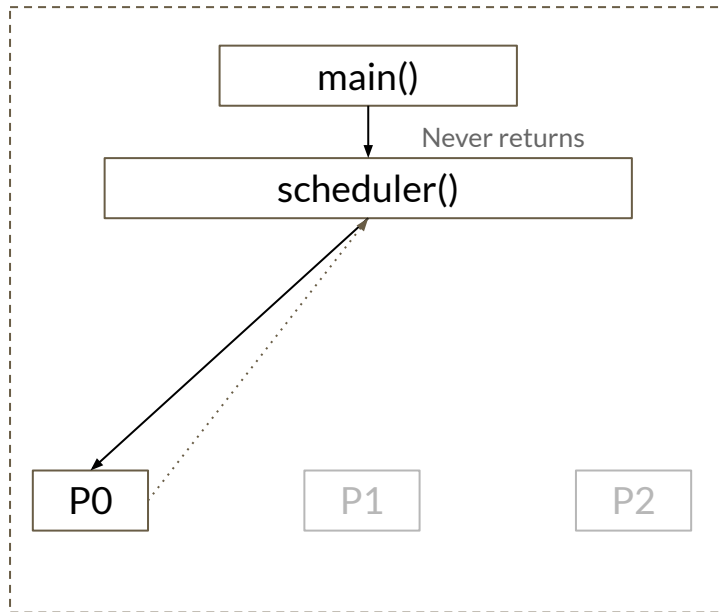
- Typically, higher average turnaround than SJF, but better *response*

# xv6 scheduling

CPU 0

CPU 1

CPU 2




# xv6 scheduling

CPU 0

CPU 1

CPU 2

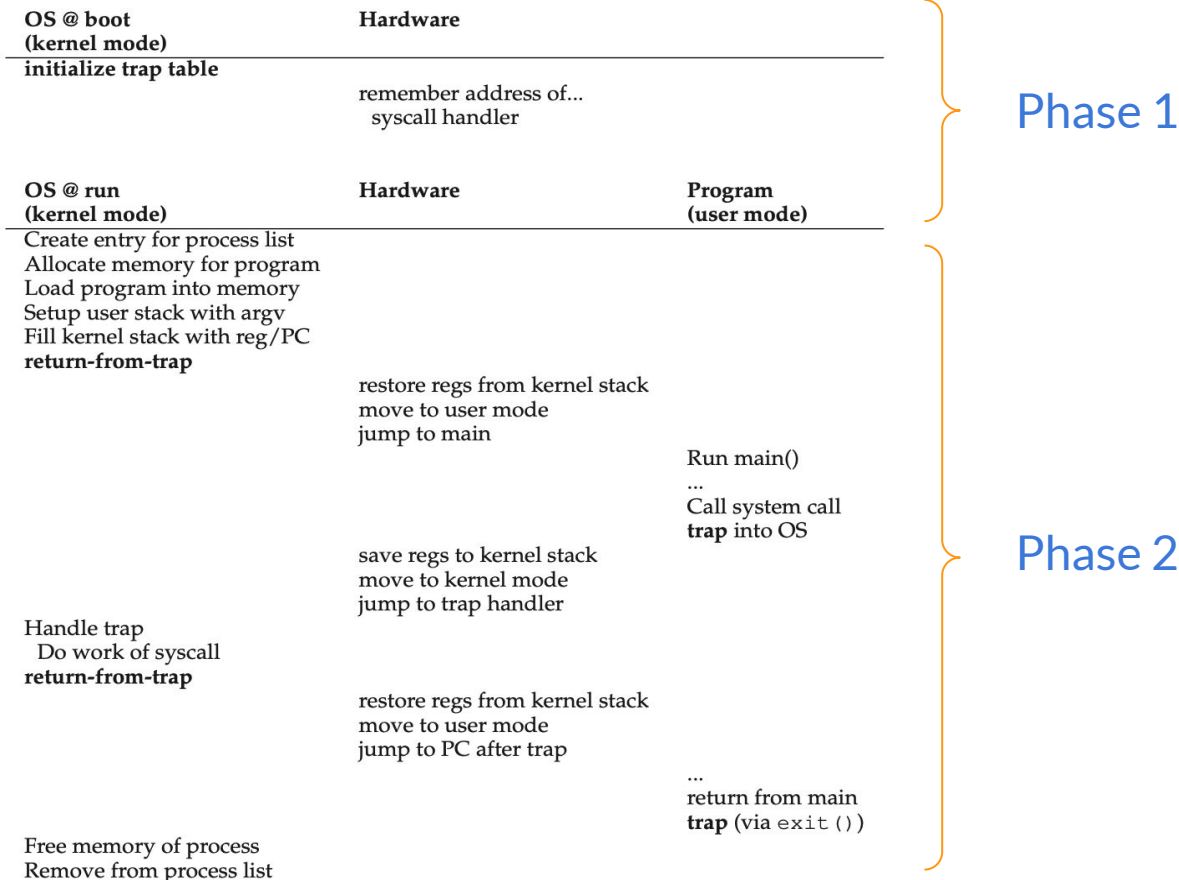


# Round Robin Scheduling of xv6

- `scheduler()` in `proc.c`
  - Invoked from `kernel/main.c`
  - **INFINITE LOOP** -> it never returns
    - Control loop that executes all user applications
  - Selects the next `RUNNABLE` state
  - Switches to that process (`swtch.S`)
    - Where it last yielded/slept, or its start
    - NEED TO SAVE THE CONTEXT OF SCHEDULER ITSELF!!!
  - Releases the lock at the end (pay attention to lock acquire releases)



# Limited Direct Execution Protocol

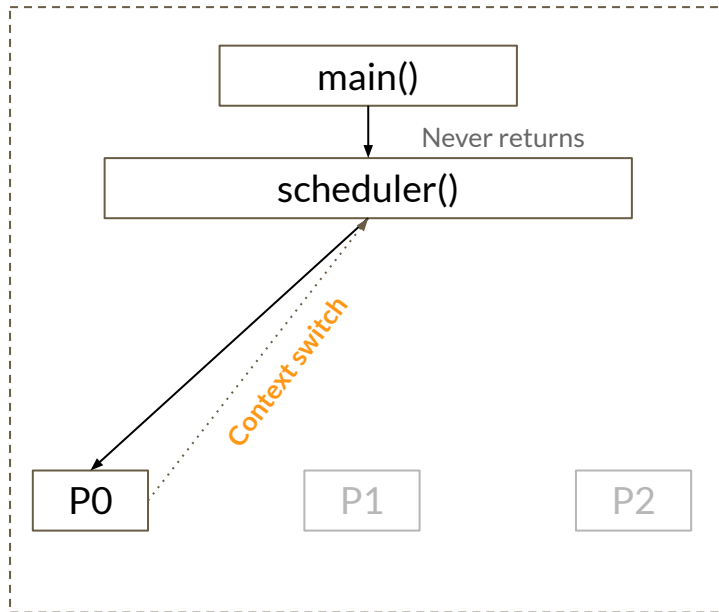


# xv6 scheduling

CPU 0

CPU 1

CPU 2



# Round Robin Scheduling of xv6

- A user process to give up CPU
  - `exit()`
  - `yield()`
  - `sleep()`
  - Preempted by the timer interrupt
    - 100ms

`sched()` in `proc.c`

# Round Robin Scheduling of xv6

- `sched()` in `proc.c`
  - `myproc()`
    - Currency process
  - Hold only `p->lock`
  - Process must not be RUNNING
    - RUNNABLE, SLEEPING, or ZOMBIE
  - Interrupt bookkeeping
  - `mycpu()->context`
    - Holds the context of the `scheduler` process


# Implementing Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (**largest**)
- Xv6 Schemes:
  - **Preemptive** – The timer interrupt yields() the running process and should provide a chance to schedule a new process with higher priority
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- **Problem: Starvation**– Low priority processes may never execute
- **Solution: Aging** – As time progresses, increase the priority of the process

# Scheduling: Priority Example

| Process        | Arrival Time | Burst Time | Priority |
|----------------|--------------|------------|----------|
| P <sub>1</sub> | 0            | 7          | 2        |
| P <sub>2</sub> | 2            | 4          | 1        |
| P <sub>3</sub> | 4            | 1          | 4        |
| P <sub>4</sub> | 5            | 4          | 3        |

- Non-preemptive:
  - $P_1 \rightarrow P_2 \rightarrow P_4 \rightarrow P_3$
- Preemptive:



# Task 1

- Implement `getpriority()` and `setpriority()` system calls
  - Where do you add the priority field?
  - When do you initialize?
- Forked child should inherit the parent priority
  - `fork()`
  - Where do you set the priority when new child is created with `fork()`?
- Modify `ps` program to show priority of each process
- Test program:
  - `task1.c`

# Task 2

- Implement priority scheduling
  - Compile time option in `param.h` to select between round robin and priority scheduling
  - `#define Scheduler_ALGO 0 or 1`
  - Which function should host the implementation of the priority scheduling?
- Test program
  - `task2.c`



# Task 3

- Keep track of process aging
  - Add a field in `proc` to store `readytime`
  - When a process is changed for other state to `RUNNABLE`
- Modify `ps` program to print age of a process
  - Update given `pstat.h`
- `age`
  - `current time - ready time`
- Test your implementation

# Task 4

- Implement and merge process aging with priority scheduling
  - Add a field in `proc` to store `readytime`
  - When a process is changed for other state to `RUNNABLE`
- Modify `ps` program to print age of a process
  - Update given `pstat.h`
- **age**
  - `current time - ready time`
- Test your implementation

# Task 4


- age
  - current time - ready time

```
#define MAXEFPRIORITY 99
#define AGING_DIV 25

effective_priority = min(MAXEFPRIORITY, priority + (currtime -
    readytime)/AGING_DIV)
```

# Scheduling Metrics (Extra credit)

- $T_a$  - **Arrival time**: Time the process became “READY” [again]
- $T_w$  - **Waiting time**: Time spent waiting for the CPU
- $T_s$  - **Service time**: Time spent executing in the CPU
- $T_r$  - **Turnaround time**: Time spent waiting and executing =  $T_w + T_s$



$$T_r / T_s = 2.5$$

# Announcement

- Quiz 3
  - On blackboard: 15 MCQs, 30 minutes single attempt.
  - Due tomorrow **11.59 PM**
- Midterm
  - **October 15, next wednesday class time**
  - MCQ and written [Based on lectures before 15th]
  - Review class on monday
- Homework 3
  - Due on Monday October 27th, 11.59 PM