

Operating Systems Concepts

Midterm Review



CS 4375, Fall 2025

Instructor: MD Armanuzzaman (*Arman*)

marmanuzzaman@utep.edu

October 13, 2025

Summary

- xv6 scheduling
 - Default round robin scheduling of xv6
 - Understand the code base
- Implement priority scheduling in xv6
 - Tasks in homework 3
- Implement aging policy

Agenda

- Midterm:
 - Logistics
 - Topics
 - Review

Midterm Exam

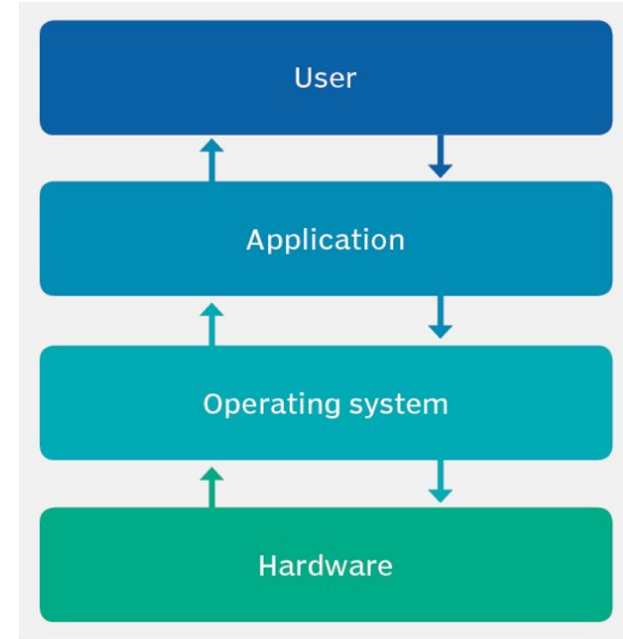
- Wednesday, October 15th
- During class time
- Handwritten one page note allowed
 - No other notes allowed
- Question pattern
 - MCQ
 - Written
- Total points **150**
- 1 hour

Topics

- Introduction to OS
- OS structure
- Processes
- Threads
- System Calls and Context Switches
- CPU Scheduling
- Address Translation

Introduction to OS

- An operating system (OS) is the systems software that **interfaces** application programs with the underlying hardware
- Applications make requests for services through a **defined APIs**
- The OS **manages** and **interfaces** to underlying hardware (e.g., **processors, memory, storage devices, network interfaces**) so that applications don't need to know about hardware details
- The OS **launches** and **manages** every application, including multiple processes or threads.



Types of Operating Systems

- General-purpose operating system
 - Run multiple applications in broad range of hardware
 - Windows, MacOS, Linux
- Mobile operating system
 - Efficient performance and resource usage and fast response time
 - Apple iOS, Google Android
- Embedded operating system
 - Usually provided on a chip that is incorporated into the device
 - ATMs, IoT devices, medical devices
 - Embedded Linux
- Real-time operating system
 - Respond quickly and predictably under time constraints
 - FreeRTOS, zephyr




Execution Steps Simplified

- Fetch instruction at PC
- Decode the instruction
- Execute (possibly using registers and modifying them)
- Write possible result to registers/memory
- Increase PC to point to the next instruction

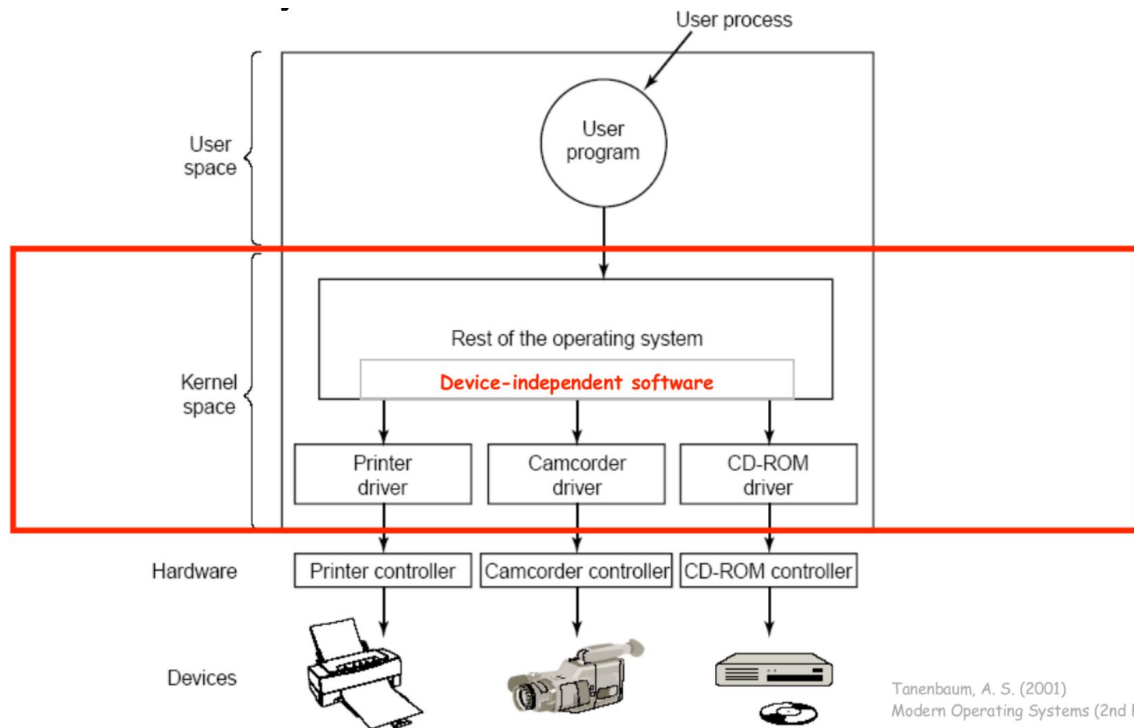
REPEAT!

Modern Computer Architecture




I/O Management

- Layers of the I/O subsystem




I/O Management

- I/O Device Queues


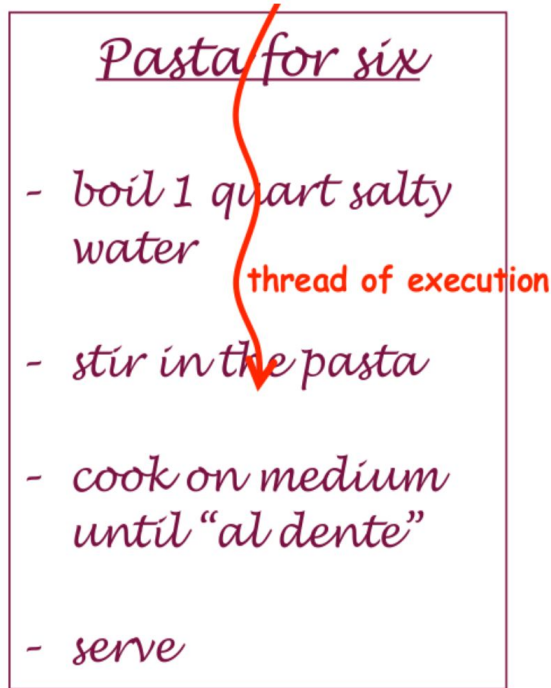


I/O Management: Two Methods



Processes Concept

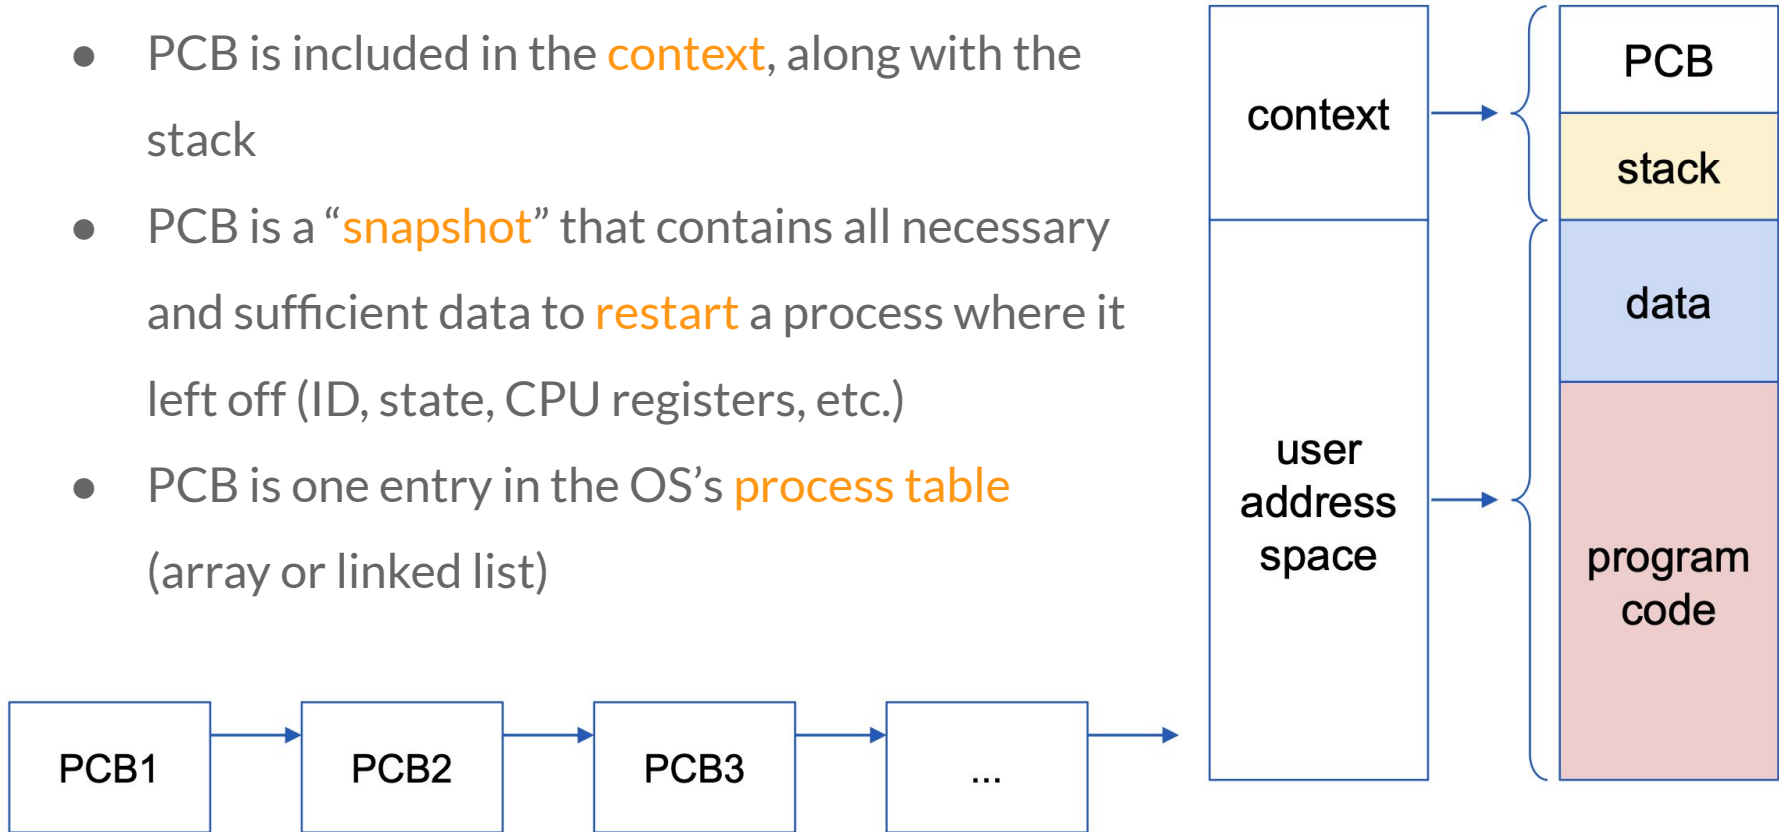
- A **process** is a program in execution.



Process


Process Control Block (PCB)

- PCB is included in the **context**, along with the stack
- PCB is a “**snapshot**” that contains all necessary and sufficient data to **restart** a process where it left off (ID, state, CPU registers, etc.)
- PCB is one entry in the OS’s **process table** (array or linked list)



Process State

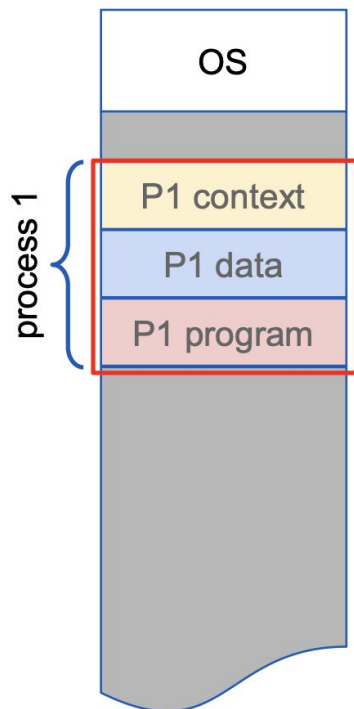
- A process changes its **state** during execution:
 - **New**: The process is being created
 - **Ready**: The process is waiting to be assigned to a processor
 - **Running**: Instructions are being executed
 - **Waiting**: The process is waiting for some event to occur
 - **Terminated**: The process has finished execution



Process Creation

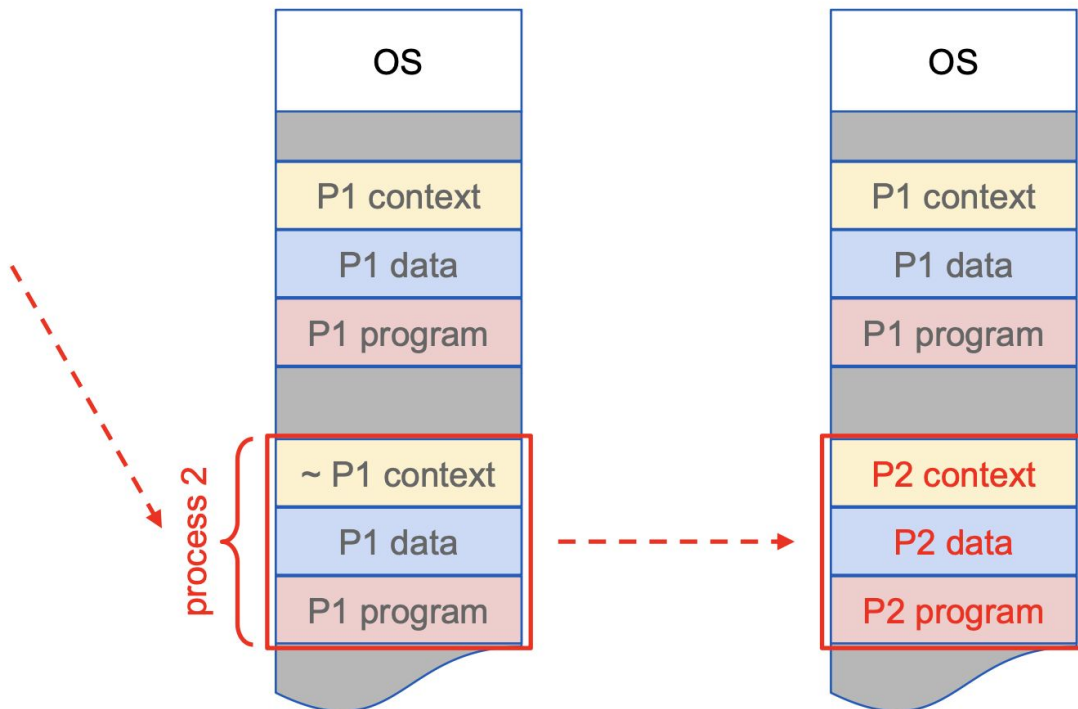
1. Clone child process

`pid = fork()`



2. Replace child's image

`execvp(name, ...)`



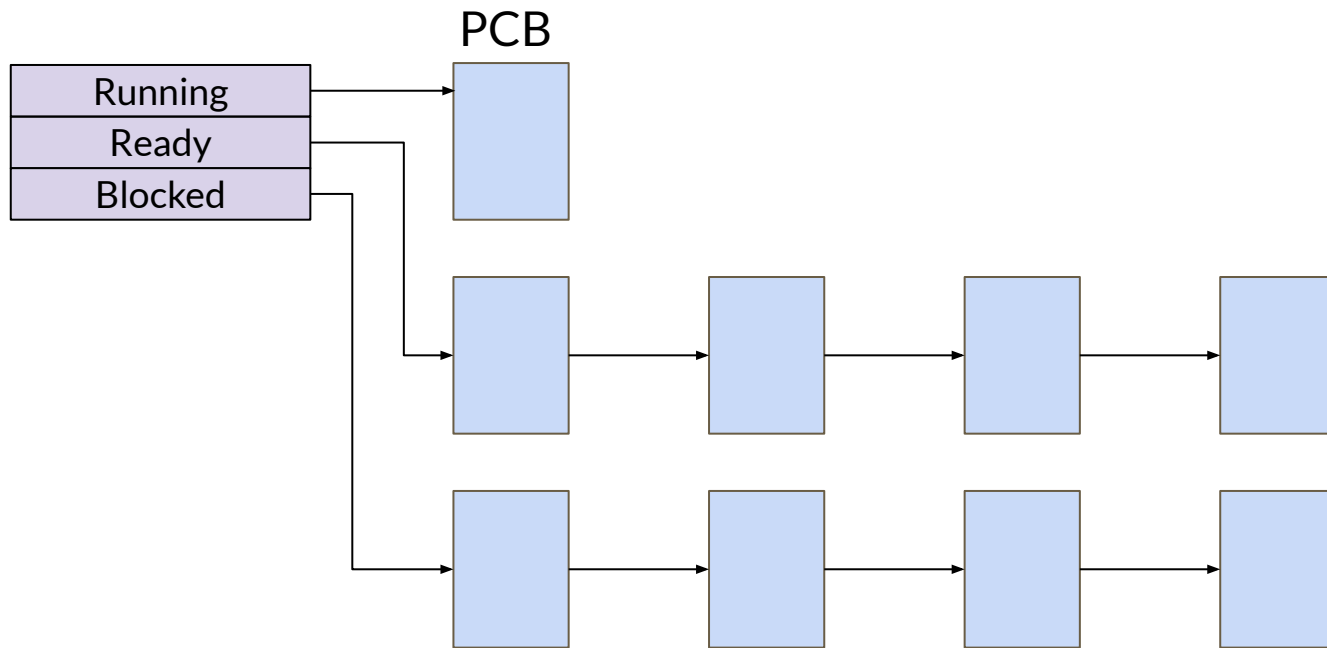
Fork Example

```
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main()
5  {
6      fork();
7      fork();
8      fork();
9      printf("Process pid is %d\n", getpid());
10     return 0;
11 }
```

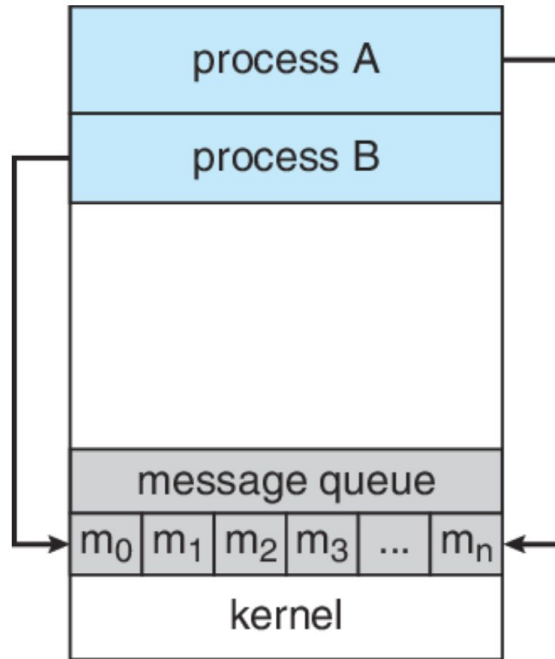
- How many lines of output will this code produce?

Process Queues

The process table can be **split** into per-state queues: PCBs can be linked together if they contain a pointer field




Communications Models



(a)

Message Passing

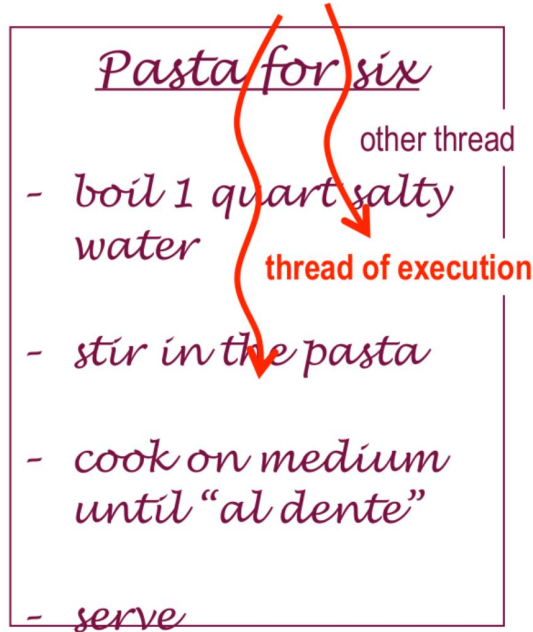


(b)


Shared Memory

Multithreading

The execution part is a “thread” that can be multiplied



Program

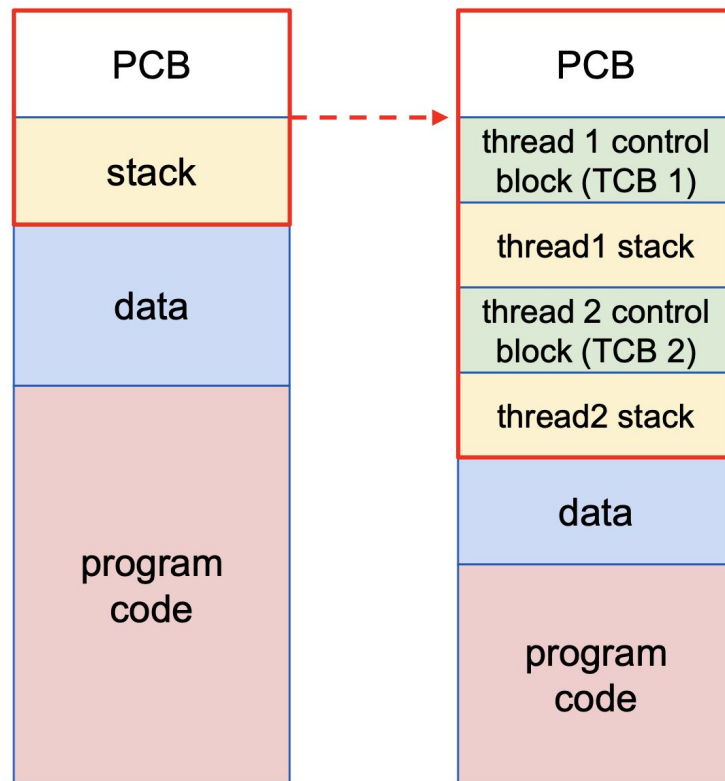


Process

New Process Description Model

Multithreading requires changes in the process description model.


- Each thread of execution receives its own control block and stack.
 - Own execution state (“Running”, “Blocked”, etc.)
 - Own copy of CPU registers
 - Own execution history (stack)
- The process keeps a global control block listing resources currently used



Multiprocessing Model

Thread spawning

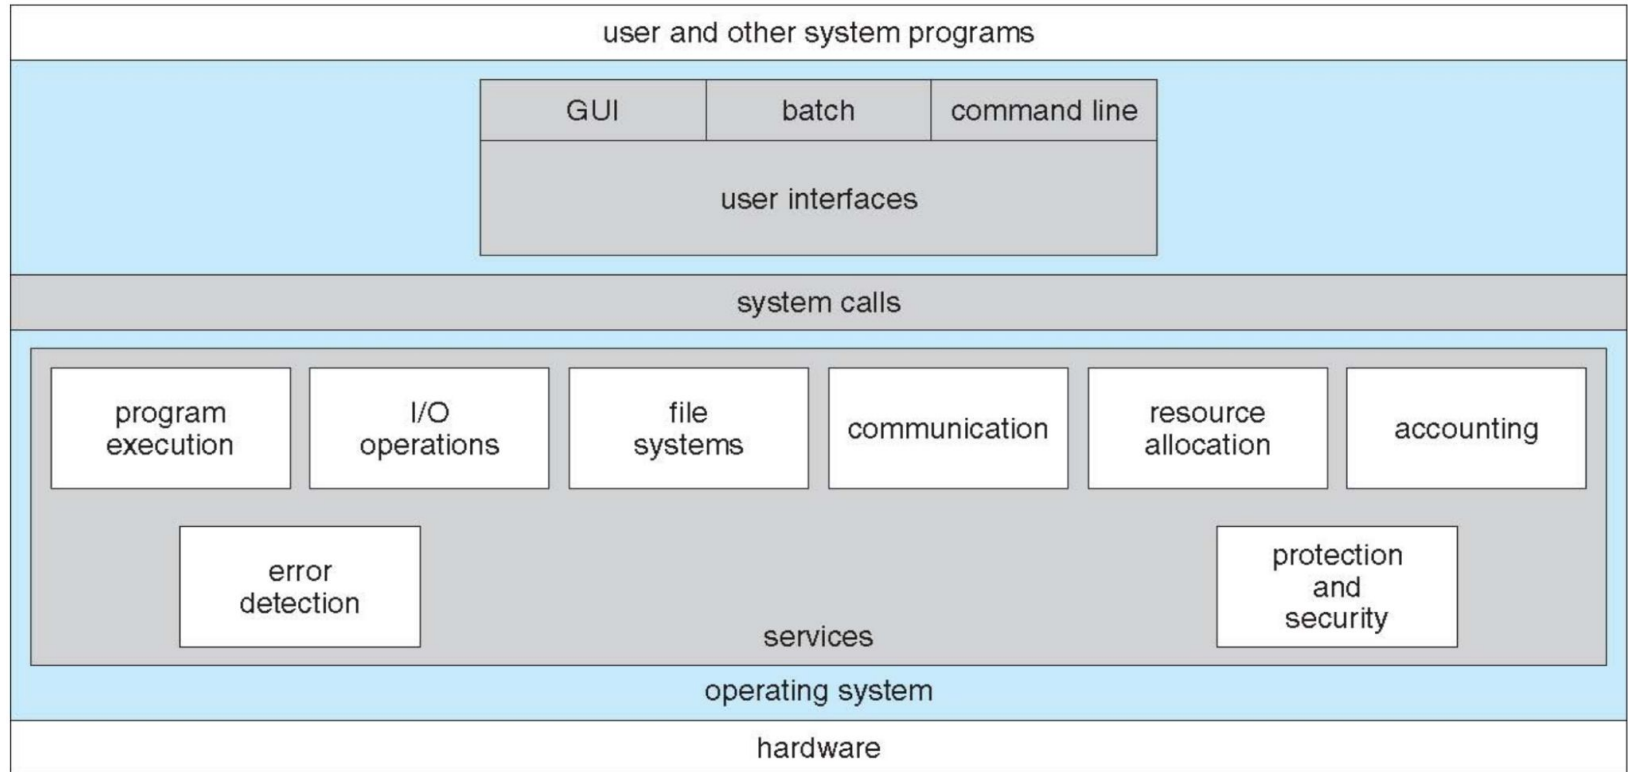
- Threads are created *within* processes, and belonging to processes
- All the threads created within one process, share the resources of the process including the address space
- Scheduling is performed on per-thread basis
- Threads have a similar lifecycle as the processes and will be managed mainly in the same way



Threads vs Processes

- A common terminology:
 - Heavyweight Process = Process
 - **Lightweight Process = Thread**
- **Advantages:**
 - Much quicker to create a thread than a process
 - spawning a new thread only involves allocating a new stack and a new CPU state block
 - Much quicker to switch between threads than to switch between processes
 - Threads share data easily
- **Disadvantages:**
 - Processes are more flexible
 - They don't have to run on the same processor
 - No security between threads: One thread can stomp on another thread's data
 - For threads which are supported by user thread package instead of the kernel:
 - If one thread blocks all threads in task block

Operating System Services



Operating System Services

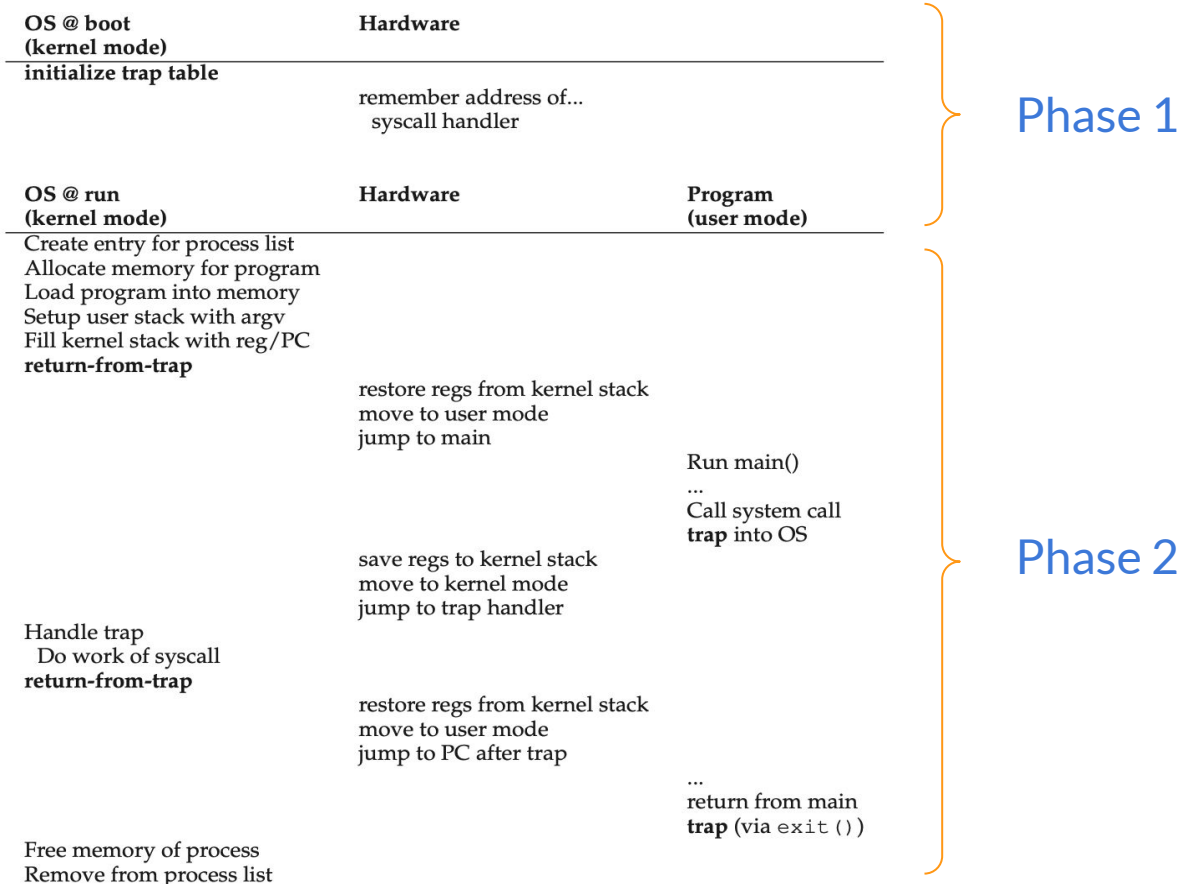
- Virtualize the CPU
 - Performance
 - Avoid overhead
 - Maximize CPU usage
 - Control
 - Provide performance while retaining control
 - CPU time, I/O resources, etc.

Solution

- OS
 - System calls
 - Allows the user program to request privileged operation
 - Trap table
 - Kernel boots first (privileged)
 - Sets up the **trap table**
 - Informs the hardware about **trap handlers** (Privileged instructions)
 - Hardware remembers until next reboot

Allowing user processes to jump to anywhere in the kernel is dangerous!!


Limited Direct Execution Protocol




System Calls

- On modern operating systems, processes do not talk to hardware directly, but must go through OS
- **System Call:** request from a process for OS to do some kind of work on its behalf
- **Application Programming Interface (API):** interface provided by OS, usually in a high-level language (C or C++), that is easier to work with than raw system calls
 - POSIX for macOS, Linux, and other Unix-like, accessible through libc.so

System Call Interface



System calls




Scheduling Criteria

- **CPU Utilization:** Keep the CPU as busy as possible **Maximize**
- **Throughput:** # of processes that complete their execution per time unit **Maximize**
- **Response Time:** Amount of time it takes from when a request was submitted until the first response (not output) is produced (for time-sharing environment) **Minimize**
- **Waiting Time:** Total amount of time a process has been waiting in the ready queue **Minimize**
- **Turnaround Time:** Amount of time passed to finish execution of a particular process i.e. execution(service) time + waiting time **Minimize**

Scheduling Metrics

- T_a - **Arrival time**: Time the process became “READY” [again]
- T_w - **Waiting time**: Time spent waiting for the CPU
- T_s - **Service time**: Time spent executing in the CPU
- T_r - **Turnaround time**: Time spent waiting and executing = $T_w + T_s$




$$T_r / T_s = 2.5$$

Scheduling: First-Come, First-Served (FCFS)

- Processes are assigned the CPU in the order they request it
- When the running process blocks, the first “READY” process is selected
- When a process gets “READY”, it is put at the end of the queue

FCFS
Scheduling
Policy



FCFS	Finish Time	A	3	B	9	C	13	D	18	E	20	Mean
	Turnaround Time (T_T)		3		7		9		12		12	8.60
	T_T/T_S		1.00		1.17		2.25		2.40		6.00	2.56


Scheduling: Non-Preemptive SJF

- Among several equally important “**READY**” jobs (or CPU bursts), the scheduler picks the one that will finish the earliest

SJF
Scheduling
Policy

Arrival times

Shortest Job
First (SJF)




SJF	Finish Time	A	3	B	9	C	15	D	20	E	11	Mean
	Turnaround Time (T_T)		3		7		11		14		3	7.60
	T_T/T_S		1.00		1.17		2.75		2.80		1.50	1.84

Scheduling: Round-Robin (RR)

- A crucial parameter is the quantum q ($\sim 10\text{-}100\text{ms}$)
 - q should be large compared to context switch latency ($\sim 10\mu\text{s}$)
 - q should be less than the longest CPU burst, or RR **degenerates** to FCFS

RR ($q = 4$)
Scheduling
Policy



RR $q = 4$	Finish Time	A 3	B 17	C 11	D 20	E 19	Mean
	Turnaround Time (T_T)	3	15	7	14	11	10.00
	T_T/T_S	1.00	2.5	1.75	2.80	5.50	2.71

Comparison: Round-Robin

- PROS:

- Great for timesharing
 - No starvation
- Does not require prior knowledge of CPU burst times
- Generally reduces average response time

- CONS:

- What if all jobs are almost the same length
 - Increases the turnaround time
- How to set the “best” time quantum?
 - If too small, it increases context-switch overhead
 - If too large, response time degrades

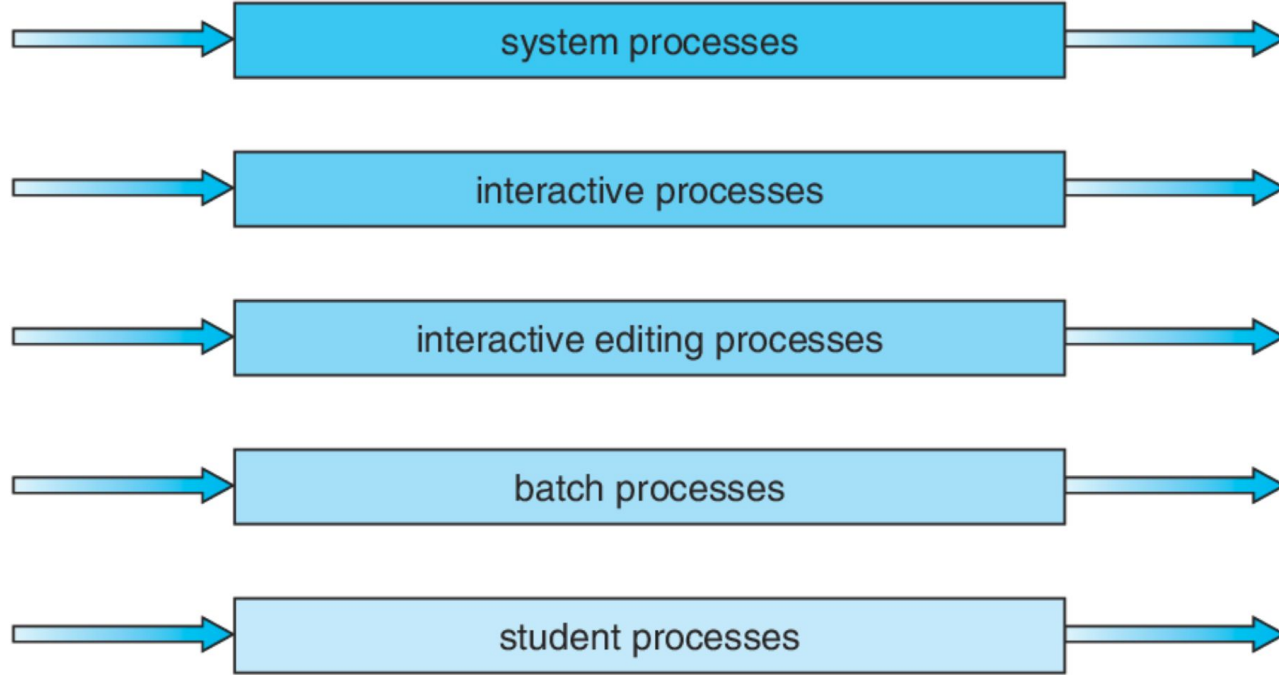
Exercise

Process	Arrival Time	Burst Time	Priority
P ₁	9	3	5
P ₂	6	4	4
P ₃	8	5	3
P ₄	7	9	2
P ₅	5	7	1

- Consider Shortest-Remaining-Time-First (SRTF) scheduling policy is used.
 - a. Draw the Gantt chart
 - b. Find the response time, waiting time, and turnaround time for each process.

Multilevel Queue Scheduler

highest priority



lowest priority


Example of Multilevel Feedback Queue Scheduler

- Three queues:

- Q_0 - RR with $q = 8$ ms
- Q_1 - RR with $q = 16$ ms
- Q_2 - FCFS


- Scheduling

- A new job enters queue Q_0 which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, **job is moved to queue Q_1**
- At Q_1 job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is **preempted and moved to queue Q_2**



Proportional Share Scheduling

- Try to guarantee that each process gets a certain percentage of the CPU time
 - Over some period of time, sometimes called the scheduling interval
 - “Fair share” may depend on process’s priority




Lottery Scheduling

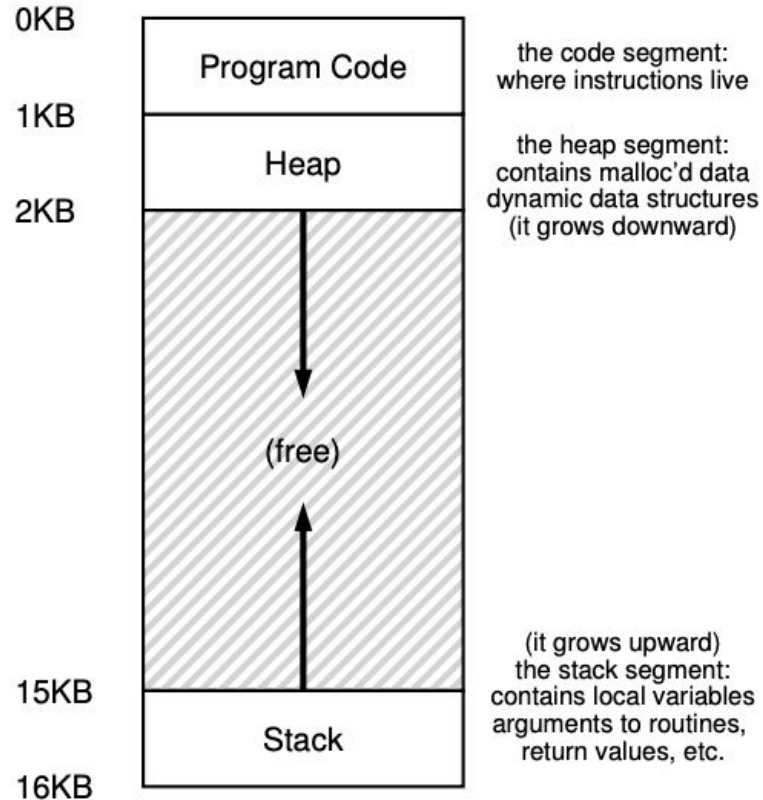
- Tickets assigned to a process represent its share
 - e.g., 10 out of 100 tickets process should get 10% of the CPU time
- Scheduling is probabilistic, using a random number generator
 - What problem does this raise?
- What about processes that do I/O?
- Ticket currency allows local authority (e.g., user) flexibility in allocating tickets
 - e.g., can do ticket inflation within local group of processes that trust each other
- Efficient implementation

Lottery Scheduling: Problems

- How To Assign Tickets?
 - Open problem
- Unfairness:
 - Two jobs with short lengths
 - Unfairness metric = $A_{\text{runtime}}/B_{\text{runtime}}$
- Not Deterministic

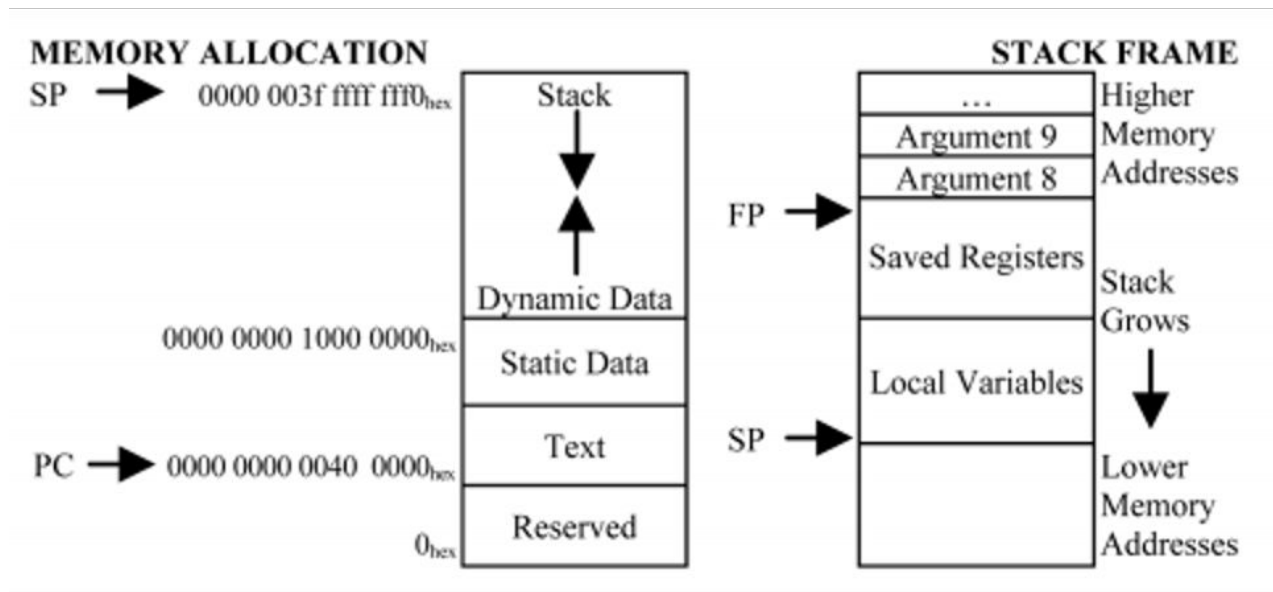


Program Address Space



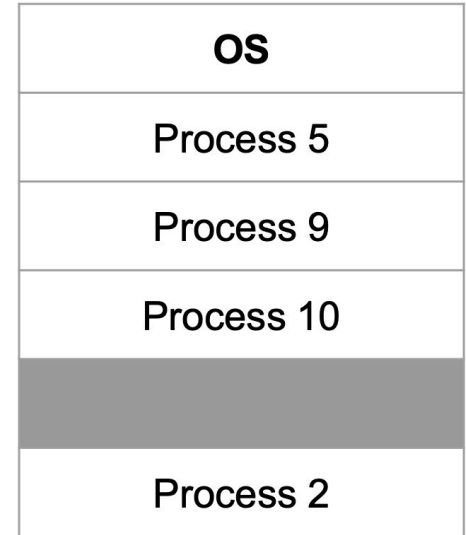
Program Address Space (RISC-V)

- RISC-V Linux 39-bit Address Space




Memory Allocation

- Fixed-partition allocation
 - Divide memory into fixed-size partitions
 - Each partition contains exactly one process
 - The degree of multiprogramming is bound by the number of partitions
 - When a process terminates, the partition becomes available for other processes
- May lead to **Internal Fragmentation**– allocated memory may be slightly larger than **requested** memory; this size difference is memory **internal** to a partition, but not being **used**



Paging Example




logical
memory

0	1
1	4
2	3
3	7

page table

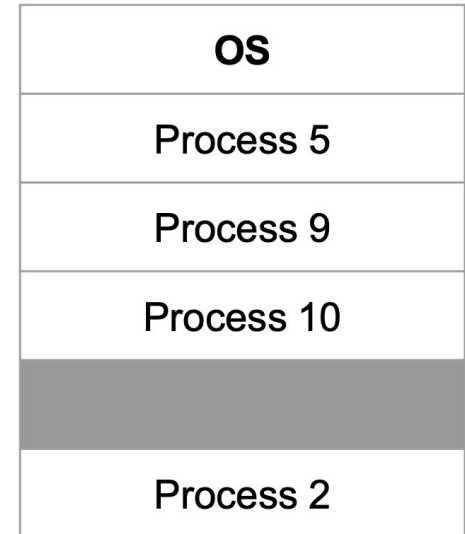
frame
number




physical
memory

Memory Allocation

- Variable-partition Scheme (Dynamic)
 - When a process arrives, search for a hole large enough for this process
 - **Hole** – block of available memory; holes of various size are scattered throughout memory
 - Allocate only as much memory as needed
 - Operating system maintains information about:
 - allocated partitions
 - free partitions (hole)



Segmentation: Address Translation Architecture



Paging vs Segmentation

Paging	Segmentation
Fixed-size division	Variable-size division
Managed by OS	Managed by compiler
Fragmentation: Internal	Fragmentation: External
Speed: Faster	Speed: Slower
Unit size by hardware	Unit size by user/programmer
Invisible to the user	Visible to the user

Modern OSes implement a hybrid approach called “**segmentation with paging**”

Midterm Exam

- Wednesday, October 15th
- During class time
- Handwritten one page note allowed
 - No other notes allowed
- Question pattern
 - MCQ
 - Written
- Total points **150**
- 1 hour