

Operating Systems Concepts

Main Memory: Paging TLB



CS 4375, Fall 2025

Instructor: MD Armanuzzaman (*Arman*)

marmanuzzaman@utep.edu

October 20, 2025

Summary

- Main Memory:
 - Fixed and Dynamic Memory Allocation
 - External and Internal Fragmentation
 - Address Binding
 - Hardware Address Protection
 - Paging
 - Segmentation

Agenda

- Main Memory:
 - Implementation of Paging
 - Effective Access Time
 - Multi-level Paging
 - Example: Intel 32-bit Architecture

Dynamic Storage-Allocation Problem

- How to satisfy a request of size n from a list of free holes?
 - **First-fit:** Allocate the first hole that is big enough
 - **Best-fit:** Allocate the smallest hole that is big enough; must search entire list, unless ordered by size. Produces the smallest leftover hole.
 - **Worst-fit:** Allocate the largest hole; must also search entire list. Produces the largest leftover hole.
- First-fit is faster.
- Best-fit is better in terms of storage utilization.
- Worst-fit may lead to less fragmentation.

Dynamic Storage-Allocation Example

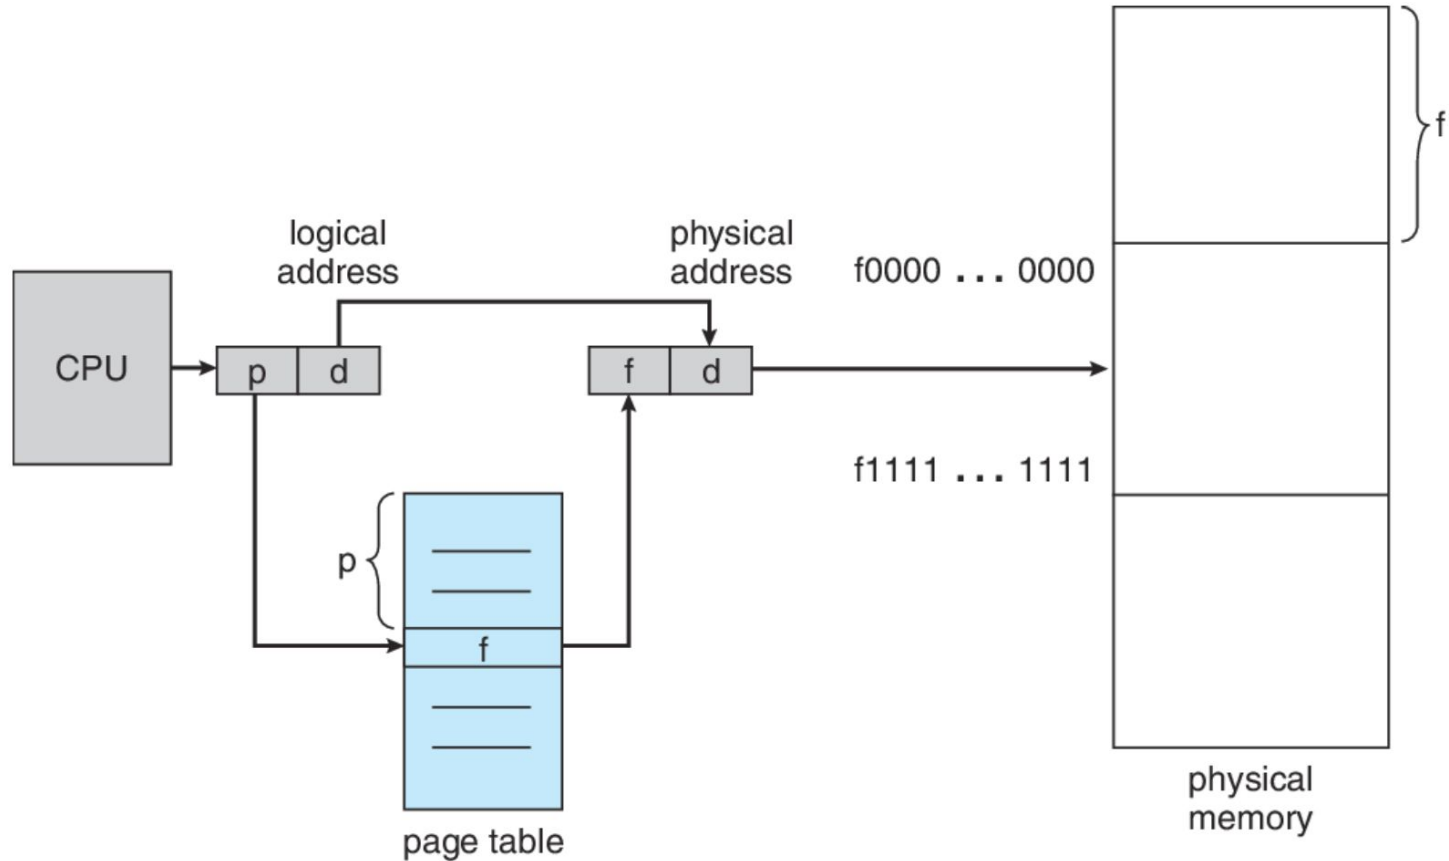
- Given five memory partitions of 100 KB, 500 KB, 200 KB, 300 KB, and 600 KB (in order), how would each of the first-fit, best-fit, and worst-fit algorithms place processes of 212 KB, 417 KB, 112 KB, and 426 KB (in order)?
- Which algorithm makes the most-efficient use of memory?

Paging Address Translation Scheme

- Address generated by CPU is divided into:
 - **Page number (p)** – used as an index into a page table which contains base address of each page in physical memory
 - **Page offset (d)** – combined with base address to define the physical memory address that is sent to the memory unit



Address Translation Architecture



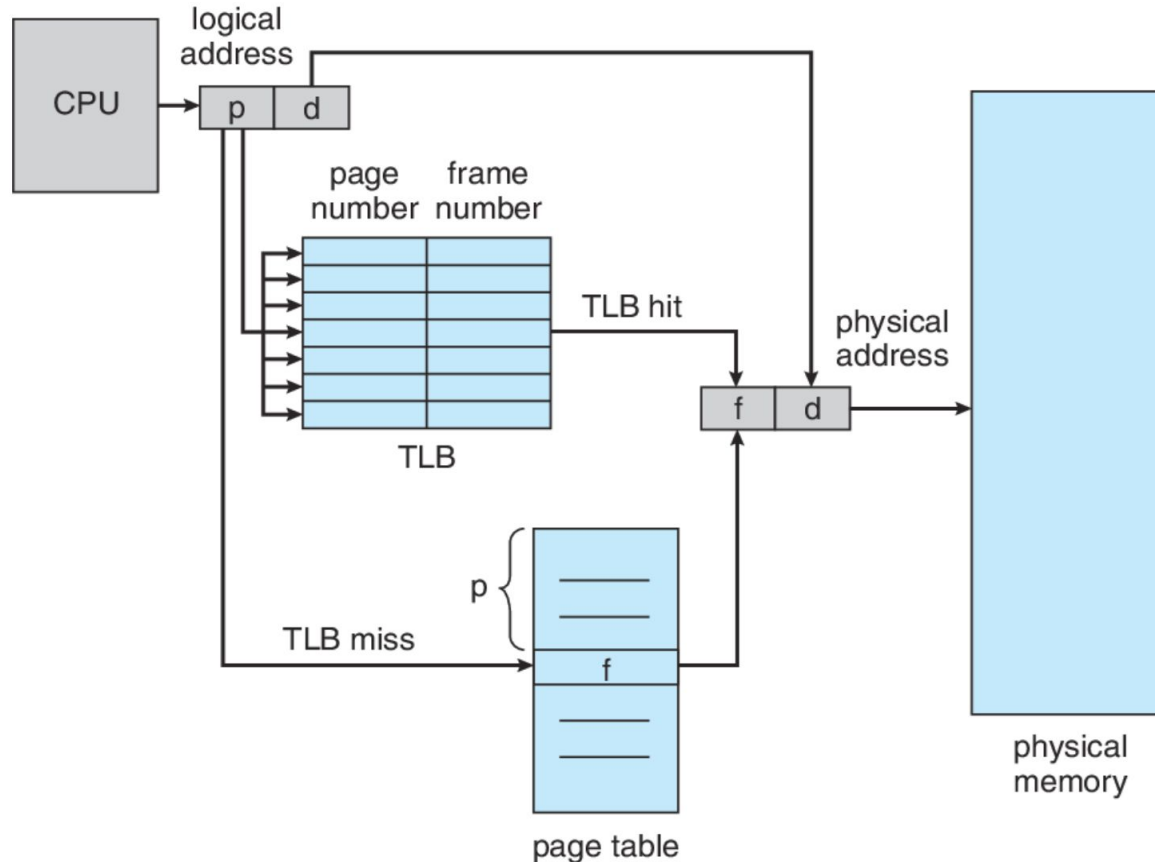
Implementation of Paging

- Page table is kept in main memory (why?)
- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PTLR)** indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
 - One for the page table and one for the data / instruction
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation lookaside buffers (TLBs)**

Implementation of Paging - TLB

- TLBs typically small (64 to 1,024 entries)
- On a **TLB miss**, value is loaded into the TLB for faster access next time
 - Replacement policies must be considered
 - Some entries can be **wired down** for permanent fast access
- Some TLBs store **Address-Space IDentifiers** (ASIDs) in each TLB entry – uniquely identifies each process to provide address-space protection for that process
 - Otherwise need to flush at every context switch
- An order of magnitude faster than memory access

Implementation of Paging - TLB



Effective Access Time

- **α** : Hit ratio - percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- **M** : memory access time
- **ϵ** : associative lookup time
- Effective Access Time (EAT):

$$\text{EAT} = (\text{Hit Rate} \times \text{Hit Time}) + (\text{Miss Rate} \times \text{Miss Time})$$

$$\text{EAT} = \alpha(M + \epsilon) + (1 - \alpha)(2M + \epsilon)$$

$$\text{If } M \gg \epsilon \rightarrow \text{EAT} = \alpha M + (1 - \alpha)(2M) = M + (1 - \alpha)M$$

- Example: consider $\alpha=80\%$ and $M=100\text{ns}$

$$\text{EAT} = 0.80 \times 100 + 0.2 \times 200 = 120\text{ns}$$


$$\text{If } \alpha=99\%, \text{EAT} = 101\text{ns}$$

Structure of the Page Table

- Memory structures for paging can get huge using straightforward methods
- Consider a 32-bit logical address space as on modern computers (2^{32} Bytes of Memory)
- Page size of 4 KB (2^{12} Bytes)
- Page table would have 1 million entries ($2^{32} / 2^{12}$)
- If each entry is 4 bytes \rightarrow 4 MB of physical address space / memory for page table alone
 - That amount of memory used to cost a lot
 - Don't want to allocate that contiguously in main memory

Hierarchical Page Tables

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table




Two-Level Paging Example

- A logical address (on 32-bit machine with 4K page size) is divided into:
 - a page number consisting of 20 bits
 - a page offset consisting of 12 bits
- Since the page table is paged, the page number is further divided into:
 - a 10-bit page table index (within the outer page directory)
 - a 10-bit page index (within the inner page table)

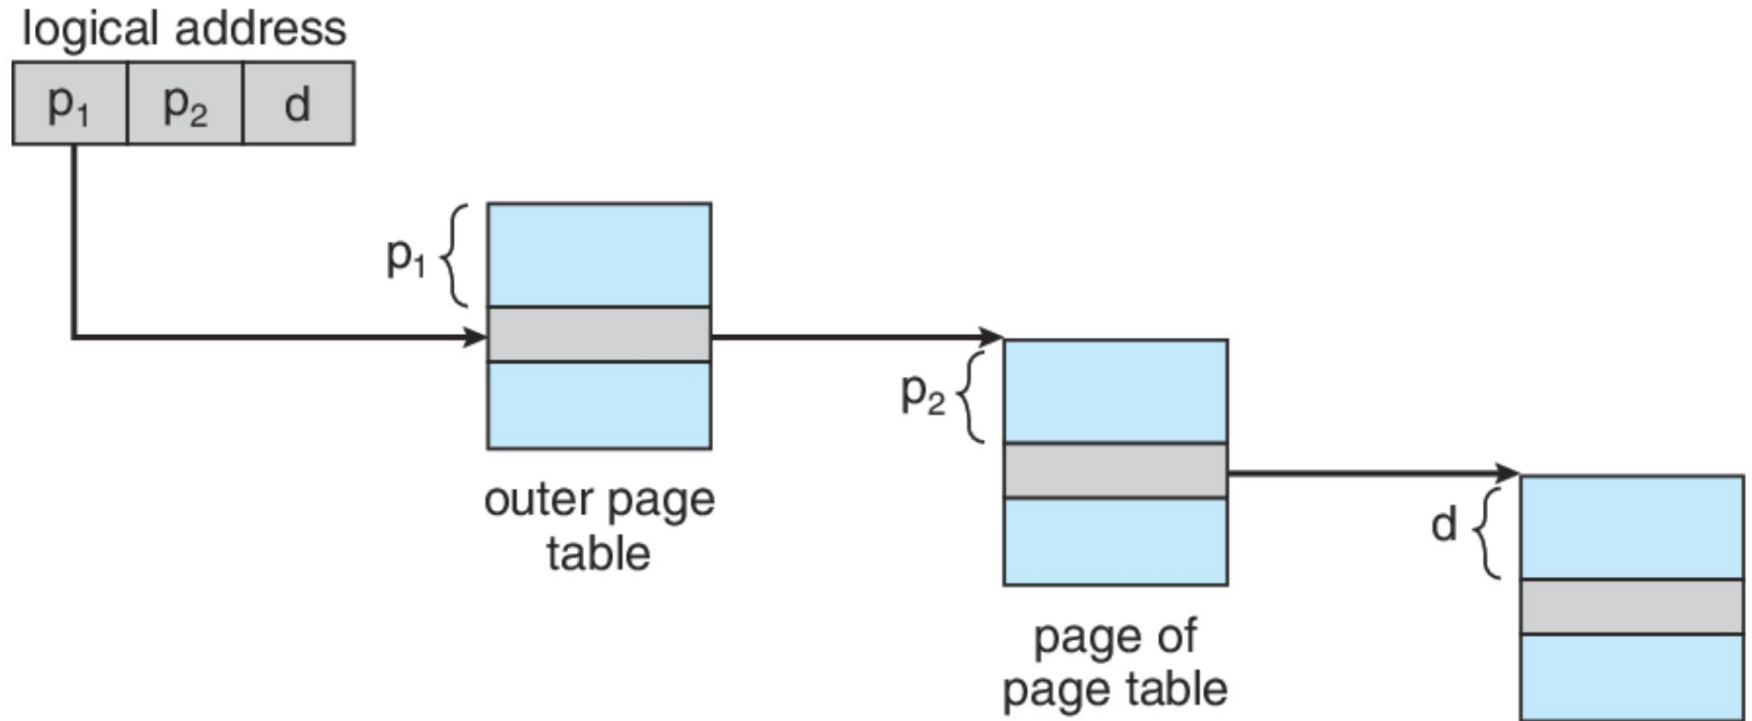
Two-Level Paging Example

- Thus, a logical address is as follows:



- where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table
- Known as **forward-mapped page table**

Two-Level Paging Example: Address Translation




Example: The Intel IA-32 Architecture

- Supports both segmentation and segmentation with paging
 - Each segment can be 4 GB
 - Up to 16 K segments per process
 - Divided into two partitions
 - First partition of up to 8 K segments are private to process **kept in local descriptor table (LDT)**
 - Second partition of up to 8K segments shared among all processes **kept in global descriptor table (GDT)**



Example: The Intel IA-32 Architecture

- CPU generates logical address
 - Selector given to segmentation unit
 - Which produces linear addresses

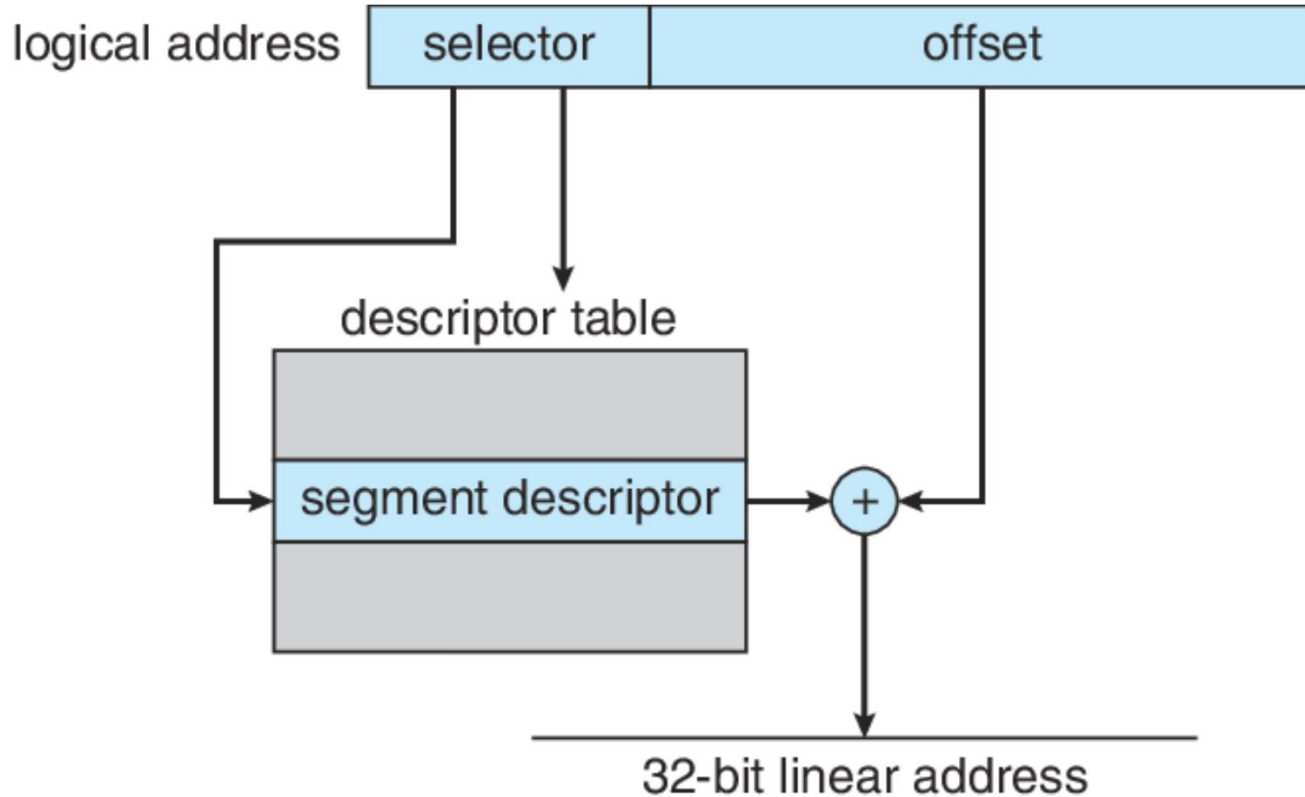


- Linear address given to paging unit
 - Which generates physical address in main memory
 - Pages sizes can be 4 KB or 4 MB
- Segmentation and paging units form equivalent of MMU


Logical to Physical Address Translation in IA-32



Intel IA-32 Segmentation



Intel IA-32 Paging Architecture



Acknowledgements

- “Operating Systems Concepts” book and supplementary material by A. Silberschatz, P. Galvin and G. Gagne
- “Operating Systems: Internals and Design Principles” book and supplementary material by W. Stallings
- “Modern Operating Systems” book and supplementary material by A. Tanenbaum
- R. Doursat and M. Yuksel from University of Nevada, Reno
- Farshad Ghanei from Illinois Tech
- T. Kosar and K. Dantu from University at Buffalo

Announcement

- Homework 3
 - **START EARLY**
 - Due on October 27th, 11.59 PM
- Quiz 4 next class