

Operating Systems Concepts

Process Synchronization: Semaphores, Mutex

CS 4375, Fall 2025

Instructor: MD Armanuzzaman (*Arman*)

marmanuzzaman@utep.edu

November 05, 2025

Summary

- Process Synchronization
 - Multiprogramming: multiple processes or threads may require consensus to work together
 - How do we achieve that? What issues may occur for shared data?
 - Race Condition
 - Critical-Section Problem
 - Different Solutions

Agenda

- Semaphores
- Classic Problems of Synchronization
 - Bounded Buffer Problem
 - Readers and Writers Problem
 - Dining Philosophers Problem
- Monitors
- Condition Variables
- Sleeping Barber Problem

Solution to Critical-Section Problem

- Summary of implementations of mutual exclusion
 - Implementation 1 - disabling hardware interrupts
 - NO: race condition avoided, but can crash the system!
 - Implementation 2 - simple lock variable (unprotected)
 - NO: still suffers from race condition
 - Implementation 3 - indivisible lock variable (TSL)
 - YES: works, but requires hardware support
 - Implementation 4 - no-TSL toggle for two threads
 - NO: race condition avoided inside, but lockup outside
 - Implementation 5 - Peterson's no-TSL, no-alternation
 - YES: works in software, but has processing overhead

This will be the
basis for “mutexes”

Solution to Critical-Section Problem

- Problem: All implementations (2-5) rely on busy waiting
 - “Busy waiting” means that the process/thread continuously executes a tight loop until some condition changes
 - Busy waiting is bad 😞
 - **Waste of CPU time**- the busy process is not doing anything useful, yet remains “READY” instead of “BLOCKED”
 - **Paradox of inverted priority**- by looping indefinitely, a higher-priority process B may starve a lower-priority process A, thus preventing A from exiting CR and liberating B! (B is working against its own interest)
- **We need for the waiting process to block, not keep idling!**

Semaphores

- Synchronization tool for critical section problem
- Semaphore S - Integer variable
- Can only be accessed through two standard operations:
 - **wait()** and **signal()**
 - **P()** and **V()**
 - Proberen/test ; Verhogen/increase (in Dutch)
- Classical implementation (using busy-waiting): (Without interruption)

```
wait (S) {  
    S--;  
    while (S <= 0); // loop  
}
```

```
signal (S) {  
    S++;  
}
```

Semaphores Without Busy-Waiting

```
wait(S) {  
    S.value--;  
    if (S.value < 0) {  
        // add this process to waiting queue  
        block();  
    }  
}  
  
signal(S) {  
    S.value++;  
    if (S.value <= 0) {  
        // remove process P from the waiting queue  
        wakeup(P);  
    }  
}
```

- We must guarantee that no two processes can execute wait() and signal() operations on the same semaphore at the same time.

Semaphores Without Busy-Waiting

```
typedef struct {

    int value;
    struct process *list;
} semaphore;

wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        // add this process to
        // waiting queue S->list
        block();
    }
}

signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        // remove process P from the
        // waiting queue S->list
        wakeup(P);
    }
}
```

Semaphores as Synchronization Tool

- **Counting semaphore** - Integer value can range over an unrestricted domain
- **Binary semaphore** - Integer value can range only between 0 and 1
 - Also known as **mutex locks**
- Provides mutual exclusion

```
Semaphore s; //initialized to N  
  
wait(s);  
  
/* critical section */  
  
signal(s);
```

```
Lock s; //initialized to 1  
  
acquire(s);  
  
/* critical section */  
  
release(s);
```

Bounded Buffer Problem

- Shared buffer with **N slots** to store at most **N items**
- **Producer** processes data items and puts into the buffer
- **Consumer** gets the data items from the buffer
- Variable **empty** keeps number of empty slots in the buffer
- Variable **full** keeps number of full items in the buffer

Bounded Buffer Problem - Solution 1

- Implementation 1 - 1 semaphore

Producer Process

```
int empty = N, full = 0;  
  
while (true) {  
  
    /* produce an item */  
  
    while (empty == 0); // loop  
  
    wait(mutex);  
  
    // add the item to the buffer  
  
    empty--; full++;  
  
    signal(mutex);  
  
}
```

Consumer Process

```
while (true) {  
  
    wait(mutex);  
  
    if (full > 0) {  
  
        // remove item from buffer  
  
        full--; empty++;  
  
    }  
  
    signal(mutex);  
  
    /* consume the item */  
  
}
```

Bounded Buffer Problem - Solution 1

- Implementation 1 - 1 semaphore

Producer Process

```
int empty = N, full = 0;  
  
while (true) {  
  
    /* produce an item */  
  
    while (empty == 0); // loop  
  
    wait(mutex);  
  
    // add the item to the buffer  
  
    empty--; full++;  
  
    signal(mutex);  
  
}
```

Consumer Process

```
while (true) {  
  
    wait(mutex);  
  
    if (full > 0) {  
  
        // remove item from buffer  
  
        full--; empty++;  
  
    }  
  
    signal(mutex);  
  
    /* consume the item */  
  
}
```

Bounded Buffer Problem - Solution 1

- Implementation 1 - 1 semaphore

Consumes non-existing item!

Consumer Process

```
while (true) {  
    wait(mutex);  
    if (full > 0) {  
        // remove item from buffer  
        full--; empty++;  
    }  
    signal(mutex);  
    /* consume the item */  
}
```

Bounded Buffer Problem - Solution 2

- Implementation 2 - 1 semaphore

Producer Process

```
int empty=N, full=0;  
  
while (true) {  
  
    /* produce an item */  
  
    while (empty == 0); // loop  
  
    wait(mutex);  
  
    // add the item to the buffer  
  
    empty--; full++;  
  
    signal(mutex);  
  
}
```

Consumer Process

```
while (true) {  
  
    while (full == 0); // loop  
  
    wait(mutex);  
  
    // remove item from buffer  
  
    full--; empty++;  
  
    signal(mutex);  
  
    /* consume the item */  
  
}
```

Bounded Buffer Problem - Solution 2

- Implementation 2 - 1 semaphore

Producer Process

```
int empty=N, full=0;  
  
while (true) {  
  
    /* produce an item */  
  
    while (empty == 0); // loop  
  
    wait(mutex);  
  
    // add the item to the buffer  
  
    empty--; full++;  
  
    signal(mutex);  
  
}
```

Consumer Process

```
while (true) {  
  
    while (full == 0); // loop  
  
    wait(mutex);  
  
    // remove item from buffer  
  
    full--; empty++;  
  
    signal(mutex);  
  
    /* consume the item */  
  
}
```

Bounded Buffer Problem - Solution 2

- Implementation 2 - 1 semaphore

Mutual exclusion is not preserved!

Consumer Process

```
while (true) {  
    while (full == 0); // loop  
    wait(mutex);  
    // remove item from buffer  
    full--; empty++;  
    signal(mutex);  
    /* consume the item */  
}
```

Bounded Buffer Problem - Solution 3

- Implementation 3 - 2 semaphore

Producer Process

```
int empty = N, full = 0;  
  
while (true) {  
  
    /* produce an item */  
  
    wait(empty);  
  
    // add the item to the buffer  
  
    signal(full);  
  
}
```

Consumer Process

```
while (true) {  
  
    wait(full);  
  
    // remove item from buffer  
  
    signal(empty);  
  
    /* consume the item */  
  
}
```

Bounded Buffer Problem - Solution 3

- Implementation 3 - 2 semaphore

Producer Process

```
int empty = N, full = 0;  
  
while (true) {  
    /* produce an item */  
  
    wait(empty);  
  
    // add the item to the buffer  
  
    signal(full);  
  
}
```

Consumer Process

```
while (true) {  
    wait(full);  
  
    // remove item from buffer  
  
    signal(empty);  
  
    /* consume the item */  
}
```

Bounded Buffer Problem - Solution 3

- Implementation 3 - 2 semaphore

Mutual exclusion is not preserved!

Consumer Process

```
while (true) {  
    wait(full);  
    // remove item from buffer  
    signal(empty);  
    /* consume the item */  
}
```

Bounded Buffer Problem - Solution 4

- Implementation 4 - 3 semaphore
 - Semaphore **mutex** to access the buffer
 - Initialized to 1
 - Semaphore **full** (number of full buffers)
 - Initialized to 0
 - Semaphore **empty** (number of empty buffers)
 - Initialized to N

Bounded Buffer Problem - Solution 4

- Implementation 4 - 3 semaphore

Producer Process

```
int empty = N, full = 0;  
  
while (true) {  
    /* produce an item */  
  
    wait(empty);  
  
    wait(mutex);  
  
    // add the item to the buffer  
  
    signal(mutex);  
  
    signal(full);  
  
}
```

Consumer Process

```
while (true) {  
    wait(full);  
  
    wait(mutex);  
  
    // remove item from buffer  
  
    signal(mutex);  
  
    signal(empty);  
  
    /* consume the item */  
}
```

Bounded Buffer Problem

- Implementation 1: Using 1 semaphore
- Implementation 2: Using 1 semaphore
- Implementation 3: Using 2 semaphores
- Implementation 4: Using 3 semaphores

Readers and Writers Problem

- Multiple *Readers* and *Writers* concurrently accessing the same database.
- Multiple *Readers* accessing at the same time → OK
- When there is a *Writer* accessing, there should be no other process accessing at the same time

Readers and Writers Problem

Reader process

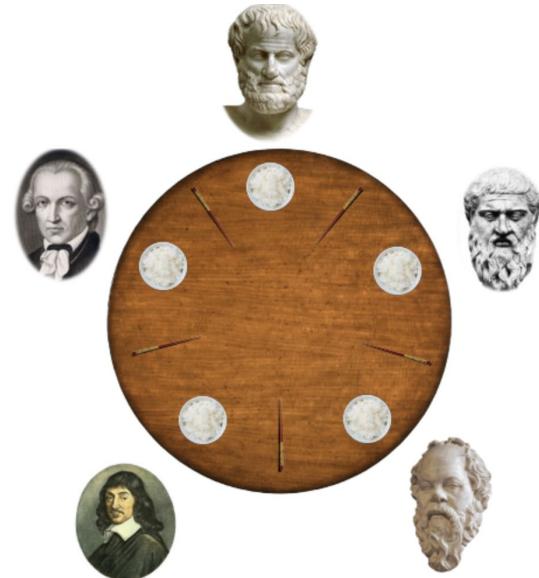
```
while (true) {  
    wait(mutex);  
    readercount++;  
    if (readercount == 1)  
        wait (wrt_mutex);  
    signal(mutex);  
    /* read from database */  
    wait(mutex);  
    readercount--;  
    if (readercount == 0)  
        signal (wrt_mutex);  
    signal(mutex);  
}
```

Writer Process

```
while (true) {  
    wait(wrt_mutex);  
    /* write to database */  
    signal(wrt_mutex);  
}
```

Dining Philosophers Problem

- Five philosophers spend their time eating and thinking
- They are sitting in front of a round table with spaghetti served.
- There are five plates at the table and five chopsticks set between the plates.
- Eating the spaghetti **requires the use of two chopsticks**, which the philosophers pick up **one at a time**.
- Philosophers do not talk to each other.
- Semaphores **chopstick [5]** are initialized to 1



Dining Philosophers Problem

Philosopher i process

```
while (true) {  
    wait( chopstick[i] );  
    wait( chopstick[(i+1)%5] );  
    /* eat */  
    signal( chopstick[(i+1)%5] );  
    signal( chopstick[i] );  
    /* think */  
}
```

Dining Philosophers Problem

Philosopher i process

```
while (true) {  
    wait( chopstick[i] );  
    wait( chopstick[(i+1)%5] );  
    /* eat */  
    signal( chopstick[(i+1)%5] );  
    signal( chopstick[i] );  
    /* think */  
}
```

Mutual exclusion is ensured, however it may result in deadlock!

Dining Philosophers Problem - Deadlock

- **Deadlock** happens when two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.
- Let **S** and **Q** be two semaphores initialized to 1.

P1

```
wait(S) ;
```

.

```
wait(Q) ;
```

.

.

```
signal(S) ;
```

```
signal(Q) ;
```

P2

```
wait(Q) ;
```

.

```
wait(S) ;
```

.

.

```
signal(Q) ;
```

```
signal(S) ;
```

Dining Philosophers Problem

- To prevent deadlock:
 - Allow a philosopher to pick up his chopsticks only if both chopsticks are available (*i.e.* in critical section)
 - Use an asymmetric solution: An odd numbered philosopher picks up first his left chopstick and then his right chopstick; An even numbered philosopher first picks up his right chopstick, then his left chopstick.
 - **Exercise: Write the algorithms for the above solutions**

Semaphores - Wrong Use of Operations

- Semaphores A and B, initialized to 1

P1

wait(A) ;

wait(B) ;

P2

wait(B) ;

wait(A) ;

- Deadlock
- signal(mutex) ... wait(mutex)
 - Violation of mutual exclusion
- wait(mutex) ... wait(mutex)
 - Deadlock
- Omitting of wait(mutex) or signal(mutex) (or both)
 - Violation of mutual exclusion or deadlock

Semaphores

- Semaphores are inadequate in dealing with deadlocks
- Do not protect the programmer from the easy mistakes of taking a semaphore that is already held by the same process, and forgetting to release a semaphore that has been taken.
- Mostly used in low level code (e.g. operating systems)
- The trends in programming language development, is towards more structured forms of synchronizations such as **monitors**.

Monitor

- A high-level abstractions that provides convenient and effective mechanism for process synchronization
- Only one process may be active within the monitor at a time

```
monitor monitor_name {  
    // shared variable declarations  
    procedure P1(...) {...}  
    ...  
    procedure Pn(...) {...}  
    initialization_code(...){...}  
    ...  
}
```

- A monitor procedure takes the lock before doing anything else, and holds it until it either finishes or waits for a condition.

Monitor Example

- As a simple example, consider a monitor for performing transactions on bank account

```
monitor account {  
    int balance := 0  
  
    function withdraw(int amount) {  
        if (amount < 0) then error "Amount may not be negative"  
        else if (balance < amount) then error "Insufficient funds"  
        else balance := balance - amount  
    }  
  
    function deposit(int amount) {  
        if (amount < 0) then error "Amount may not be negative"  
        else balance := balance + amount  
    }  
}
```

Monitor Example

```
class Account {  
    private lock myLock;  
    private int balance := 0  
    invariant (balance >= 0)  
    public method boolean withdraw(int amount)  
        precondition (amount >= 0) {  
            myLock.acquire();  
            try:  
                if (balance < amount) then return false  
                else { balance := balance - amount ; return  
                    true }  
            finally:  
                myLock.release();  
        }  
    public method deposit(int amount)  
        precondition (amount >= 0) {  
            myLock.acquire();  
            try:  
                balance := balance + amount  
            finally:  
                myLock.release();  
        }  
}
```

By hiding the details of the synchronization code from the programmer, mutual exclusion is inherently provided by monitors, so programmer does not need to manage these locks manually.

Condition Variables

- For many applications, mutual exclusion is not sufficient
- A thread may need to wait until some condition holds true
- We don't want to use busy waiting...
- Condition variables queue threads until a certain condition is met (non-blocking, non-busy-waiting)
- Two operations on a condition variable x:
 - `x.wait()`– a thread invoking this operation is suspended
 - `x.signal()`– resumes one of the threads (if any) that invoked `x.wait()`
- If no thread was suspended, `x.signal()` operation has no effect (unlike semaphores)

Dining Philosophers Problem - Solution with Monitor

- Implementation using Monitors and condition variables

```
monitor DP {  
    enum {THINKING; HUNGRY, EATING} state[5] ;  
    condition self[5]; // to delay philosopher when he is hungry  
                      // but unable to get chopsticks  
    initialization_code() {  
        for (int i = 0; i < 5; i++)  
            state[i] = THINKING;  
    }  
    void pickup (int i) {  
        state[i] = HUNGRY;  
        test(i); // only if both neighbors are not eating  
        if (state[i] != EATING) self[i].wait();  
    }  
}
```

Dining Philosophers Problem - Solution with Monitor

```
void test (int i) {  
    if ((state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) &&  
        (state[(i + 4) % 5] != EATING) ) {  
        state[i] = EATING;  
        self[i].signal();  
    }  
}  
  
void putdown (int i) {  
    state[i] = THINKING; // test left and right neighbors  
    test((i + 4) % 5);  
    test((i + 1) % 5);  
}
```

Dining Philosophers Problem - Solution with Monitor

- Main thread will run the initialization code

...

```
DiningPhilosophers.initialization_code();
```

...

- Each Dining Philosopher thread will run the **pickup** and **putdown** functions for Philosopher thread *i*:

...

```
DiningPhilosophers.pickup(i);
```

...

```
eat
```

...

```
DiningPhilosophers.putdown(i);
```

...

Dining Philosophers Problem - Solution with Monitor

- No two philosophers eat at the same time
- No deadlock
- But **starvation** can occur!
 - Indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.
 - How can we prevent this? **Exercise!**

Sleeping Barber Problem

- Based upon a hypothetical barber shop with **one barber, one barber chair, and a number of chairs** for waiting customers
- When there are no customers, the barber sits in his chair and sleeps
- As soon as a customer arrives, he either **awakens** the barber or, if the barber is cutting someone else's hair, **sits down** in one of the vacant chairs
- If all of the chairs are **occupied**, the newly arrived customer simply **leaves**

Sleeping Barber Problem - Solution

- Use **three semaphores**: one for any **waiting customers**, one for the **barber** (to see if he is idle), and a **mutex**
- When a customer arrives, he attempts to **acquire the mutex**, and waits until he has succeeded.
- The customer then checks to see if there is an empty chair for him (either one in the waiting room or the barber chair), and if **none** of these are **empty**, the customer checks again later (in a **loop**).
- Otherwise the customer takes a seat – thus **reducing** the number available (a **critical section**).

Sleeping Barber Problem - Solution

- The customer then signals the barber to **awaken** through his **semaphore**, and the mutex is released to allow other customers (or the barber) the ability to **acquire** it
- If the barber is not free, the customer then waits. The barber sits in a perpetual waiting loop, being awakened by any **waiting** customers. Once he is awoken, he signals the **waiting** customers through their **semaphore**, allowing them to get their hair cut one at a time.

Sleeping Barber Problem - Solution

```
Semaphore Customers          // to wait for available customers
Semaphore Barber             // to wait for available barber
Lock accessSeats             // mutex lock, to change seat availability
int NumberOfFreeSeats

The Barber(Thread):
while (true) {                // runs in an infinite loop
    Customers.wait()          // tries to acquire a customer
                               // if none is available he's going to sleep
    accessSeats.acquire() // at this time he has been awoken
                           // -> want to modify the number of available seats
    NumberOfFreeSeats++ // one chair becomes free
    Barber.signal()        // the barber is ready to cut
    accessSeats.release() // we don't need the lock on the chairs anymore
    /* barber cuts hair */
}
```

Sleeping Barber Problem - Solution

The Customer(Thread):

```
needCut = true                                // the customer decides that they need haircut
while (needCut) {                               // as long as the customer is not cut
    accessSeats.acquire()                      // tries to get access to the chairs
    if (NumberOfFreeSeats > 0) {               // if there are any free seats
        NumberOfFreeSeats --                   // sitting down on a chair
        Customers.signal()                   // notify the barber, who's waiting
                                                // for a customer
        accessSeats.release()                 // don't need to lock the chairs anymore
        Barber.wait()                        // now it's this customers turn, but
                                                // needs to wait if the barber is busy
        needCut = false
    } else {                                    // there are no free seats (tough luck!)
        accessSeats.release()                 // but don't forget to
                                                // release the lock on the seats
    }
}
```

Acknowledgements

- “Operating Systems Concepts” book and supplementary material by A. Silberschatz, P. Galvin and G. Gagne
- “Operating Systems: Internals and Design Principles” book and supplementary material by W. Stallings
- “Modern Operating Systems” book and supplementary material by A. Tanenbaum
- R. Doursat and M. Yuksel from University of Nevada, Reno
- Farshad Ghanei from Illinois Tech
- T. Kosar and K. Dantu from University at Buffalo

Announcement

- Homework 4
 - Due on November 10th at 11.59PM
- Quiz 5
 - Released on black board
 - Due tomorrow November 6th, 11.59PM
- Please check the final exam schedule
 - December 12th, 10 AM - 12.45 PM