

Operating Systems Concepts

Dead Locks (cont.)

CS 4375, Fall 2025

Instructor: MD Armanuzzaman (*Arman*)

marmanuzzaman@utep.edu

November 12, 2025

Summary

- The Deadlock Problem
 - Characterization of Deadlock
 - Resource Allocation Graph
 - Handling Deadlocks
 - Deadlock Prevention
 - Deadlock Avoidance
 - Deadlock Detection
 - Deadlock Recovery

Agenda

- The Deadlock Problem
 - Characterization of Deadlock
 - Resource Allocation Graph
 - Handling Deadlocks
 - Deadlock Prevention
 - **Deadlock Avoidance**
 - **Deadlock Detection**
 - **Deadlock Recovery**

Deadlock Avoidance

- **Deadlock Prevention:** Prevents deadlocks by restraining resources and making sure **one of 4 necessary conditions for a deadlock does not hold.** (system design)
 - Possible side effect: **low device utilization and reduced system throughput**
- **Deadlock Avoidance:** Requires that the system has some additional a priori information available. (dynamic request check)
 - E.g. request disk and then printer
 - E.g. request at most n resources.
 - Allows more concurrency
- Similar to the difference between a **traffic light** and a **police officer** directing the traffic!

Deadlock Avoidance

- Simplest and most useful model requires that each process declare the **maximum number** of resources of each type that it may need.
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a **circular-wait** condition.
- **Resource-allocation state** is defined by the number of available and allocated resources, and the maximum demands of the processes.

Deadlock Avoidance: Safe State

- A state is **safe** if the system can **allocate** resources to each process (up to its maximum) in some order and can still **avoid** a deadlock.
- When a process requests an available resource, system **must decide** if immediate allocation leaves the system in a safe state.
- System is in safe state if there exists a **safe sequence** of all processes.

Deadlock Avoidance: Safe State

- Sequence $\langle P_1, P_2, \dots, P_n \rangle$ is **safe** if for each P_i , the resources that P_i can still request, can be satisfied by currently available resources + resources held by all the P_j , with $j < i$
 - If P_i resource needs are not **immediately available**, then P_i can **wait until all P_j have finished**.
 - When P_j is **finished**, P_i can obtain **needed resources**, execute, return allocated resources, and terminate.
 - When P_i **terminates**, P_{i+1} **can obtain** its needed resources, and so on.
- If no such sequence exists, the state is **unsafe!**

Deadlock Avoidance: Safe State Example

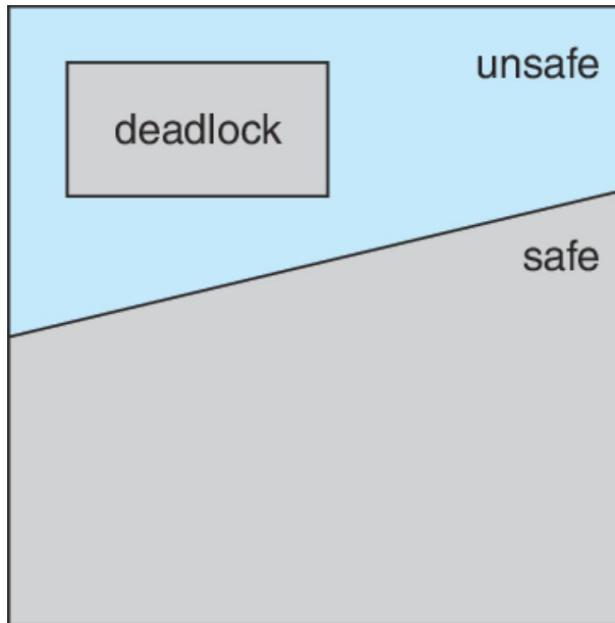
- We have five processes P_1 through P_5 , three resource types are available: A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time T_0 looks like:

	Allocation			Request			Available			Work		
	A	B	C	A	B	C	A	B	C	A	B	C
P_1	0	1	0	0	0	0	0	0	0	0	0	0
P_2	2	0	0	2	0	2						
P_3	3	0	3	0	0	0						
P_4	2	1	1	1	0	0						
P_5	0	0	2	0	0	2						

- Sequence $\langle P_1, P_3, P_4, P_2, P_5 \rangle$ represents a safe sequence for the current state

Deadlock Avoidance: Safe State Facts

- If a system is in safe state → There is no deadlock
- If a system is in unsafe state → Possibility of deadlock
- Deadlock Avoidance → Ensure that system will never enter unsafe state



Deadlock Avoidance: Safe State Example

- Consider a system with 3 processes and 12 disks.
- Snapshot at time T_0 looks like:

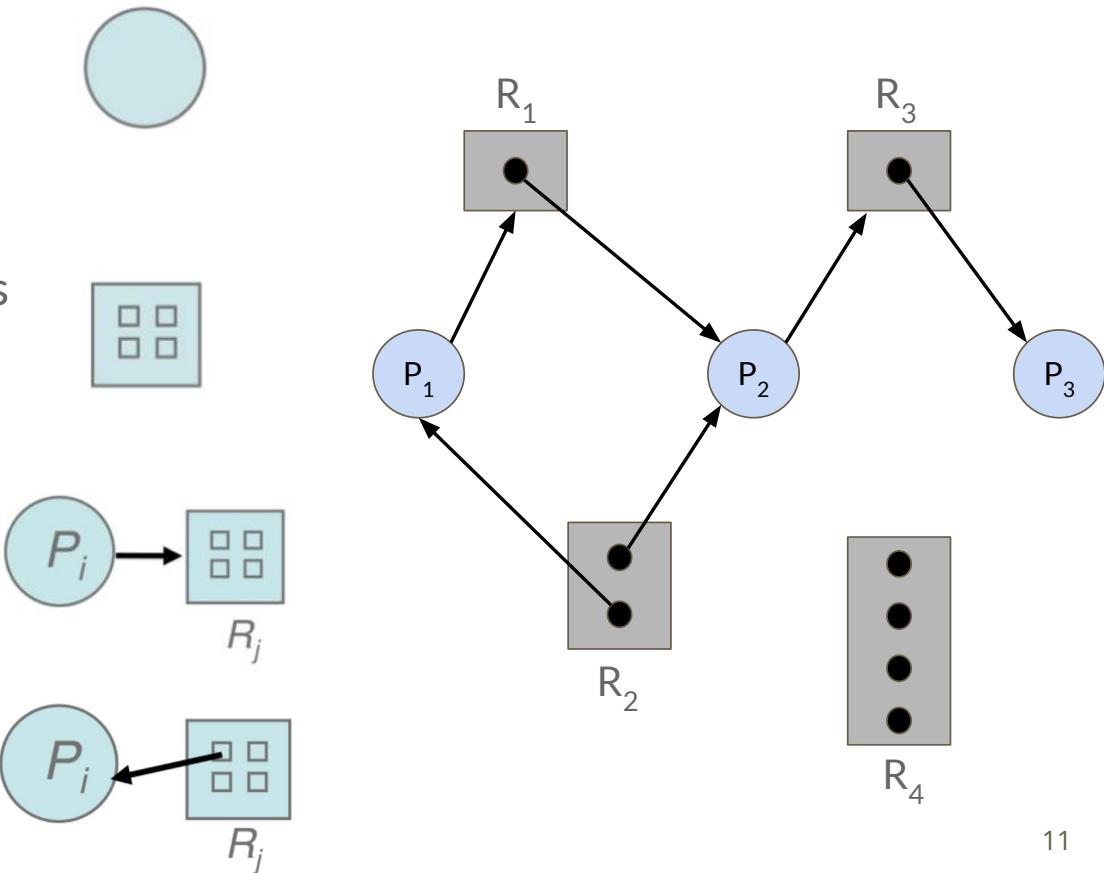
	Maximum Needs	Current Allocation
P_1	10	5
P_2	4	2
P_3	9	2

- Snapshot at time T_1 looks like:

	Maximum Needs	Current Allocation
P_1	10	5
P_2	4	2
P_3	9	3

Resource Allocation Graph

- Process
- Resource type with 4 instances

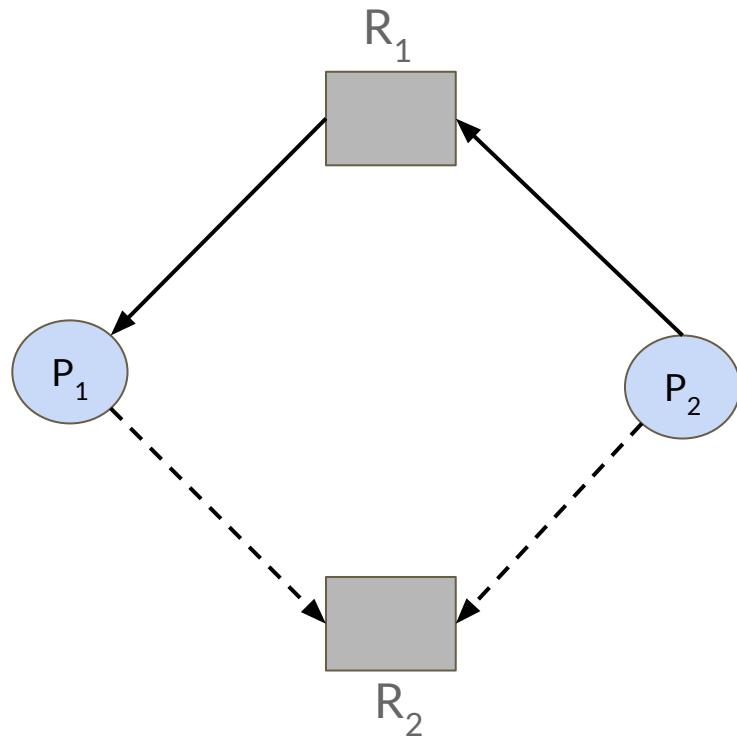


- P_i requests instance of R_j
- P_i is holding an instance of R_j

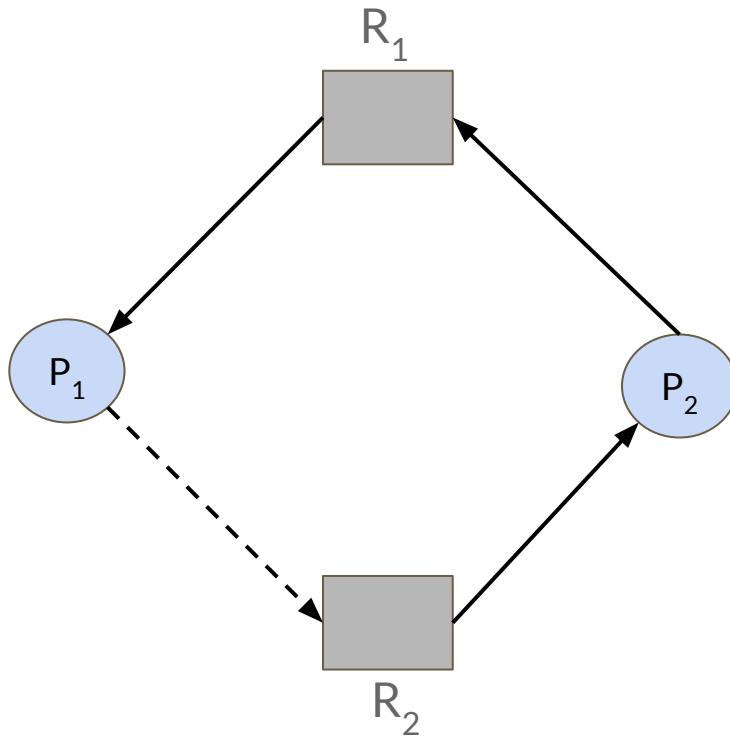
Deadlock Avoidance: Resource Allocation Graph

- Claim edge $P_i \rightarrow R_j$ indicated that process P_i may request resource R_j (dashed line)
- Claim edge converts to request edge when a process requests a resource.
- When a resource is released by a process, assignment edge reconverts to a claim edge.
- If a new allocation results in a cycle in the graph (unsafe state), we cannot allow that allocation.
- Resources must be claimed a priori in the system

Deadlock Avoidance: Resource Allocation Graph



Deadlock Avoidance: Resource Allocation Graph



Deadlock Avoidance: Banker's Algorithm

- Works for multiple resource instances.
- Each process declares maximum # of resources it may need.
- When a process requests a resource, it may have to wait if this leads to an unsafe state.
- When a process gets all its resources, it must return them in a finite amount of time.

Banker's Algorithm: Data Structure

- Let n = number of processes, and let m = number of resources types
- **Available:** Vector of length m . If $\text{Available}[j] = k$, there are k instances of resource type R_j available
- **Max:** $n \times m$ matrix. If $\text{Max}[i, j] = k$, then process P_i may request at most k instances of resource type R_j
- **Allocation:** $n \times m$ matrix. If $\text{Allocation}[i, j] = k$, then process P_i is currently allocated k instances of resource type R_j
- **Need:** $n \times m$ matrix. If $\text{Need}[i, j] = k$, then process P_i may need k more instances of resource type R_j to complete its task

$$\text{Need}[i, j] = \text{Max}[i, j] - \text{Allocation}[i, j]$$

Banker's Algorithm: Safety Algorithm

1. Let Work and Finish be vectors of length m and n , respectively. Initialize:

$\text{Work} = \text{Available}$

$\text{Finish}[i] = \text{false}$ for $i = 0, 1, \dots, n-1$

2. Find an i such that both:

a. $\text{Finish}[i] == \text{false}$

b. $\text{Need}_i \leq \text{Work}$

If no such i exists, go to step 4.

3. $\text{Work} = \text{Work} + \text{Allocation}_i$

$\text{Finish}[i] = \text{true}$

Go to step 2.

4. If $\text{Finish}[i] == \text{true}$ for all i , then the system is in a safe state

Banker's Algorithm: Safety Algorithm

1. Let Work and Finish be vectors of length m and n , respectively. Initialize:

$\text{Work} = \text{Available}$

$\text{Finish}[i] = \text{false}$ for $i = 0, 1, \dots, n-1$

2. Find an i such that both:

a. $\text{Finish}[i] == \text{false}$

b. $\text{Need}_i \leq \text{Work}$

If no such i exists, go to step 4.

3. $\text{Work} = \text{Work} + \text{Allocation}_i$

$\text{Finish}[i] = \text{true}$

Go to step 2.

4. If $\text{Finish}[i] == \text{true}$ for all i , then the system is in a safe state

Algorithm requires an order of $O(m \times n^2)$ operations to detect whether the system is in safe state.

Banker's Algorithm: Resource-Request for Process P_i

- Let Request_i be the request vector for process P_i .
If $\text{Request}_i[j] = k$ then process P_i wants k instances of resource type R_j .
 - If $\text{Request}_i \leq \text{Need}_i$, go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
 - If $\text{Request}_i \leq \text{Available}$, go to step 3. Otherwise P_i must wait, since resources are not available.
 - Pretend to allocate requested resources to P_i by modifying the state as follows:
 - $\text{Available} = \text{Available} - \text{Request}_i$,
 - $\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i$,
 - $\text{Need}_i = \text{Need}_i - \text{Request}_i$;
 - If safe \rightarrow the resources are allocated to P_i ,
 - If unsafe $\rightarrow P_i$ must wait, and the old resource-allocation state is restored

Banker's Algorithm: Example

- We have five processes P_1 through P_5 , three resource types are available: A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T_0 looks like:

	Allocation			Max			Available		
	A	B	C	A	B	C	A	B	C
P_1	0	1	0	7	5	3	3	3	2
P_2	2	0	0	3	2	2			
P_3	3	0	2	9	0	2			
P_4	2	1	1	2	2	2			
P_5	0	0	2	4	3	3			

Banker's Algorithm: Example

- We have five processes P_1 through P_5 , three resource types are available: A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T_0 looks like:

	Allocation			Max			Available			Need		
	A	B	C	A	B	C	A	B	C	A	B	C
P_1	0	1	0	7	5	3	3	3	2	7	4	3
P_2	2	0	0	3	2	2				1	2	2
P_3	3	0	2	9	0	2				6	0	0
P_4	2	1	1	2	2	2				0	1	1
P_5	0	0	2	4	3	3				4	3	1

- The content of the matrix **Need** is defined to be **Max - Allocation**

Banker's Algorithm: Example

- We have five processes P_1 through P_5 , three resource types are available: A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T_0 looks like:

	Allocation			Max			Available			Need		
	A	B	C	A	B	C	A	B	C	A	B	C
P_1	0	1	0	7	5	3	3	3	2	7	4	3
P_2	2	0	0	3	2	2				1	2	2
P_3	3	0	2	9	0	2				6	0	0
P_4	2	1	1	2	2	2				0	1	1
P_5	0	0	2	4	3	3				4	3	1

- The system is in a safe state since the sequence $\langle P_2, P_4, P_5, P_3, P_1 \rangle$ satisfies safety

Banker's Algorithm: Example, P2 Requests (1, 0, 2)

- Check that Request \leq Available. i.e $(1, 0, 2) \leq (3, 3, 2) \rightarrow \text{true}$
- Snapshot at time T_0 looks like:

	Allocation			Max			Available			Need		
	A	B	C	A	B	C	A	B	C	A	B	C
P ₁	0	1	0	7	5	3	3	3	2	7	4	3
P ₂	2	0	0	3	2	2				1	2	2
P ₃	3	0	2	9	0	2				6	0	0
P ₄	2	1	1	2	2	2				0	1	1
P ₅	0	0	2	4	3	3				4	3	1

Banker's Algorithm: Example, P2 Requests (1, 0, 2)

- Check that Request \leq Available. i.e $(1, 0, 2) \leq (3, 3, 2) \rightarrow \text{true}$
- Snapshot at time T_1 looks like:

	Allocation			Max			Available			Need		
	A	B	C	A	B	C	A	B	C	A	B	C
P ₁	0	1	0	7	5	3	2	3	0	7	4	3
P ₂	3	0	2	3	2	2				0	2	0
P ₃	3	0	2	9	0	2				6	0	0
P ₄	2	1	1	2	2	2				0	1	1
P ₅	0	0	2	4	3	3				4	3	1

Banker's Algorithm: Example

- Executing safety algorithm shows that sequence $\langle P_2, P_4, P_5, P_3, P_1 \rangle$ satisfies safety requirement.
- Can request for $(3, 3, 0)$ by P_5 be granted?
- Can request for $(0, 2, 0)$ by P_1 be granted?

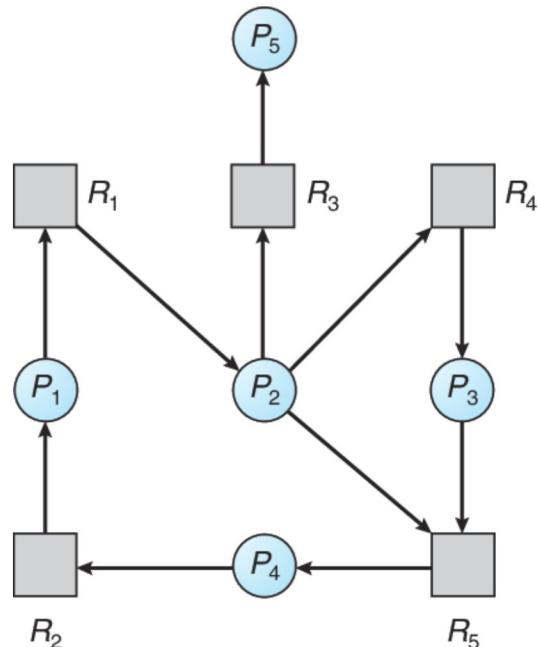
Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme

Deadlock Detection: Single-Instance Resources

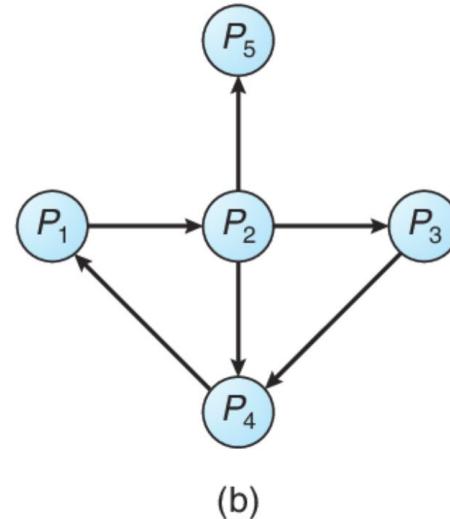
- Maintain ***wait-for*** graph
 - Nodes are processes
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j

Deadlock Detection: Single-Instance Resources



(a)

Resource Allocation Graph



(b)

Corresponding wait-for graph

Deadlock Detection: Single-Instance Resources

- Periodically invoke an algorithm that searches for a cycle in the graph.
- An algorithm to detect a cycle in a graph requires an order of n^2 operations, where n is the number of vertices in the graph.
- Only good for single-instance resource allocation systems.

Deadlock Detection: Several-Instance Resources

- **Available:** Vector of length m , indicates the number of available resources of each type.
- **Allocation:** Matrix of size $n \times m$, defines the number of resources of each type currently allocated to each process.
- **Request:** Matrix of size $n \times m$, indicates the current request of each process.

If $\text{Request}[i, j] = k$, then process P_i is requesting k more instances of resource type R_j

Deadlock Detection: Detection Algorithm

1. Let Work and Finish be vectors of length m and n , respectively. Initialize:

$\text{Work} = \text{Available}$

$\text{Finish}[i] = \text{false}$ for $i = 0, 1, \dots, n-1$

2. Find an i such that both:

a. $\text{Finish}[i] == \text{false}$

b. $\text{Request}_i \leq \text{Work}$

If no such i exists, go to step 4.

3. $\text{Work} = \text{Work} + \text{Allocation}_i$

$\text{Finish}[i] = \text{true}$

Go to step 2.

4. If $\text{Finish}[i] == \text{false}$ for some $i, 0 \leq i < n$, then the system is in deadlock (with all P_i s)

Deadlock Detection: Detection Algorithm

1. Let Work and Finish be vectors of length m and n , respectively. Initialize:

$\text{Work} = \text{Available}$

$\text{Finish}[i] = \text{false}$ for $i = 0, 1, \dots, n-1$

2. Find an i such that both:

a. $\text{Finish}[i] == \text{false}$

b. $\text{Request}_i \leq \text{Work}$

If no such i exists, go to step 4.

3. $\text{Work} = \text{Work} + \text{Allocation}_i$

$\text{Finish}[i] = \text{true}$

Go to step 2.

4. If $\text{Finish}[i] == \text{false}$ for some $i, 0 \leq i < n$, then the system is in deadlock (with all P_i s)

Algorithm requires an order of $O(m \times n^2)$ operations to detect whether the system is in deadlocked state.

Deadlock Detection: Detection Algorithm Example

- We have five processes P_1 through P_5 , three resource types are available: A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time T_0 looks like:

	Allocation			Request			Available			Work		
	A	B	C	A	B	C	A	B	C	A	B	C
P_1	0	1	0	0	0	0	0	0	0	0	0	0
P_2	2	0	0	2	0	2						
P_3	3	0	3	0	0	0						
P_4	2	1	1	1	0	0						
P_5	0	0	2	0	0	2						

- Sequence $\langle P_1, P_3, P_4, P_2, P_5 \rangle$ will result in $Finish[i] = true$ for all i

Deadlock Detection: Detection Algorithm Example

- P_3 requests an additional instance of type C
- What is the state of system?

	Allocation			Request			Available			Work		
	A	B	C	A	B	C	A	B	C	A	B	C
P_1	0	1	0	0	0	0	0	0	0	0	0	0
P_2	2	0	0	2	0	2						
P_3	3	0	3	0	0	1						
P_4	2	1	1	1	0	0						
P_5	0	0	2	0	0	2						

- Deadlock exists, consisting of processes P_2, P_3, P_4 , and P_5

Deadlock Recovery

- Process termination
- Resource preemption

Deadlock Recovery: Process Termination

- Abort all deadlocked processes
 - **Expensive!**
- Abort **one process** at a time until the deadlock cycle is eliminated
 - **Overhead** of deadlock detection algorithm.
- In which order should we choose to abort?
 - Priority of the process
 - How long the process has computed, and how much longer to completion
 - Resources the process has used
 - Resources the process needs to complete
 - How many processes will need to be terminated?
 - Is the process interactive or batch?

Deadlock Recovery: Resource preemption

- Selecting a victim
 - Minimize cost
- Rollback
 - Return to some **safe state**, restart process from that state
- Starvation
 - Same process may always be picked as victim
 - Include number of rollback in cost factor

Acknowledgements

- “Operating Systems Concepts” book and supplementary material by A. Silberschatz, P. Galvin and G. Gagne
- “Operating Systems: Internals and Design Principles” book and supplementary material by W. Stallings
- “Modern Operating Systems” book and supplementary material by A. Tanenbaum
- R. Doursat and M. Yuksel from University of Nevada, Reno
- Farshad Ghanei from Illinois Tech
- T. Kosar and K. Dantu from University at Buffalo

Announcement

- Reduced homeworks!!
 - Good news for everyone
- Homework 5
 - Will be released next week
- Quiz 6
 - Next Class- Lecture 20 - 22