

System Security - Attack and Defense for Binaries

CS 4390/5390, Spring 2026

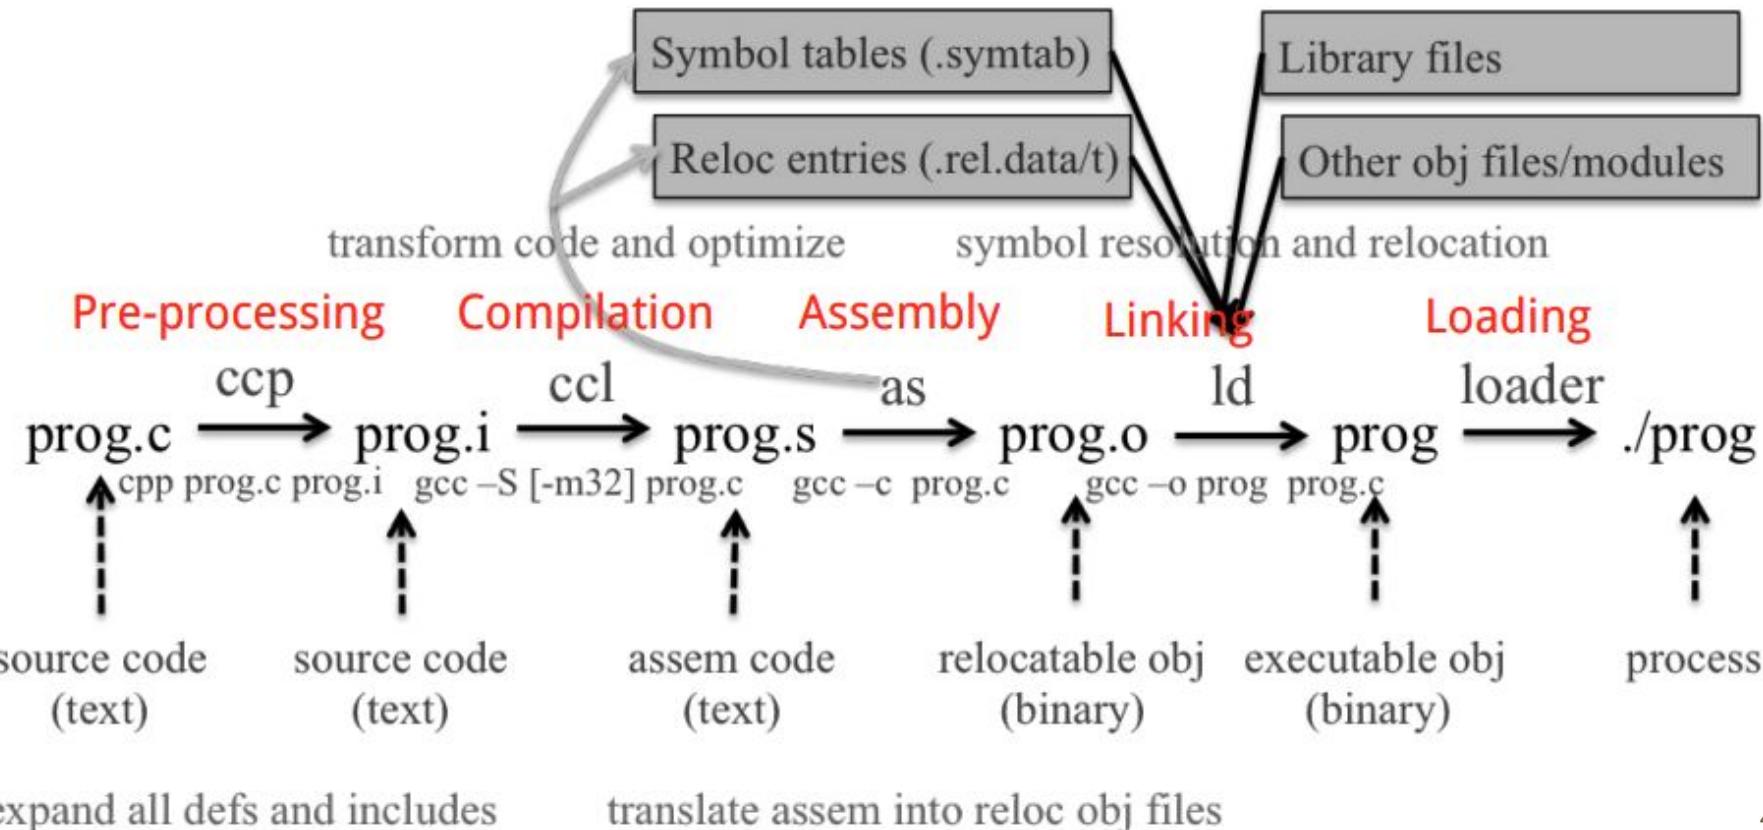
Instructor: MD Armanuzzaman (*Arman*)

Agenda

- Background knowledge
 - Compiler, linker, loader
 - x86 and x86-64 architectures and ISA
 - Linux fundamentals
 - Linux file permissions
 - Set-UID programs
 - Memory map of a Linux process
 - System calls
 - Environment and Shell variables
 - ELF files
 - Reverse engineering tools

Background Knowledge: Compiler, linker, and loader

From a C program to a process



A Shell in a Nutshell

```
int pid = fork();

if (pid == 0) {
    // I am the child process
    exec("ls");
}

else if (pid == -1) {
    // fork failed
}

else {
    // I am the parent; continue my business being a cool program
    // I could wait for the child to finish if I want
}
```

Loading and Executing a Binary Program on Linux

Validation (permissions, memory requirements etc.)

Operating system starts by setting up a new process for the program to run in, including a virtual address space.

The operating system maps an interpreter into the process's virtual memory.

Interpreter, e.g., `/lib/ld-linux.so` in Linux

The interpreter loads the binary into its virtual address space (the same space in which the interpreter is loaded).

It then parses the binary to find out (among other things) which dynamic libraries the binary uses.

The interpreter maps these into the virtual address space (using `mmap` or an equivalent function) and then performs any necessary last-minute relocations in the binary's code sections to fill in the correct addresses for references to the dynamic libraries.

1. Copying the command-line arguments on the stack
2. Initializing registers (e.g., the stack pointer)
3. Jumping to the program entry point (`_start`)

Compiling a C program behind the scene (add_32 add_64)

add.c

```
#include "add.h"
#define BASE 50

int add(int a, int b)
{ return a + b + BASE;}
```

add.h

```
#ifndef ADD_H
#define ADD_H

int add(int, int);

#endif
```

```
gcc -Wall -fno-stack-protector -m32 -O2 add.c main.c -o add_32
```

```
gcc -Wall -fno-stack-protector -m64 -O2 add.c main.c -o add_64
```

main.c

```
/* This program has an integer overflow vulnerability. */
#include "add.h"
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#define USAGE "Add two integers with 50. Usage: add a b\n"

int main(int argc, char *argv[])
{
    int a = 0;
    int b = 0;

    if (argc != 3)
    {
        printf(USAGE);
        return 0; }

    a = atoi(argv[1]);
    b = atoi(argv[2]);
    printf("%d + %d + 50 = %d\n", a, b, add(a, b));
}
```

Background Knowledge: x86 architecture

Data Types

There are 5 integer data types:

Byte – 8 bits.

Word – 16 bits.

Dword, Doubleword – 32 bits.

Quadword – 64 bits.

Double quadword – 128 bits.

Endianness

- Little Endian (Intel, ARM)

Least significant byte has lowest address

Dword address: 0x0

Value: 0x78563412

- Big Endian

Least significant byte has highest address

Dword address: 0x0

Value: 0x12345678

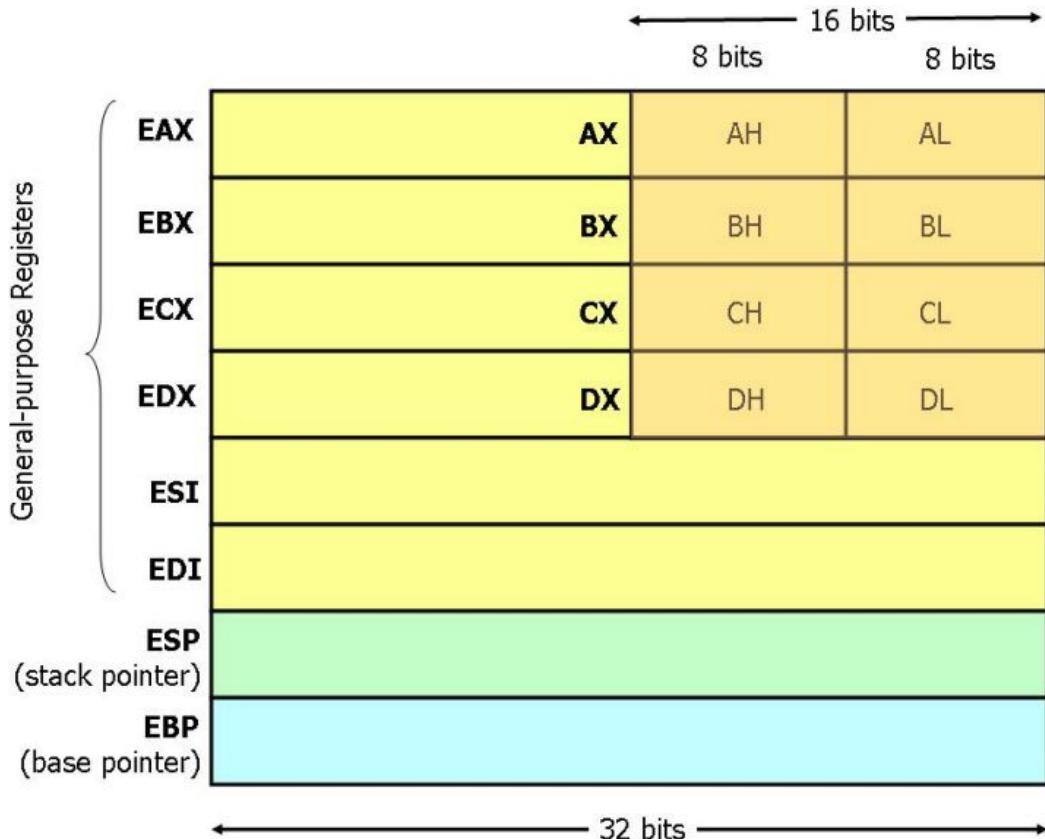
Address 0	0x12
Address 1	0x34
Address 2	0x56
Address 3	0x78

Base Registers

There are

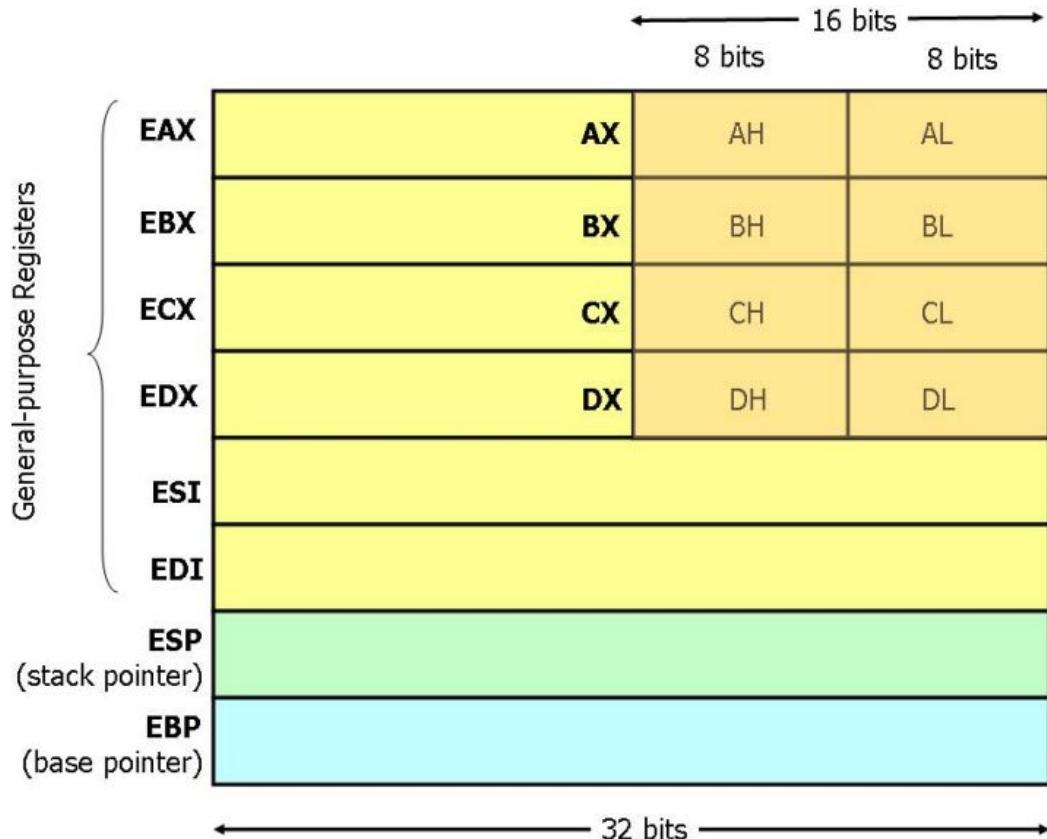
- Eight 32-bit “general-purpose” registers,
- One 32-bit EFLAGS register,
- One 32-bit instruction pointer register (eip), and
- Other special-purpose registers.

The General-Purpose Registers



- 8 general-purpose registers
- **esp** is the stack pointer
- **ebp** is the base pointer
- **esi** and **edi** are source and destination index registers for array and string operations

The General-Purpose Registers



- The registers **eax**, **ebx**, **ecx**, and **edx** may be accessed as 32-bit, 16-bit, or 8-bit registers.
- The other four registers can be accessed as 32-bit or 16-bit.

EFLAGS Register

The various bits of the 32-bit EFLAGS register are set (1) or reset/clear (0) according to the results of certain operations.

We will be interested in, at most, the bits

CF – carry flag

PF – parity flag

ZF – zero flag

SF – sign flag

Instruction Pointer (EIP)

Finally, there is the **EIP** register, which is the instruction pointer (program counter).

Register **EIP** holds the address of the **next** instruction to be executed.

Registers on x86 and amd64

ZMM0	YMM0	XMM0	ZMM1	YMM1	XMM1	ST(0)	MM0	ST(1)	MM1	AL	AH	AX	EAX	RAX	R8B	R8W	R8D	R8	R12B	R12W	R12D	R12	MSW	CR0	CR4
ZMM2	YMM2	XMM2	ZMM3	YMM3	XMM3	ST(2)	MM2	ST(3)	MM3	BL	BH	BX	EBX	RBX	R9B	R9W	R9D	R9	R13B	R13W	R13D	R13	CR1	CR5	
ZMM4	YMM4	XMM4	ZMM5	YMM5	XMM5	ST(4)	MM4	ST(5)	MM5	CL	CH	CX	ECX	RCX	R10B	R10W	R10D	R10	R14B	R14W	R14D	R14	CR2	CR6	
ZMM6	YMM6	XMM6	ZMM7	YMM7	XMM7	ST(6)	MM6	ST(7)	MM7	DL	DH	DX	EDX	RDX	R11B	R11W	R11D	R11	R15B	R15W	R15D	R15	CR3	CR7	
ZMM8	YMM8	XMM8	ZMM9	YMM9	XMM9					BPL	BP	EBP	RBP		DIL	DI	EDI	RDI		IP	EIP	RIP	MXCSR	CR8	
ZMM10	YMM10	XMM10	ZMM11	YMM11	XMM11	CW	FP_IP	FP_DP	FP_CS	SIL	SI	ESI	RSI		SPL	SP	ESP	RSP					CR9		
ZMM12	YMM12	XMM12	ZMM13	YMM13	XMM13	TW																	CR10		
ZMM14	YMM14	XMM14	ZMM15	YMM15	XMM15	FP_DS																	CR11		
ZMM16	ZMM17	ZMM18	ZMM19	ZMM20	ZMM21	ZMM22	ZMM23	FP_OPC	FP_DP	FP_IP	CS	SS	DS		GDTR	IDTR		DR0	DR6			CR12			
ZMM24	ZMM25	ZMM26	ZMM27	ZMM28	ZMM29	ZMM30	ZMM31				ES	FS	GS		TR	LDTR		DR1	DR7			CR13			
														FLAGS	EFLAGS	RFLAGS		DR2	DR8			CR14			
																		DR3	DR9			CR15			
																		DR4	DR10	DR12	DR14				
																		DR5	DR11	DR13	DR15				

Instructions

Each instruction is of the form

label: mnemonic operand1, operand2, operand3

The label is optional.

The number of operands is 0, 1, 2, or 3, depending on the mnemonic .

Each operand is either

- An immediate value,
- A register, or
- A memory address.

Source and Destination Operands

Each operand is either a source operand or a destination operand.

A source operand, in general, may be

- An immediate value,
- A register, or
- A memory address.

A destination operand, in general, may be

- A register, or
- A memory address.

Instructions

hlt - 0 operands

halts the central processing unit (CPU) until the next external interrupt is fired

inc - 1 operand; inc <reg>, inc <mem>

add - 2 operands; add <reg>,<reg>

imul - 1, 2, or 3 operands; imul <reg32>,<reg32>,<con>

In Intel syntax the first operand is the destination

Intel Syntax Assembly and Disassembly

Machine instructions generally fall into three categories: **data movement**, **arithmetic/logic**, and **control-flow**.

<reg32> Any 32-bit register (eax, ebx, ecx, edx, esi, edi, esp, or ebp)

<reg16> Any 16-bit register (ax, bx, cx, or dx)

<reg8> Any 8-bit register (ah, bh, ch, dh, al, bl, cl, or dl)

<reg> Any register

<mem> A memory address (e.g., [eax] or [eax + ebx*4]); [] square brackets

<con32> Any 32-bit immediate

<con16> Any 16-bit immediate

<con8> Any 8-bit immediate

<con> Any 8-, 16-, or 32-bit immediate

Addressing Memory

Move from **source** (operand 2) to **destination** (operand 1)

Square bracket [] represents memory location.

mov [eax], ebx Copy 4 bytes from register EBX into memory address specified in EAX.

mov eax, [esi - 4] Move 4 bytes at memory address ESI - 4 into EAX.

mov [esi + eax * 1], cl Move the contents of CL into the byte at address ESI+EAX*1.

mov edx, [esi + ebx*4] Move the 4 bytes of data at address ESI+4*EBX into EDX.

Addressing Memory

The size directives **BYTE PTR**, **WORD PTR**, and **DWORD PTR** serve this purpose, indicating sizes of 1, 2, and 4 bytes respectively.

mov [ebx], 2 isn't this ambiguous? We can have a default.

mov BYTE PTR [ebx], 2 Move 2 into the single byte at the address stored in EBX.

mov WORD PTR [ebx], 2 Move the 16-bit integer representation of 2 into the 2 bytes starting at the address in EBX.

mov DWORD PTR [ebx], 2 Move the 32-bit integer representation of 2 into the 4 bytes starting at the address in EBX.

Data Movement Instructions

mov – Move

Syntax

mov <reg>, <reg>

mov <reg>, <mem>

mov <mem>, <reg>

mov <reg>, <con>

mov <mem>, <con>

Examples

mov eax, ebx – copy the value in EBX into EAX

mov byte ptr [var], 5 – store the value 5 into the byte at location var

Data Movement Instructions

push – Push on stack; decrements ESP by 4, then places the operand at the location ESP points to.

Syntax

push <reg32>

push <mem>

push <con32>

Examples

push eax – push eax on the stack

push [var] – push the 4 bytes at address var onto the stack

Data Movement Instructions

pop – Pop from stack

Syntax

`pop <reg32>`

`pop <mem>`

Examples

`pop edi` – pop the top element of the stack into EDI.

`push [var]` – pop the top element of the stack into memory at the four bytes starting at location EBX.

LEA Instructions

lea – Load effective address; used for quick calculation

Syntax

```
lea <reg32>, <mem>
```

Examples

Lea edi, [ebx+4*esi] – the quantity EBX+4*ESI is placed in EDI.

Arithmetic and Logic Instructions

add eax, 10 – EAX is set to EAX + 10

addb byte ptr [eax], 10 – add 10 to the single byte stored at memory address stored in EAX

sub al, ah – AL is set to AL - AH

sub eax, 216 – subtract 216 from the value stored in EAX

dec eax – subtract one from the contents of EAX

imul eax, [ebx] – multiply the contents of EAX by the 32-bit contents of the memory at location EBX. Store the result in EAX.

shr ebx, cl – Store in EBX the floor of result of dividing the value of EBX by 2^n where n is the value in CL.

Control Flow Instructions

jmp – Jump

Transfers program control flow to the instruction at the memory location indicated by the operand.

Syntax

`jmp <label>` # direct jump

`jmp <reg32>` # indirect jump

Examples

jmp begin – Jump to the instruction labeled begin.

Control Flow Instructions

jcondition – Conditional jump

Syntax

je <label> (jump when equal)

jne <label> (jump when not equal)

jz <label> (jump when last result was zero)

jg <label> (jump when greater than)

jge <label> (jump when greater than or equal to)

jl <label> (jump when less than)

jle <label> (jump when less than or equal to)

Examples

cmp ebx, eax

jle done

Control Flow Instructions

cmp – Compare

Syntax

cmp <reg>, <reg>

cmp <mem>, <reg>

cmp <reg>, <mem>

cmp <con>, <reg>

Examples

cmp byte ptr [ebx], 10

jeq loop

If the byte stored at the memory location in EBX is equal to the integer constant 10, jump to the location labeled loop.

Control Flow Instructions

call – Subroutine call

The call instruction first **pushes the current code location onto the hardware supported stack** in memory, and then performs an **unconditional jump** to the code location indicated by the label operand. *Unlike the simple jump instructions, the call instruction saves the location to return to when the subroutine completes.*

Syntax

call <label>

call <reg32>

call <mem>

Control Flow Instructions

ret – Subroutine return

The ret instruction implements a subroutine return mechanism. This instruction pops a code location off the hardware supported in-memory stack to the program counter.

Syntax

```
ret
```

The Run-time Stack

The run-time stack supports procedure calls and the passing of parameters between procedures.

The stack is located in memory.

The stack grows towards **low memory**.

When we **push** a value, **esp** is decremented.

When we **pop** a value, **esp** is incremented.

Stack Instructions

enter – Create a function frame

Equivalent to:

push ebp

mov ebp, esp

sub esp, Imm

Stack Instructions

leave – Releases the function frame set up by an earlier ENTER instruction.

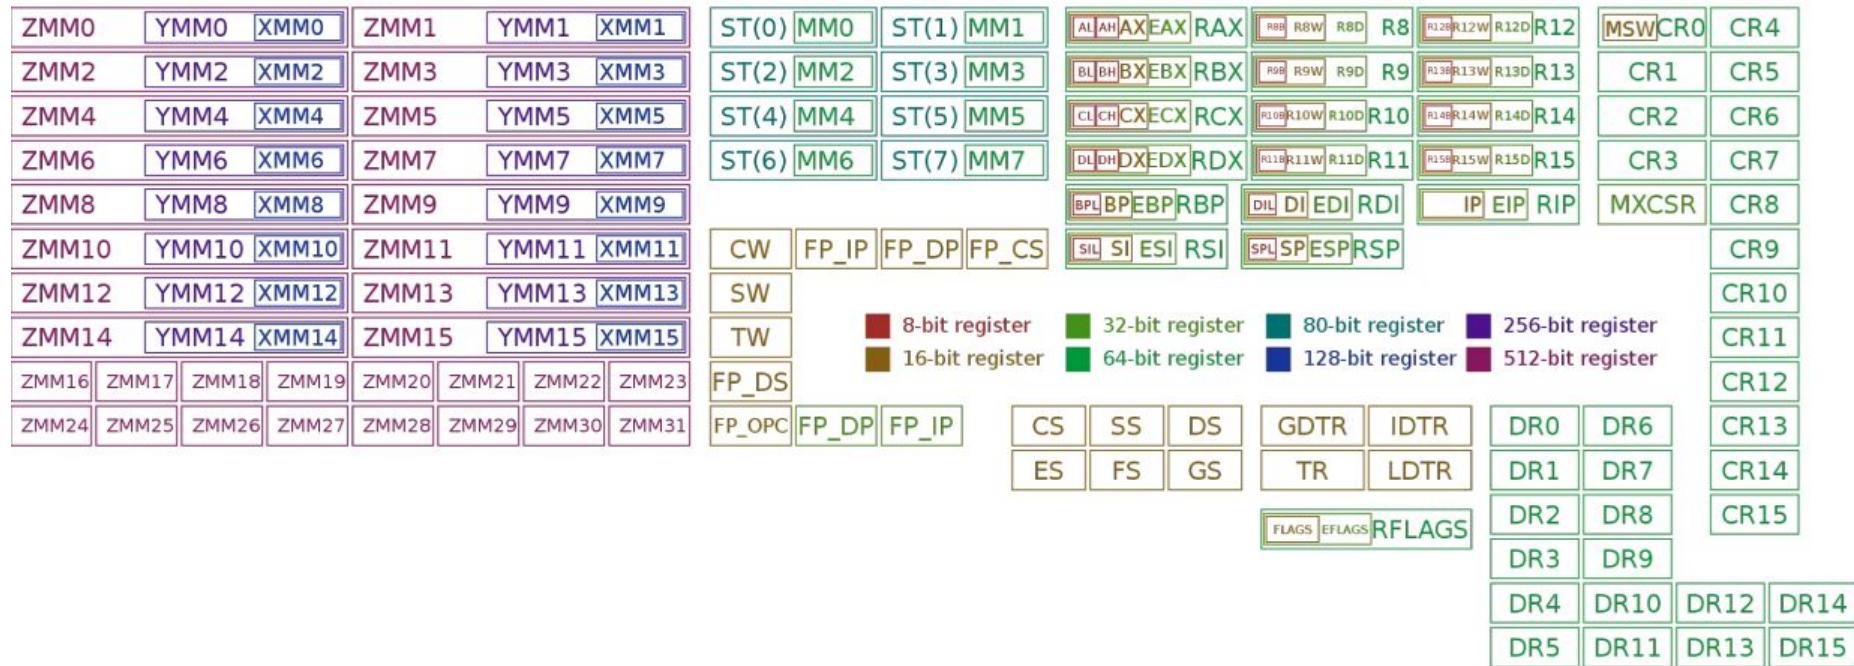
Equivalent to:

mov esp, ebp

pop ebp

Background Knowledge: x86-64/amd64 architecture

Registers on x86 and amd64



x86 vs. x86-64 (code/ladd)

main.c

```
/*
This program has an integer overflow vulnerability.
*/
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
long long ladd(long long *xp, long long y)
{
    long long t = *xp + y;
    return t;
}
```

```
gcc -Wall -m32 -O2 main.c -o ladd
```

```
gcc -Wall -O2 main.c -o ladd64
```

```
int main(int argc, char *argv[])
{
    long long a = 0;
    long long b = 0;

    if (argc != 3)
    {
        printf("Usage: ladd a b\n");
        return 0;
    }
    printf("The sizeof(long long) is %d\n", sizeof(long long));
    a = atoll(argv[1]);
    b = atoll(argv[2]);
    printf("%lld + %lld = %lld\n", a, b, ladd(&a, b));
}
```

x86 vs. x86-64 (code/ladd)

x86

```
000012c0 <ladd>:  
12c0: f3 0f 1e fb    endbr32  
12c4: 8b 44 24 04    mov  0x4(%esp),%eax  
12c8: 8b 50 04        mov  0x4(%eax),%edx  
12cb: 8b 00            mov  (%eax),%eax  
12cd: 03 44 24 08    add   0x8(%esp),%eax  
12d1: 13 54 24 0c    adc   0xc(%esp),%edx  
12d5: c3                ret
```

x86-64

```
00000000001220 <ladd>:  
1220: f3 0f 1e fa endbr64  
1224: 48 8b 07 mov rax,QWORD PTR [rdi]  
1227: 48 01 f0 add rax,rsi  
122a: c3 ret
```

```
objdump -M intel -d ladd_32  
objdump -M intel -d ladd_64
```

Background Knowledge: ARM Cortex-A/M Architecture

Cortex-A 64 bit

X0/W0				
X1/W1				
X2/W2				
X3/W3				
X4/W4				
X5/W5				
X6/W6				
X7/W7				
X8/W8				
X9/W9				
X10/W10				
X11/W11				
X12/W12				
X13/W13				
X14/W14				
X15/W15				
X16/W16				
X17/W17				
X18/W18				
X19/W19				
X20/W20				
X21/W21				
X22/W22				
X23/W23				
X24/W24				
X25/W25				
X26/W26				
X27/W27				
X28/W28				
X29/W29				
X30/W30				
Frame pointer				
Procedure link register				
EL0, EL1, EL2, EL3				

Special registers

Zero register XZR/WZR

Program counter PC

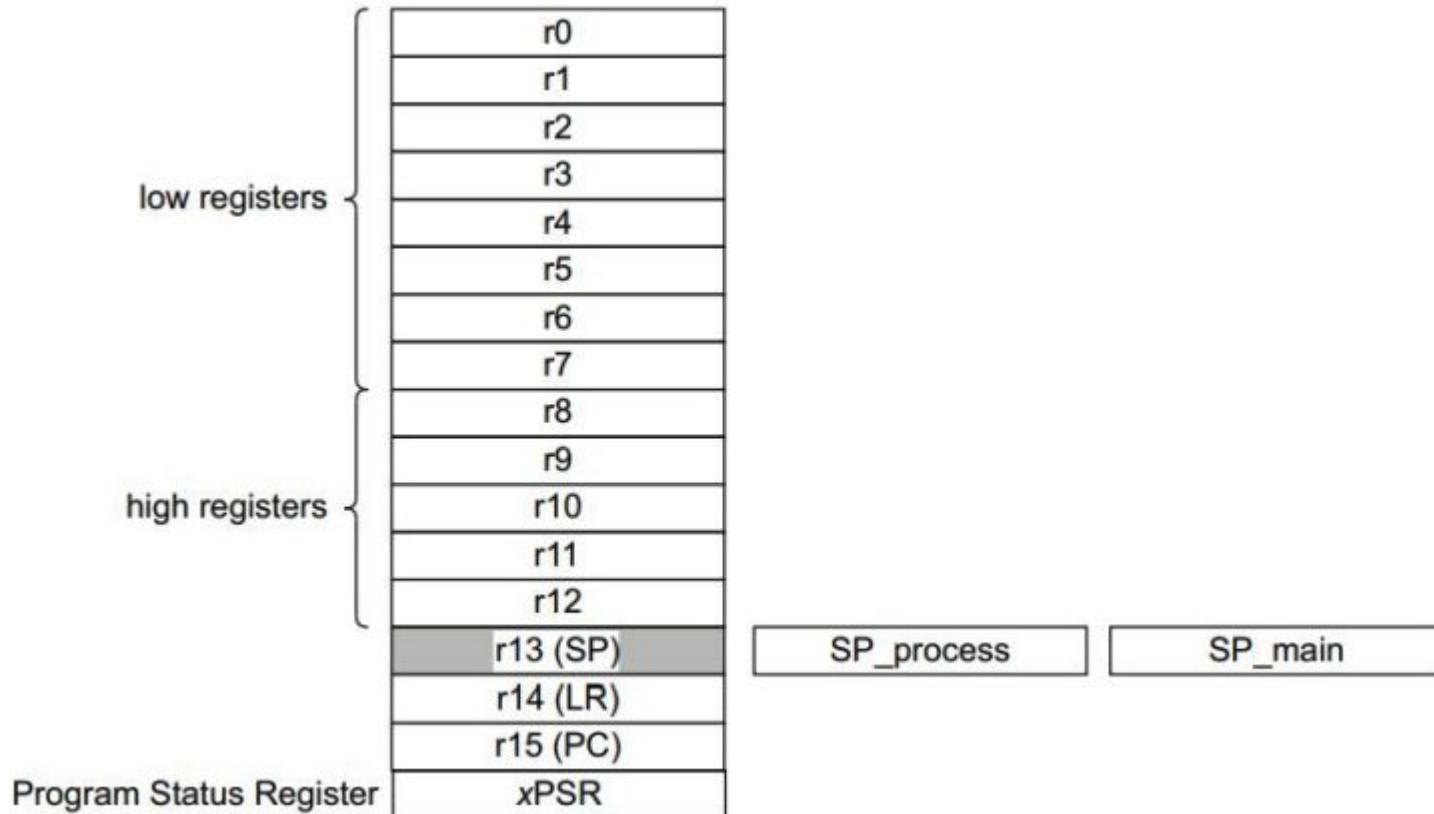
Stack pointer SP_EL0 SP_EL1 SP_EL2 SP_EL3

Program Status Register SPSR_EL1 SPSR_EL2 SPSR_EL3

Exception Link Register ELR_EL1 ELR_EL2 ELR_EL3

EL0 EL1 EL2 EL3

Cortex-M 32 bit



Background Knowledge: Linux File Permissions

Permission Groups

Each file and directory has three user-based permission groups:

Owner – A user is the owner of the file. By default, the person who created a file becomes its owner. The Owner permissions apply only to the owner of the file or directory

Group – A group can contain multiple users. All users belonging to a group will have the same access permissions to the file. The Group permissions apply only to the group that has been assigned to the file or directory

Others – The Others permissions apply to all other users on the system.

Permission Types

Each file or directory has three basic permission types defined for all the 3 user types:

Read – The Read permission refers to a user's capability to read the contents of the file.

Write – The Write permissions refer to a user's capability to write or modify a file or directory.

Execute – The Execute permission affects a user's capability to execute a file or view the contents of a directory.

File type

First field in the output is file type. If there is a **-** it means it is a plain file. If there is **d** it means it is a directory, **c** represents a character device, **b** represents a block device.

```
arman@aserver:~/STC/software-security$ ls -l
total 80
-rwxrwxr-x  1 arman arman 1401 Jan 13 18:09 auth.py
drwxrwxr-x 15 arman arman 4096 Jan 13 18:09 challenges
drwxrwxr-x  3 arman arman 4096 Jan 13 18:09 conf
-rw-rw-r--  1 arman arman  264 Jan 14 02:18 config.env
drwxrwxr-x  7 arman arman 4096 Jan 13 18:09 containers
drwxrwxr-x  7 arman arman 4096 Jan 13 18:09 CTFd
drwxrwxr-x  3 arman arman 4096 Jan 13 18:09 CTFd_plugin
-rw-rw-r--  1 arman arman 2526 Jan 17 23:24 docker-compose.yml
-rwxrwxr-x  1 arman arman  771 Jan 13 18:09 enter.py
drwxrwxr-x  2 arman arman 4096 Jan 13 18:09 home_daemon
-rw-rw-r--  1 arman arman 3074 Jan 13 18:10 README.md
-rwxrwxr-x  1 arman arman 3003 Jan 13 18:31 resetup.sh
-rwxrwxr-x  1 arman arman 1373 Jan 13 18:09 restart.sh
-rwxrwxr-x  1 arman arman 1211 Jan 14 02:07 run.sh
drwxrwxr-x  2 arman arman 4096 Jan 13 18:09 scripts
-rw-rw-r--  1 arman arman 1986 Jan 13 18:09 script.sh
-rwxrwxr-x  1 arman arman 2996 Jan 14 01:40 setup.sh
-rw-rw-r--  1 arman arman 7196 Jan 13 18:09 tips.md
-rw-rw-r--  1 arman arman     0 Jan 13 18:09 todo-list.md
-rwxrwxr-x  1 arman arman  830 Jan 13 18:09 update.sh
```

Permissions for owner, group, and others

```
arman@aserver:~/STC/software-security$ ls -l
total 80
-rwxrwxr-x 1 arman arman 1401 Jan 13 18:09 auth.py
drwxrwxr-x 15 arman arman 4096 Jan 13 18:09 challenges
drwxrwxr-x 3 arman arman 4096 Jan 13 18:09 conf
-rw-rw-r-- 1 arman arman 264 Jan 14 02:18 config.env
drwxrwxr-x 7 arman arman 4096 Jan 13 18:09 containers
drwxrwxr-x 7 arman arman 4096 Jan 13 18:09 CTFd
drwxrwxr-x 3 arman arman 4096 Jan 13 18:09 CTFd_plugin
-rw-rw-r-- 1 arman arman 2526 Jan 17 23:24 docker-compose.yml
-rwxrwxr-x 1 arman arman 771 Jan 13 18:09 enter.py
drwxrwxr-x 2 arman arman 4096 Jan 13 18:09 home_daemon
-rw-rw-r-- 1 arman arman 3074 Jan 13 18:10 README.md
-rwxrwxr-x 1 arman arman 3003 Jan 13 18:31 resetup.sh
-rwxrwxr-x 1 arman arman 1373 Jan 13 18:09 restart.sh
-rwxrwxr-x 1 arman arman 1211 Jan 14 02:07 run.sh
drwxrwxr-x 2 arman arman 4096 Jan 13 18:09 scripts
-rw-rw-r-- 1 arman arman 1986 Jan 13 18:09 script.sh
-rwxrwxr-x 1 arman arman 2996 Jan 14 01:40 setup.sh
-rw-rw-r-- 1 arman arman 7196 Jan 13 18:09 tips.md
-rw-rw-r-- 1 arman arman 0 Jan 13 18:09 todo-list.md
-rwxrwxr-x 1 arman arman 830 Jan 13 18:09 update.sh
```

Link count

```
arman@aserver:~/STC/software-security$ ls -l
total 80
-rwxrwxr-x 1 arman arman 1401 Jan 13 18:09 auth.py
drwxrwxr-x 15 arman arman 4096 Jan 13 18:09 challenges
drwxrwxr-x 3 arman arman 4096 Jan 13 18:09 conf
-rw-rw-r-- 1 arman arman 264 Jan 14 02:18 config.env
drwxrwxr-x 7 arman arman 4096 Jan 13 18:09 containers
drwxrwxr-x 7 arman arman 4096 Jan 13 18:09 CTFd
drwxrwxr-x 3 arman arman 4096 Jan 13 18:09 CTFd_plugin
-rw-rw-r-- 1 arman arman 2526 Jan 17 23:24 docker-compose.yml
-rwxrwxr-x 1 arman arman 771 Jan 13 18:09 enter.py
drwxrwxr-x 2 arman arman 4096 Jan 13 18:09 home_daemon
-rw-rw-r-- 1 arman arman 3074 Jan 13 18:10 README.md
-rwxrwxr-x 1 arman arman 3003 Jan 13 18:31 resetup.sh
-rwxrwxr-x 1 arman arman 1373 Jan 13 18:09 restart.sh
-rwxrwxr-x 1 arman arman 1211 Jan 14 02:07 run.sh
drwxrwxr-x 2 arman arman 4096 Jan 13 18:09 scripts
-rw-rw-r-- 1 arman arman 1986 Jan 13 18:09 script.sh
-rwxrwxr-x 1 arman arman 2996 Jan 14 01:40 setup.sh
-rw-rw-r-- 1 arman arman 7196 Jan 13 18:09 tips.md
-rw-rw-r-- 1 arman arman 0 Jan 13 18:09 todo-list.md
-rwxrwxr-x 1 arman arman 830 Jan 13 18:09 update.sh
```

Owner

This field provide info about the creator of the file.

```
arman@aserver:~/STC/software-security$ ls -l
total 80
-rwxrwxr-x 1 arman arman 1401 Jan 13 18:09 auth.py
drwxrwxr-x 15 arman arman 4096 Jan 13 18:09 challenges
drwxrwxr-x 3 arman arman 4096 Jan 13 18:09 conf
-rw-rw-r-- 1 arman arman 264 Jan 14 02:18 config.env
drwxrwxr-x 7 arman arman 4096 Jan 13 18:09 containers
drwxrwxr-x 7 arman arman 4096 Jan 13 18:09 CTFd
drwxrwxr-x 3 arman arman 4096 Jan 13 18:09 CTFd_plugin
-rw-rw-r-- 1 arman arman 2526 Jan 17 23:24 docker-compose.yml
-rwxrwxr-x 1 arman arman 771 Jan 13 18:09 enter.py
drwxrwxr-x 2 arman arman 4096 Jan 13 18:09 home_daemon
-rw-rw-r-- 1 arman arman 3074 Jan 13 18:10 README.md
-rwxrwxr-x 1 arman arman 3003 Jan 13 18:31 resetup.sh
-rwxrwxr-x 1 arman arman 1373 Jan 13 18:09 restart.sh
-rwxrwxr-x 1 arman arman 1211 Jan 14 02:07 run.sh
drwxrwxr-x 2 arman arman 4096 Jan 13 18:09 scripts
-rw-rw-r-- 1 arman arman 1986 Jan 13 18:09 script.sh
-rwxrwxr-x 1 arman arman 2996 Jan 14 01:40 setup.sh
-rw-rw-r-- 1 arman arman 7196 Jan 13 18:09 tips.md
-rw-rw-r-- 1 arman arman 0 Jan 13 18:09 todo-list.md
-rwxrwxr-x 1 arman arman 830 Jan 13 18:09 update.sh
```

Group

```
arman@aserver:~/STC/software-security$ ls -l
total 80
-rwxrwxr-x 1 arman arman 1401 Jan 13 18:09 auth.py
drwxrwxr-x 15 arman arman 4096 Jan 13 18:09 challenges
drwxrwxr-x 3 arman arman 4096 Jan 13 18:09 conf
-rw-rw-r-- 1 arman arman 264 Jan 14 02:18 config.env
drwxrwxr-x 7 arman arman 4096 Jan 13 18:09 containers
drwxrwxr-x 7 arman arman 4096 Jan 13 18:09 CTFd
drwxrwxr-x 3 arman arman 4096 Jan 13 18:09 CTFd_plugin
-rw-rw-r-- 1 arman arman 2526 Jan 17 23:24 docker-compose.yml
-rwxrwxr-x 1 arman arman 771 Jan 13 18:09 enter.py
drwxrwxr-x 2 arman arman 4096 Jan 13 18:09 home_daemon
-rw-rw-r-- 1 arman arman 3074 Jan 13 18:10 README.md
-rwxrwxr-x 1 arman arman 3003 Jan 13 18:31 resetup.sh
-rwxrwxr-x 1 arman arman 1373 Jan 13 18:09 restart.sh
-rwxrwxr-x 1 arman arman 1211 Jan 14 02:07 run.sh
drwxrwxr-x 2 arman arman 4096 Jan 13 18:09 scripts
-rw-rw-r-- 1 arman arman 1986 Jan 13 18:09 script.sh
-rwxrwxr-x 1 arman arman 2996 Jan 14 01:40 setup.sh
-rw-rw-r-- 1 arman arman 7196 Jan 13 18:09 tips.md
-rw-rw-r-- 1 arman arman 0 Jan 13 18:09 todo-list.md
-rwxrwxr-x 1 arman arman 830 Jan 13 18:09 update.sh
```

File size

```
arman@aserver:~/STC/software-security$ ls -l
total 80
-rwxrwxr-x  1 arman arman 1401 Jan 13 18:09 auth.py
drwxrwxr-x 15 arman arman 4096 Jan 13 18:09 challenges
drwxrwxr-x  3 arman arman 4096 Jan 13 18:09 conf
-rw-rw-r--  1 arman arman  264 Jan 14 02:18 config.env
drwxrwxr-x  7 arman arman 4096 Jan 13 18:09 containers
drwxrwxr-x  7 arman arman 4096 Jan 13 18:09 CTFd
drwxrwxr-x  3 arman arman 4096 Jan 13 18:09 CTFd_plugin
-rw-rw-r--  1 arman arman 2526 Jan 17 23:24 docker-compose.yml
-rwxrwxr-x  1 arman arman  771 Jan 13 18:09 enter.py
drwxrwxr-x  2 arman arman 4096 Jan 13 18:09 home_daemon
-rw-rw-r--  1 arman arman 3074 Jan 13 18:10 README.md
-rwxrwxr-x  1 arman arman 3003 Jan 13 18:31 resetup.sh
-rwxrwxr-x  1 arman arman 1373 Jan 13 18:09 restart.sh
-rwxrwxr-x  1 arman arman 1211 Jan 14 02:07 run.sh
drwxrwxr-x  2 arman arman 4096 Jan 13 18:09 scripts
-rw-rw-r--  1 arman arman 1986 Jan 13 18:09 script.sh
-rwxrwxr-x  1 arman arman 2996 Jan 14 01:40 setup.sh
-rw-rw-r--  1 arman arman 7196 Jan 13 18:09 tips.md
-rw-rw-r--  1 arman arman     0 Jan 13 18:09 todo-list.md
-rwxrwxr-x  1 arman arman  830 Jan 13 18:09 update.sh
```

Last modify time

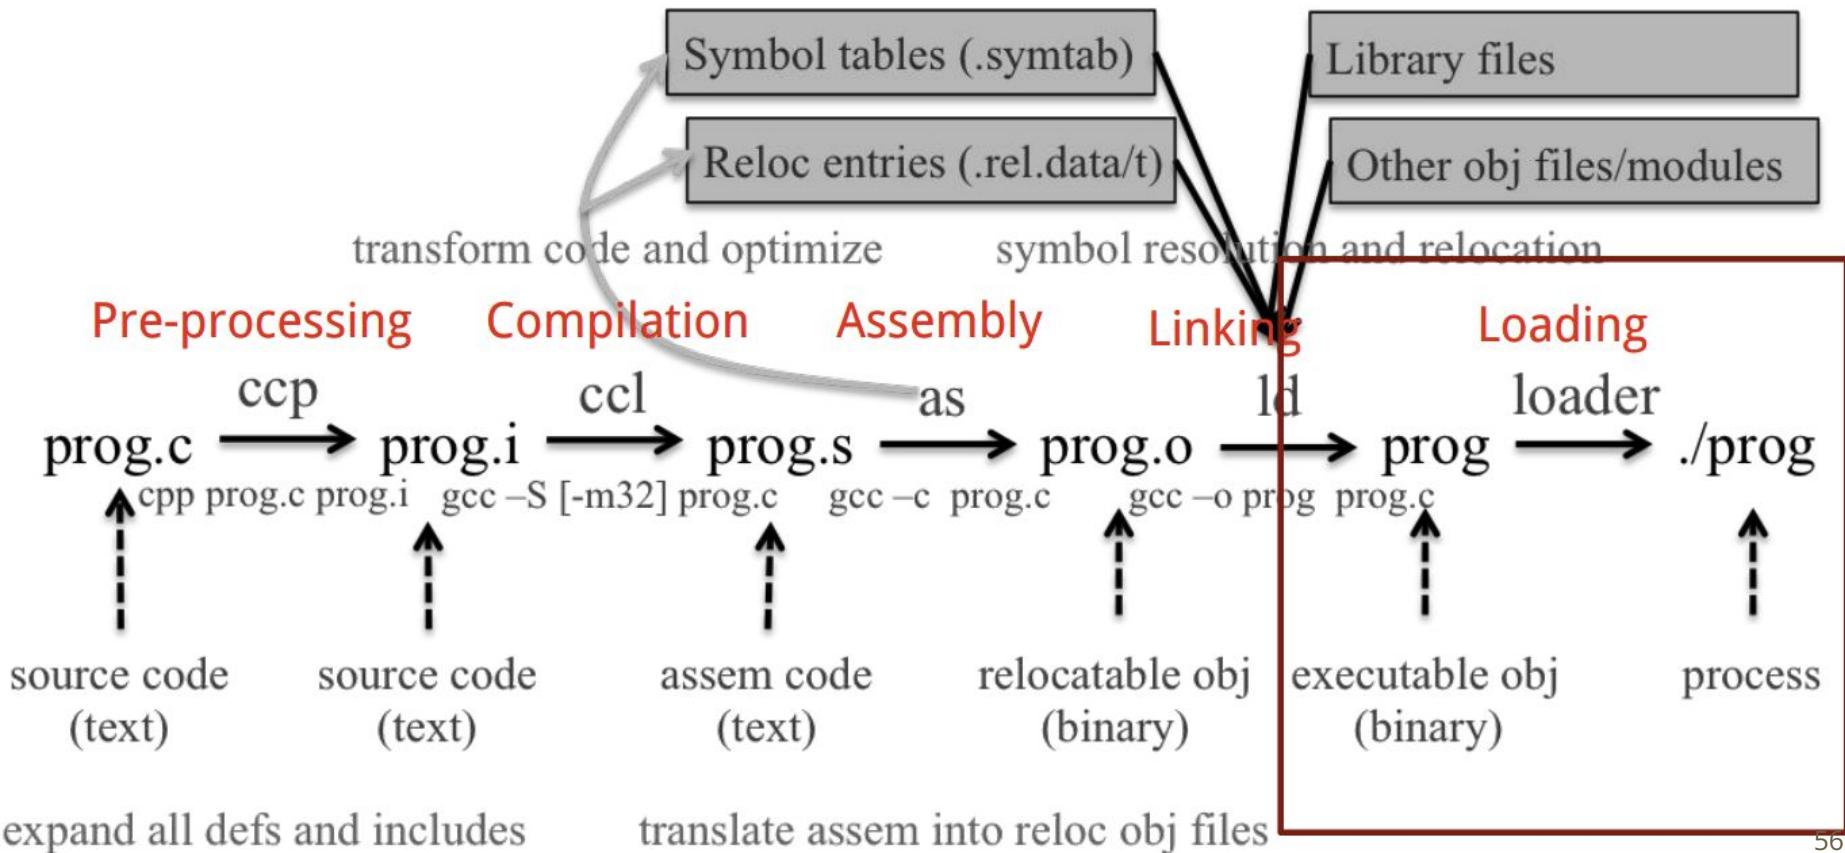
```
arman@aserver:~/STC/software-security$ ls -l
total 80
-rwxrwxr-x  1 arman arman 1401 Jan 13 18:09 auth.py
drwxrwxr-x 15 arman arman 4096 Jan 13 18:09 challenges
drwxrwxr-x  3 arman arman 4096 Jan 13 18:09 conf
-rw-rw-r--  1 arman arman   264 Jan 14  02:18 config.env
drwxrwxr-x  7 arman arman 4096 Jan 13 18:09 containers
drwxrwxr-x  7 arman arman 4096 Jan 13 18:09 CTFd
drwxrwxr-x  3 arman arman 4096 Jan 13 18:09 CTFd_plugin
-rw-rw-r--  1 arman arman 2526 Jan 17 23:24 docker-compose.yml
-rwxrwxr-x  1 arman arman   771 Jan 13 18:09 enter.py
drwxrwxr-x  2 arman arman 4096 Jan 13 18:09 home_daemon
-rw-rw-r--  1 arman arman 3074 Jan 13 18:10 README.md
-rwxrwxr-x  1 arman arman 3003 Jan 13 18:31 resetup.sh
-rwxrwxr-x  1 arman arman 1373 Jan 13 18:09 restart.sh
-rwxrwxr-x  1 arman arman 1211 Jan 14  02:07 run.sh
drwxrwxr-x  2 arman arman 4096 Jan 13 18:09 scripts
-rw-rw-r--  1 arman arman 1986 Jan 13 18:09 script.sh
-rwxrwxr-x  1 arman arman 2996 Jan 14  01:40 setup.sh
-rw-rw-r--  1 arman arman 7196 Jan 13 18:09 tips.md
-rw-rw-r--  1 arman arman     0 Jan 13 18:09 todo-list.md
-rwxrwxr-x  1 arman arman  830 Jan 13 18:09 update.sh
```

filename

```
arman@aserver:~/STC/software-security$ ls -l
total 80
-rwxrwxr-x  1 arman arman 1401 Jan 13 18:09 auth.py
drwxrwxr-x 15 arman arman 4096 Jan 13 18:09 challenges
drwxrwxr-x  3 arman arman 4096 Jan 13 18:09 conf
-rw-rw-r--  1 arman arman  264 Jan 14 02:18 config.env
drwxrwxr-x  7 arman arman 4096 Jan 13 18:09 containers
drwxrwxr-x  7 arman arman 4096 Jan 13 18:09 CTFd
drwxrwxr-x  3 arman arman 4096 Jan 13 18:09 CTFd_plugin
drwxrwxr-x  1 arman arman 2526 Jan 17 23:24 docker-compose.yml
-rwxrwxr-x  1 arman arman  771 Jan 13 18:09 enter.py
drwxrwxr-x  2 arman arman 4096 Jan 13 18:09 home_daemon
-rw-rw-r--  1 arman arman 3074 Jan 13 18:10 README.md
-rwxrwxr-x  1 arman arman 3003 Jan 13 18:31 resetup.sh
-rwxrwxr-x  1 arman arman 1373 Jan 13 18:09 restart.sh
-rwxrwxr-x  1 arman arman 1211 Jan 14 02:07 run.sh
drwxrwxr-x  2 arman arman 4096 Jan 13 18:09 scripts
-rw-rw-r--  1 arman arman 1986 Jan 13 18:09 script.sh
-rwxrwxr-x  1 arman arman 2996 Jan 14 01:40 setup.sh
-rw-rw-r--  1 arman arman  7196 Jan 13 18:09 tips.md
-rw-rw-r--  1 arman arman     0 Jan 13 18:09 todo-list.md
-rwxrwxr-x  1 arman arman   830 Jan 13 18:09 update.sh
```

Background Knowledge: Set-UID Programs

From a C program to a process



Real UID, Effective UID, and Saved UID

Each Linux/Unix process has 3 UIDs associated with it.

Real UID (RUID): This is the UID of the user/process that created THIS process. It can be changed only if the running process has EUID=0.

Effective UID (EUID): This UID is used to evaluate privileges of the process to perform a particular action. EUID can be changed either to RUID, or SUID if EUID!=0. If EUID=0, it can be changed to anything.

Saved UID (SUID): If the binary image file, that was launched has a Set-UID bit on, SUID will be the UID of the owner of the file. Otherwise, SUID will be the RUID.

Set-UID Program

The kernel makes the decision whether a process has the privilege by looking on the **EUID** of the process.

For non Set-UID programs, the effective uid and the real uid are the same. For Set-UID programs, **the effective uid is the owner of the program**, while the real uid is the user of the program.

What will happen is when a setuid binary executes, **the process changes its Effective User ID (EUID) from the default RUID to the owner of this special binary executable file which in this case is - root**.

```
arman@aserver:~/STC/software-security$ ls -al --color=always /usr/bin/ | head -n 100
total 407872
drwxr-xr-x 2 root root      36864 Jan 17 23:11 .
drwxr-xr-x 14 root root     4096 Jan 14 22:41 ..
-rw xr-xr-x 1 root root    55744 Apr  5 2024 [REDACTED]
-rw xr-xr-x 1 root root     94 Nov 12 12:15 2to3
-rw xr-xr-x 1 root root   18744 Mar 19 2025 aa-enabled
-rw xr-xr-x 1 root root   18744 Mar 19 2025 aa-exec
-rw xr-xr-x 1 root root   18736 Mar 19 2025 aa-features-abi
-rw xr-xr-x 1 root root    1622 Nov 18 11:26 acpidbg
-rw xr-xr-x 1 root root   16422 Feb 18 2025 add-apt-repository
-rw xr-xr-x 1 root root   14720 Jun  5 2025 addpart
lrwxrwxrwx 1 root root      26 Dec  3 15:01 addr2line -> x86_64-linux-gnu-addr2line
-rw xr-xr-x 1 root root   2322 Apr 18 2024 apport-bug
-rw xr-xr-x 1 root root   13625 Jul  8 2025 apport-cli
lrwxrwxrwx 1 root root      10 Jul  8 2025 apport-collect -> apport-bug
-rw xr-xr-x 1 root root   3790 Jul  8 2025 apport-unpack
-rw xr-xr-x 1 root root   141544 Apr  8 2024 appstreamcli
lrwxrwxrwx 1 root root      6 Aug  5 17:14 apropos -> whatis
-rw xr-xr-x 1 root root   18824 Oct 22 2024 apt
lrwxrwxrwx 1 root root      18 Feb 18 2025 apt-add-repository -> add-apt-repository
-rw xr-xr-x 1 root root   88544 Oct 22 2024 apt-cache
-rw xr-xr-x 1 root root   27104 Oct 22 2024 apt-cdrom
-rw xr-xr-x 1 root root   31120 Oct 22 2024 apt-config
```

-rwxr-xr-x	1	root	root	59912	Apr	5	2024	chcon	
-rwsr-xr-x	1	root	root	72792	May	30	2024	chfn	
-rwxr-xr-x	1	root	root	59912	Apr	5	2024	chgrp	
-rwxr-xr-x	1	root	root	55816	Apr	5	2024	chmod	
-rwxr-xr-x	1	root	root	22912	Jun	5	2025	choom	
-rwxr-xr-x	1	root	root	59912	Apr	5	2024	chown	
-rwxr-xr-x	1	root	root	31104	Jun	5	2025	chrt	
-rwsr-xr-x	1	root	root	44760	May	30	2024	chsh	
-rwxr-xr-x	1	root	root	14712	Mar	31	2024	chvt	
-rwxr-xr-x	1	root	root	27080	Aug	5	17:14	cifsiostat	
-rwxr-xr-x	1	root	root	150674	Feb	26	2024	ckbcomp	
-rwxr-xr-x	1	root	root		227	Aug	5	17:14	ckeygen3
-rwxr-xr-x	1	root	root	104984	Apr	5	2024	cksum	
-rwxr-xr-x	1	root	root	14656	Apr	8	2024	clear	
-rwxr-xr-x	1	root	root	14568	Mar	31	2024	clear_console	
-rwxr-xr-x	1	root	root		972	Jun	24	2025	cloud-id
-rwxr-xr-x	1	root	root		976	Jun	24	2025	cloud-init
-rwxr-xr-x	1	root	root	2108	Jun	24	2025	cloud-init-per	
-rwxr-xr-x	1	root	root	43408	Apr	8	2024	cmp	
-rwxr-xr-x	1	root	root	14640	Mar	31	2024	codepage	
-rwxr-xr-x	1	root	root	22920	Aug	5	17:14	col	
-rwxr-xr-x	1	root	root		963	Aug	5	17:14	col1
lrwxrwxrwx	1	root	root		4	Aug	5	17:14	col2 -> col1
lrwxrwxrwx	1	root	root		4	Aug	5	17:14	col3 -> col1

-rwxr-xr-x	1	root	root	80408	Apr	5	2024	stty
-rwsr-xr-x	1	root	root	55680	Jun	5	2025	su
-rwsr-xr-x	1	root	root	277936	Jun	25	2025	sudo
lrwxrwxrwx	1	root	root	4	Jun	25	2025	sudoedit -> sudo
-rwxr-xr-x	1	root	root	98256	Jun	25	2025	sudoreplay
-rwxr-xr-x	1	root	root	35240	Apr	5	2024	sum
-rwxr-xr-x	1	root	root	35240	Apr	5	2024	sync
-rwxr-xr-x	1	root	root	1501304	Jul	2	2025	systemctl
lrwxrwxrwx	1	root	root	22	Jul	2	2025	systemd -> ./lib/systemd/systemd
-rwxr-xr-x	1	root	root	14792	Jul	2	2025	systemd-ac-power
-rwxr-xr-x	1	root	root	203624	Jul	2	2025	systemd-analyze
-rwxr-xr-x	1	root	root	19024	Jul	2	2025	systemd-ask-password
-rwxr-xr-x	1	root	root	18896	Jul	2	2025	systemd-cat
-rwxr-xr-x	1	root	root	23112	Jul	2	2025	systemd-cgls
-rwxr-xr-x	1	root	root	39392	Jul	2	2025	systemd-cgtop
lrwxrwxrwx	1	root	root	14	Jul	2	2025	systemd-confext -> systemd-sysext
-rwxr-xr-x	1	root	root	43744	Jul	2	2025	systemd-creds
-rwxr-xr-x	1	root	root	72624	Jul	2	2025	systemd-cryptenroll
-rwxr-xr-x	1	root	root	80840	Jul	2	2025	systemd-cryptsetup
-rwxr-xr-x	1	root	root	27080	Jul	2	2025	systemd-delta
-rwxr-xr-x	1	root	root	18888	Jul	2	2025	systemd-detect-virt
-rwxr-xr-x	1	root	root	22984	Jul	2	2025	systemd-escape
-rwxr-xr-x	1	root	root	60232	Jul	2	2025	systemd-firstboot
-rwxr-xr-x	1	root	root	158456	Jul	2	2025	systemd-hwdb

Example: rdsecret

main.c

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <pwd.h>
int main(int argc, char *argv[])
{
    FILE *fp = NULL;
    char buffer[100] = {0};
    // get ruid and euid
    uid_t uid = getuid();
    struct passwd *pw = getpwuid(uid);
    if (pw)
    {
        printf("UID: %d, USER: %s.\n", uid, pw->pw_name);
    }
    uid_t euid = geteuid();
    pw = getpwuid(euid);
```

```
if (pw)
{
    printf("EUID: %d, EUSER: %s.\n", euid, pw->pw_name);
}
print_flag();

return(0);
}

void print_flag()
{
    FILE *fp;
    char buff[MAX_FLAG_SIZE];
    fp = fopen("flag","r");
    fread(buff, MAX_FLAG_SIZE, 1, fp);
    printf("flag is : %s\n", buff);
    fclose(fp);
}
```

Background Knowledge: ELF Binary Files

ELF Files

The **Executable and Linkable Format (ELF)** is a common standard file format for executable files, object code, shared libraries, and core dumps. Filename extension none, .axf, .bin, .elf, .o, .prx, .puff, .ko, .mod and .so

Contains the program and its data. Describes how the program should be loaded (program/segment headers). Contains metadata describing program components (section headers).

Command *file*

```
arman@aserver:~$ file /bin/ls
/bin/ls: ELF 64-bit LSB pie executable, x86-64, version 1 (SYSV), dynamically linked, interpreter /lib64/ld-linux-x86-
64.so.2, BuildID[sha1]=3eca7e3905b37d48cf0a88b576faa7b95cc3097b, for GNU/Linux 3.2.0, stripped
```

file /bin/ls

```
arman@aserver:~$ readelf -a /bin/ls
ELF Header:
  Magic: 7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class: ELF64
  Data: 2's complement, little endian
  Version: 1 (current)
  OS/ABI: UNIX - System V
  ABI Version: 0
  Type: DYN (Position-Independent Executable file)
  Machine: Advanced Micro Devices X86-64
  Version: 0x1
  Entry point address: 0x6d30
  Start of program headers: 64 (bytes into file)
  Start of section headers: 140328 (bytes into file)
  Flags: 0x0
  Size of this header: 64 (bytes)
  Size of program headers: 56 (bytes)
  Number of program headers: 13
  Size of section headers: 64 (bytes)
  Number of section headers: 31
  Section header string table index: 30

Section Headers:
[Nr] Name           Type            Address          Offset
    Size           EntSize        Flags Link Info Align
[ 0] NULL           PROGBITS        0000000000000000 0000000000000000
[ 1] .interp        PROGBITS        0000000000000318 0000000000000318
  0000000000000001c 0000000000000000 A   0   0   1
[ 2] .note.gnu.pr [...] NOTE          0000000000000338 0000000000000338
  0000000000000030 0000000000000000 A   0   0   8
[ 3] .note.gnu.bu [...] NOTE          0000000000000368 0000000000000368
  0000000000000024 0000000000000000 A   0   0   4
[ 4] .note.ABI-tag NOTE          000000000000038c 000000000000038c
  0000000000000020 0000000000000000 A   0   0   4
[ 5] .gnu.hash      GNU_HASH        00000000000003b0 00000000000003b0
  0000000000000050 0000000000000000 A   6   0   8
[ 6] .dynsym        DYNSYM         00000000000000400 00000000000000400
  00000000000000c00 0000000000000018 A   7   1   8
[ 7] .dynstr        STRTAB         0000000000001000 0000000000001000
  00000000000000614 0000000000000000 A   0   0   1
```

INTERP: defines the library that should be used to load this ELF into memory.

LOAD: defines a part of the file that should be loaded into memory.

Sections:

.text: the executable code of your program.

.plt and .got: used to resolve and dispatch library calls.

.data: used for pre-initialized global writable data (such as global arrays with initial values)

.rodata: used for global read-only data (such as string constants)

.bss: used for uninitialized global writable data (such as global arrays without initial values)

Tools for ELF

gcc to make your ELF.

readelf to parse the ELF header.

objdump to parse the ELF header and disassemble the source code.

nm to view your ELF's symbols.

patchelf to change some ELF properties.

objcopy to swap out ELF sections.

strip to remove otherwise-helpful information (such as symbols).

kaitai struct (<https://ide.kaitai.io/>) to look through your ELF interactively.

Background Knowledge: Memory Map of a Linux Process

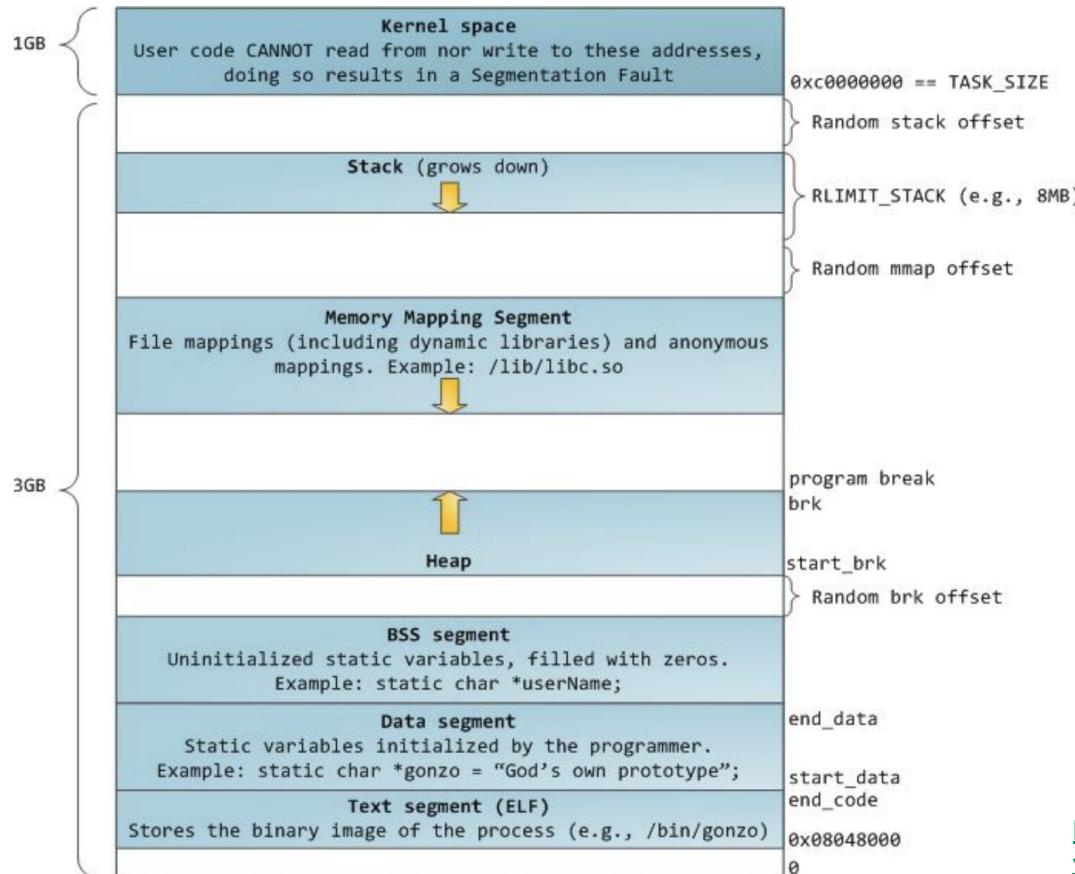
Memory Map of Linux Process (32 bit)

Each process in a multi-tasking OS runs in its own memory sandbox.

This sandbox is the **virtual address space**, which in 32-bit mode is **always a 4GB block of memory addresses**.

These virtual addresses are mapped to physical memory by **page tables**, which are maintained by the operating system kernel and consulted by the processor

Memory Map of Linux Process (32 bit)



NULL Pointer in C/C++

```
int * pInt = NULL;
```

In possible definitions of NULL in C/C++:

```
#define NULL ((char *)0)
```

```
#define NULL 0
```

```
//since C++11
```

```
#define NULL nullptr
```

/proc/pid_of_process/maps

Example processmap.c

```
#include <stdio.h>
#include <stdlib.h>
int main()
{
    getchar();
    return 0;
}
```

cat /proc/pid/maps

pmap -X pid

pmap -X `pidof pm`

/proc/pid_of_process/maps

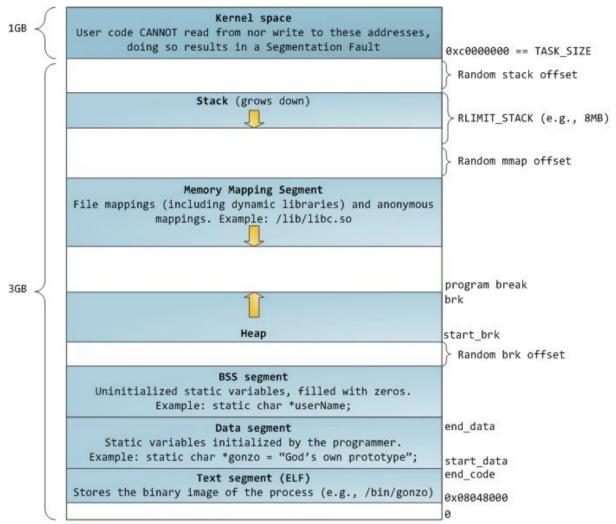
Example processmap.c

```
#include <stdio.h>
#include <stdlib.h>
int main()
{
    getchar();
    return 0;
}
```

cat /proc/pid/maps

pmap -X pid

pmap -X `pidof pm`



```
arman@aserver:~$ pmap -X 746491
746491: ./processmap_32
```

Address	Perm	Offset	Device	Inode	Size	Rss	Pss	Pss_Dirty	Referenced	Anonymous	KSM	LazyFree	ShmemPmdMapped	FilePmdMapped	Shared_Hugetlb	Private_Hugetlb	Swap	SwapPss	Locked	THPEligible	ProtectionKey	Mapping
5d978000	r--p	00000000	fc:00	3145956	4	4	4	0	4	0	0	0	0	0	0	0	0	0	0	0	0	0 processmap_32
5d979000	r-xp	00001000	fc:00	3145956	4	4	4	0	4	0	0	0	0	0	0	0	0	0	0	0	0	0 processmap_32
5d97a000	r--p	00002000	fc:00	3145956	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0 processmap_32
5d97b000	r--p	00002000	fc:00	3145956	4	4	4	4	4	4	0	0	0	0	0	0	0	0	0	0	0	0 processmap_32
5d97c000	rw-p	00003000	fc:00	3145956	4	4	4	4	4	4	0	0	0	0	0	0	0	0	0	0	0	0 processmap_32
5e74c000	rw-p	00000000	00:00	0	136	4	4	4	4	4	4	0	0	0	0	0	0	0	0	0	0	[heap]
efde9000	r--p	00000000	fc:00	5544451	140	140	140	0	140	0	0	0	0	0	0	0	0	0	0	0	0	0 libc.so.6
efe0c000	r-xp	00023000	fc:00	5544451	1532	784	784	0	784	0	0	0	0	0	0	0	0	0	0	0	0	0 libc.so.6
eff8b000	r--p	001a2000	fc:00	5544451	532	64	64	0	64	0	0	0	0	0	0	0	0	0	0	0	0	0 libc.so.6
f0010000	r--p	00226000	fc:00	5544451	8	8	8	8	8	8	0	0	0	0	0	0	0	0	0	0	0	0 libc.so.6
f0012000	rw-p	00228000	fc:00	5544451	4	4	4	4	4	4	0	0	0	0	0	0	0	0	0	0	0	0 libc.so.6
f0013000	rw-p	00000000	00:00	0	40	16	16	16	16	0	0	0	0	0	0	0	0	0	0	0	0	0
f0024000	rw-p	00000000	00:00	0	8	8	8	8	8	8	0	0	0	0	0	0	0	0	0	0	0	0
f0026000	r--p	00000000	00:00	0	16	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	[vvar]
f002a000	r-xp	00000000	00:00	0	8	4	4	0	4	0	0	0	0	0	0	0	0	0	0	0	0	[vdso]
f002c000	r--p	00000000	fc:00	5544448	4	4	4	0	4	0	0	0	0	0	0	0	0	0	0	0	0	ld-linux.so.2
f002d000	r-xp	00001000	fc:00	5544448	140	140	140	0	140	0	0	0	0	0	0	0	0	0	0	0	0	0 ld-linux.so.2
f0050000	r--p	00024000	fc:00	5544448	56	56	56	0	56	0	0	0	0	0	0	0	0	0	0	0	0	0 ld-linux.so.2
f005e000	r--p	00031000	fc:00	5544448	8	8	8	8	8	8	0	0	0	0	0	0	0	0	0	0	0	0 ld-linux.so.2
f0060000	rw-p	00033000	fc:00	5544448	4	4	4	4	4	4	0	0	0	0	0	0	0	0	0	0	0	0 ld-linux.so.2
ff9bc000	rw-p	00000000	00:00	0	132	12	12	12	12	12	0	0	0	0	0	0	0	0	0	0	0	[stack]
	====	====	====	====	====	====	====	====	====	====	====	====	====	====	====	====	====	====	====	====	====	0 KB
	2788	1272	1272		72	1272	72	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

Memory Map of Linux Process (64 bit system)

Address	Perm	Offset	Device	Inode	Size	Rss	Pss	Pss_Dirty	Referenced	Anonymous	KSM	LazyFree	ShmemPmdMapped	FilePmdMapped	Shared_Hugetlb	Private_Hugetlb	Swap	SwapPss	Locked	THPeligible	ProtectionKey	Mapping
56e4fc8ed000	r--p	00000000	fc:00	3146755	4	4	4	0	4	0	0	0	0	0	0	0	0	0	0	0	0 [processmap]	
56e4fc8ee000	r-xp	00001000	fc:00	3146755	4	4	4	0	4	0	0	0	0	0	0	0	0	0	0	0	0 [processmap]	
56e4fc8ef000	r--p	00002000	fc:00	3146755	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0 [processmap]	
56e4fc8f0000	r--p	00002000	fc:00	3146755	4	4	4	4	4	4	0	0	0	0	0	0	0	0	0	0	0 [processmap]	
56e4fc8f1000	rw-p	00003000	fc:00	3146755	4	4	4	4	4	4	0	0	0	0	0	0	0	0	0	0	0 [processmap]	
56e51f43c000	rw-p	00000000	00:00	0	132	4	4	4	4	4	0	0	0	0	0	0	0	0	0	0	0 [heap]	
7d0ef6400000	r--p	00000000	fc:00	5540086	160	160	3	0	160	0	0	0	0	0	0	0	0	0	0	0	0 [libc.so.6]	
7d0ef6428000	r-xp	00028000	fc:00	5540086	1568	864	19	0	864	0	0	0	0	0	0	0	0	0	0	0	0 [libc.so.6]	
7d0ef65b0000	r--p	001b0000	fc:00	5540086	316	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0 [libc.so.6]	
7d0ef65ff000	r--p	001fe000	fc:00	5540086	16	16	16	16	16	16	0	0	0	0	0	0	0	0	0	0	0 [libc.so.6]	
7d0ef6603000	rw-p	00202000	fc:00	5540086	8	8	8	8	8	8	0	0	0	0	0	0	0	0	0	0	0 [libc.so.6]	
7d0ef6605000	rw-p	00000000	00:00	0	52	20	20	20	20	20	0	0	0	0	0	0	0	0	0	0	0	
7d0ef67fa000	rw-p	00000000	00:00	0	12	8	8	8	8	8	8	0	0	0	0	0	0	0	0	0	0	
7d0ef6804000	rw-p	00000000	00:00	0	8	4	4	4	4	4	4	0	0	0	0	0	0	0	0	0	0	
7d0ef6806000	r--p	00000000	fc:00	5540083	4	4	0	0	4	0	0	0	0	0	0	0	0	0	0	0	0 [ld-linux-x86-64.so.2]	
7d0ef6807000	r-xp	00001000	fc:00	5540083	172	172	3	0	172	0	0	0	0	0	0	0	0	0	0	0	0 [ld-linux-x86-64.so.2]	
7d0ef6832000	r--p	0002c000	fc:00	5540083	40	40	0	0	40	0	0	0	0	0	0	0	0	0	0	0	0 [ld-linux-x86-64.so.2]	
7d0ef683c000	r--p	00036000	fc:00	5540083	8	8	8	8	8	8	0	0	0	0	0	0	0	0	0	0	0 [ld-linux-x86-64.so.2]	
7d0ef683e000	rw-p	00038000	fc:00	5540083	8	8	8	8	8	8	0	0	0	0	0	0	0	0	0	0	0 [ld-linux-x86-64.so.2]	
7ffe0eaec000	rw-p	00000000	00:00	0	132	12	12	12	12	12	0	0	0	0	0	0	0	0	0	0	0 [stack]	
7ffe0fb9000	r--p	00000000	00:00	0	16	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0 [var]	
7ffe0afbd000	r-xp	00000000	00:00	0	8	4	0	0	4	0	0	0	0	0	0	0	0	0	0	0	0 [vds]	
ffffffffffff600000	--xp	00000000	00:00	0	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0 [vsyscall]	
<hr/>																						
					2684	1348	129	96	1348	96	0	0	0	0	0	0	0	0	0	0	0 KB	

Background Knowledge: System Calls

What is System Call?

When a process needs to invoke a kernel service, it invokes a procedure call in the operating system interface using special instructions (not a **call** instruction in x86). Such a procedure is called a system call.

The system call enters the kernel; the kernel performs the service and returns. Thus a process alternates between executing in user space and kernel space.

System calls are generally not invoked directly by a program, but rather via wrapper functions in glibc (or perhaps some other library).

Popular System Call

On Unix, Unix-like and other POSIX-compliant operating systems, popular system calls are open, read, write, close, wait, exec, fork, exit, and kill.

Many modern operating systems have hundreds of system calls. For example, Linux and OpenBSD each have over 300 different calls, FreeBSD has over 500, Windows 7 has close to 700.

Glibc interfaces

Often, but not always, the name of the wrapper function is the same as the name of the system call that it invokes

For example, glibc contains a function `chdir()` which invokes the underlying "chdir" system call.

Tools: strace & ltrace

misc/firstflag main.c

```
int main(int argc, char *argv[])
```

1

```
printf("Congratulations on getting your first flag!!\n");
print_flag();
```

3

flag.h

```
int print_flag()
```

{

```
FILE *fp = NULL;
```

```
char buff[MAX_FLAG_SIZE] = {0}
```

```
fp = fopen("/flag", "r")
```

if (fp == NULL)

1

```
printf("Error: Cannot open the flag file!!!\n");  
return 1;
```

10

```
fread(buff, MAX_FLAG_SIZE - 2, 1, fp);
```

```
printf("The flag is: %s\n", buff)
```

`fclose(fp)`

```
return 0
```

Tools: strace & ltrace

Execve - first system call

Access - check file permission

Brk - check data segment/heap

Arch_prctl - set architecture-specific thread state

Fcntl - manipulate file descriptor

Openat - similar to open

Fstat - get file status

Mmap

Close Read

Read
Broad/4 - similar to read

Protect - set protection on a region of memory

Mprotect - set protection on a region of memory
Munmap - map files or devices into memory

Plain Write

Write

Use “man 2 syscall name” to check out its usage

Making a System Call in x86/64 Assembly

On x86/x86-64, most system calls rely on the software interrupt.

A software interrupt is caused either by an **exceptional condition** in the processor itself, or a **special instruction** (the `int 0x80` instruction or `syscall` instruction)

For example: a divide-by-zero exception will be thrown if the processor's arithmetic logic unit is commanded to divide a number by zero as this instruction is in error and impossible.

Making a System Call in x86 Assembly (INT 0x80)

x86 (32-bit)

Compiled from Linux 4.14.0 headers.

NR	syscall name	references	%eax	arg0 (%ebx)	arg1 (%ecx)	arg2 (%edx)	arg3 (%esi)	arg4 (%edi)	arg5 (%ebp)
0	restart_syscall	man/cs/	0x00	-	-	-	-	-	-
1	exit	man/cs/	0x01	int error_code	-	-	-	-	-
2	fork	man/cs/	0x02	-	-	-	-	-	-
3	read	man/cs/	0x03	unsigned int fd	char *buf	size_t count	-	-	-
4	write	man/cs/	0x04	unsigned int fd	const char *buf	size_t count	-	-	-
5	open	man/cs/	0x05	const char *filename	int flags	umode_t mode	-	-	-
6	close	man/cs/	0x06	unsigned int fd	-	-	-	-	-
7	waitpid	man/cs/	0x07	pid_t pid	int *stat_addr	int options	-	-	-
8	creat	man/cs/	0x08	const char *pathname	umode_t mode	-	-	-	-
9	link	man/cs/	0x09	const char *oldname	const char *newname	-	-	-	-
10	unlink	man/cs/	0x0a	const char *pathname	-	-	-	-	-
11	execve	man/cs/	0x0b	const char *filename	const char *const *argv	const char *const *envp	-	-	-
12	chdir	man/cs/	0x0c	const char *filename	-	-	-	-	-
13	time	man/cs/	0x0d	time_t *tloc	-	-	-	-	-
14	mknod	man/cs/	0x0e	const char *filename	umode_t mode	unsigned dev	-	-	-
15	chmod	man/cs/	0x0f	const char *filename	umode_t mode	-	-	-	-
16	lchown	man/cs/	0x10	const char *filename	uid_t user	gid_t group	-	-	-
17	break	man/cs/	0x11	?	?	?	?	?	?

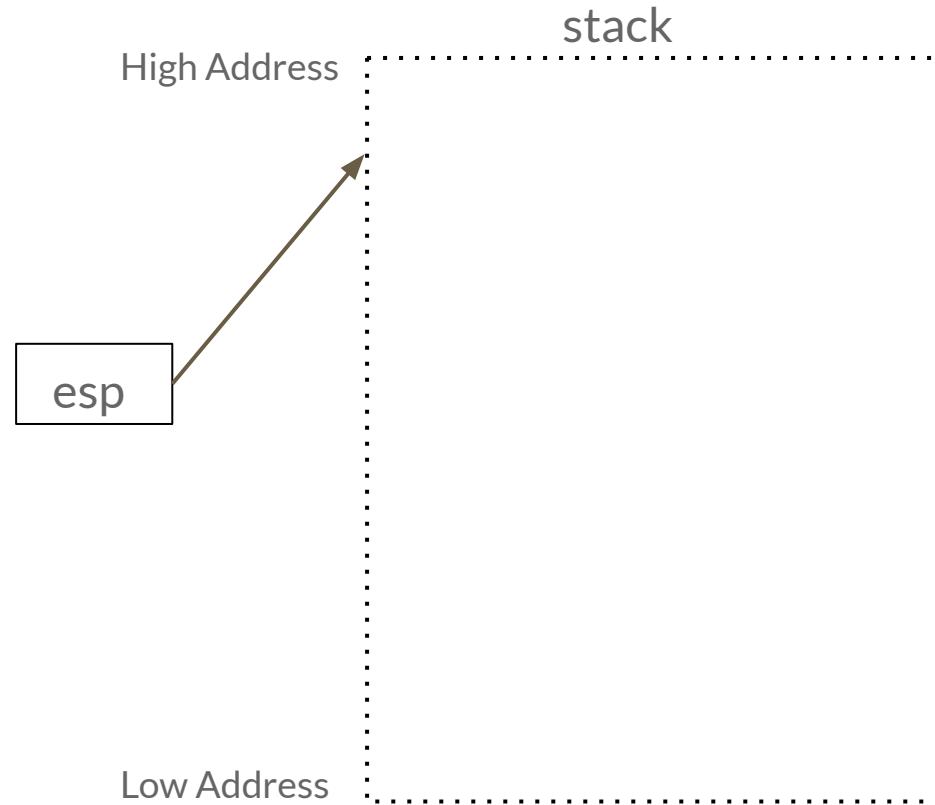
Making a System Call in x86 Assembly

```
xor eax,eax  
push eax  
push 0x68732f2f  
push 0x6e69622f  
mov ebx,esp  
push eax  
push ebx  
mov ecx,esp  
mov al,0xb  
int 0x80
```

Dec	Hx	Oct	Char	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr
0	0 000	NUL	(null)	32	20 040	 	Space		64	40 100	@	Ø	96	60 140	`	`		
1	1 001	SOH	(start of heading)	33	21 041	!	!	!	65	41 101	A	A	97	61 141	a	a		
2	2 002	STX	(start of text)	34	22 042	"	"	"	66	42 102	B	B	98	62 142	b	b		
3	3 003	ETX	(end of text)	35	23 043	#	#	#	67	43 103	C	C	99	63 143	c	c		
4	4 004	EOT	(end of transmission)	36	24 044	$	\$	\$	68	44 104	D	D	100	64 144	d	d		
5	5 005	ENQ	(enquiry)	37	25 045	%	%	%	69	45 105	E	E	101	65 145	e	e		
6	6 006	ACK	(acknowledge)	38	26 046	&	&	&	70	46 106	F	F	102	66 146	f	f		
7	7 007	BEL	(bell)	39	27 047	'	'	'	71	47 107	G	G	103	67 147	g	g		
8	8 010	BS	(backspace)	40	28 050	(((72	48 110	H	H	104	68 150	h	h		
9	9 011	TAB	(horizontal tab)	41	29 051)))	73	49 111	I	I	105	69 151	i	i		
10	A 012	LF	(NL line feed, new line)	42	2A 052	*	*	*	74	4A 112	J	J	106	6A 152	j	j		
11	B 013	VT	(vertical tab)	43	2B 053	+	+	+	75	4B 113	K	K	107	6B 153	k	k		
12	C 014	FF	(NP form feed, new page)	44	2C 054	,	,	,	76	4C 114	L	L	108	6C 154	l	l		
13	D 015	CR	(carriage return)	45	2D 055	-	-	-	77	4D 115	M	M	109	6D 155	m	m		
14	E 016	SO	(shift out)	46	2E 056	.	.	.	78	4E 116	N	N	110	6E 156	n	n		
15	F 017	SI	(shift in)	47	2F 057	/	/	/	79	4F 117	O	O	111	6F 157	o	o		
16	10 020	DLE	(data link escape)	48	30 060	0	0	0	80	50 120	P	P	112	70 160	p	p		
17	11 021	DC1	(device control 1)	49	31 061	1	1	1	81	51 121	Q	Q	113	71 161	q	q		
18	12 022	DC2	(device control 2)	50	32 062	2	2	2	82	52 122	R	R	114	72 162	r	r		
19	13 023	DC3	(device control 3)	51	33 063	3	3	3	83	53 123	S	S	115	73 163	s	s		
20	14 024	DC4	(device control 4)	52	34 064	4	4	4	84	54 124	T	T	116	74 164	t	t		
21	15 025	NAK	(negative acknowledge)	53	35 065	5	5	5	85	55 125	U	U	117	75 165	u	u		
22	16 026	SYN	(synchronous idle)	54	36 066	6	6	6	86	56 126	V	V	118	76 166	v	v		
23	17 027	ETB	(end of trans. block)	55	37 067	7	7	7	87	57 127	W	W	119	77 167	w	w		
24	18 030	CAN	(cancel)	56	38 070	8	8	8	88	58 130	X	X	120	78 170	x	x		
25	19 031	EM	(end of medium)	57	39 071	9	9	9	89	59 131	Y	Y	121	79 171	y	y		
26	1A 032	SUB	(substitute)	58	3A 072	:	:	:	90	5A 132	Z	Z	122	7A 172	z	z		
27	1B 033	ESC	(escape)	59	3B 073	;	;	;	91	5B 133	[[123	7B 173	{	{		
28	1C 034	FS	(file separator)	60	3C 074	<	<	<	92	5C 134	\	\	124	7C 174	|			
29	1D 035	GS	(group separator)	61	3D 075	=	=	=	93	5D 135]]	125	7D 175	}	}		
30	1E 036	RS	(record separator)	62	3E 076	>	>	>	94	5E 136	^	^	126	7E 176	~	~		
31	1F 037	US	(unit separator)	63	3F 077	?	?	?	95	5F 137	_	_	127	7F 177		DEL		

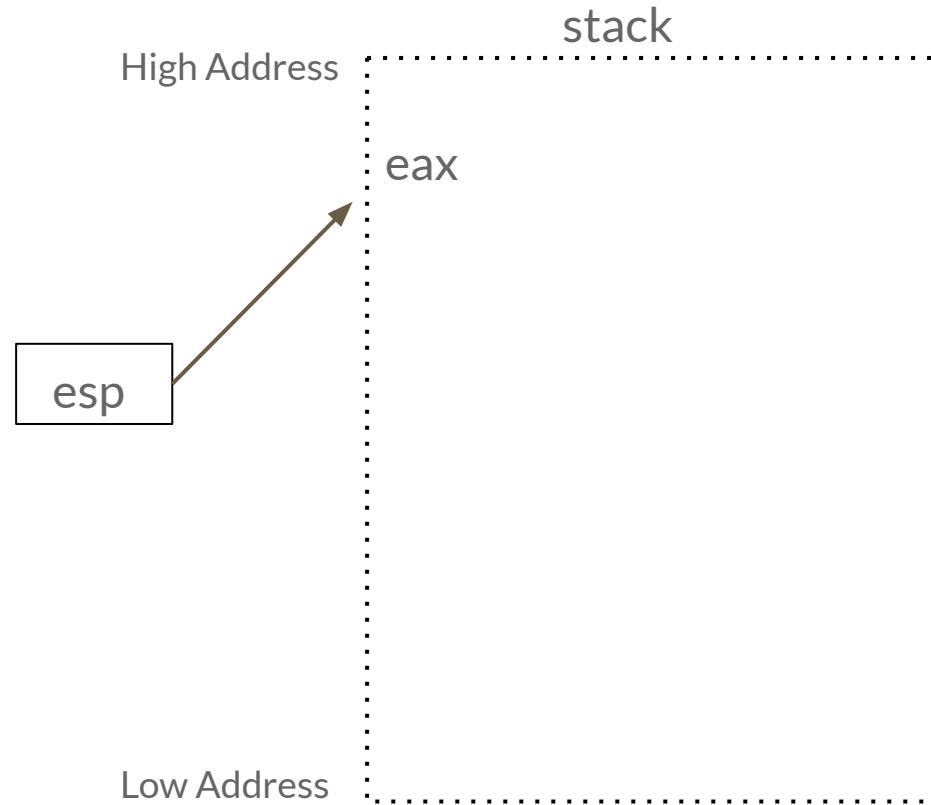
Making a System Call in x86 Assembly

```
xor eax,eax  
push eax  
push 0x68732f2f  
push 0x6e69622f  
mov ebx,esp  
push eax  
push ebx  
mov ecx,esp  
mov al,0xb  
int 0x80
```



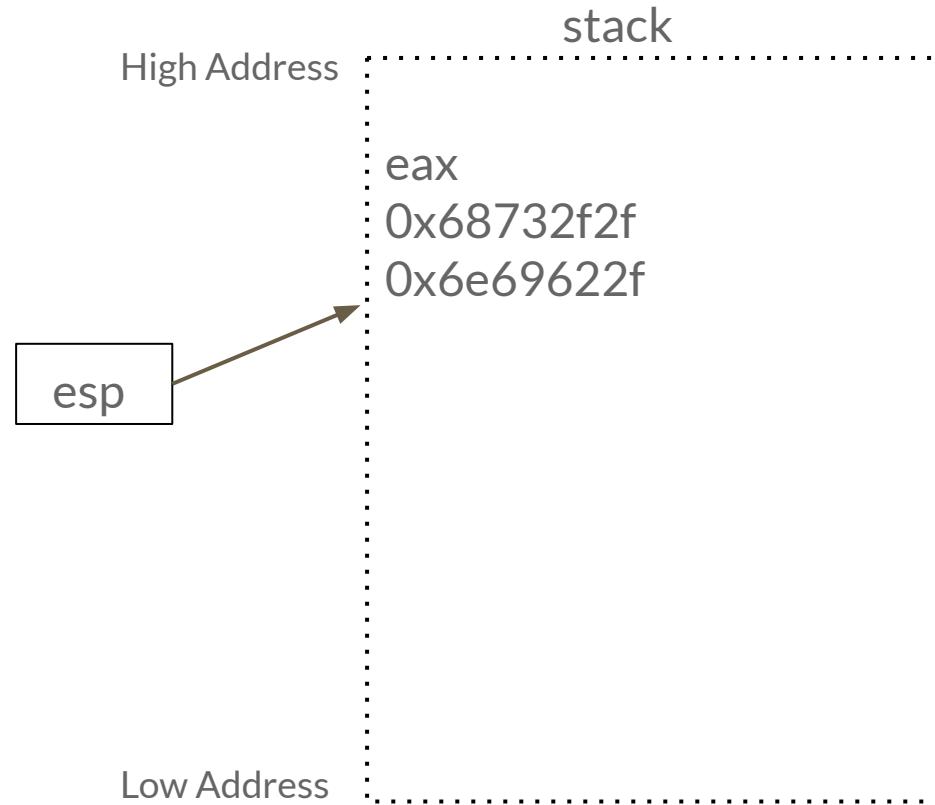
Making a System Call in x86 Assembly

```
xor eax,eax  
push eax  
push 0x68732f2f  
push 0x6e69622f  
mov ebx,esp  
push eax  
push ebx  
mov ecx,esp  
mov al,0xb  
int 0x80
```



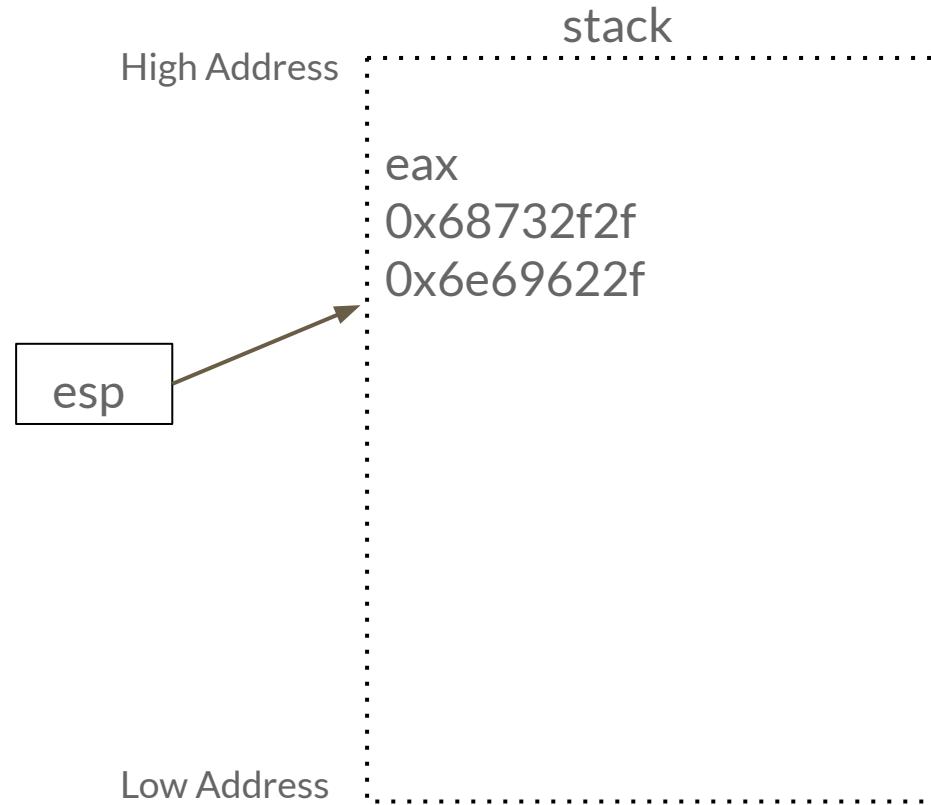
Making a System Call in x86 Assembly

```
xor eax,eax  
push eax  
push 0x68732f2f  
push 0x6e69622f  
mov ebx,esp  
push eax  
push ebx  
mov ecx,esp  
mov al,0xb  
int 0x80
```



Making a System Call in x86 Assembly

```
xor eax,eax  
push eax  
push 0x68732f2f  
push 0x6e69622f  
mov ebx,esp  
push eax  
push ebx  
mov ecx,esp  
mov al,0xb  
int 0x80
```



Making a System Call in x86 Assembly

[execve\(2\)](#)

System Calls Manual

NAME

execve - execute program

LIBRARY

Standard C library ([libc](#), [-lc](#))

SYNOPSIS

```
#include <unistd.h>

int execve(const char *pathname, char *const _Nullable argv[],  
          char *const _Nullable envp[]);
```

/bin/sh, 0x0

0x00000000

Address of /bin/sh, 0x00000000

execve("/bin/sh", address of string "/bin/sh", 0)

Making a System Call in x86_64 (64-bit) Assembly

x86_64 (64-bit)

Compiled from [Linux 4.14.0 headers](#).

NR	syscall name	references	%rax	arg0 (%rdi)	arg1 (%rsi)	arg2 (%rdx)	arg3 (%r10)	arg4 (%r8)	arg5 (%r9)
0	read	man/cs/	0x00	unsigned int fd	char *buf	size_t count	-	-	-
1	write	man/cs/	0x01	unsigned int fd	const char *buf	size_t count	-	-	-
2	open	man/cs/	0x02	const char *filename	int flags	umode_t mode	-	-	-
3	close	man/cs/	0x03	unsigned int fd	-	-	-	-	-
4	stat	man/cs/	0x04	const char *filename	struct __old_kernel_stat *statbuf	-	-	-	-
5	fstat	man/cs/	0x05	unsigned int fd	struct __old_kernel_stat *statbuf	-	-	-	-
6	lstat	man/cs/	0x06	const char *filename	struct __old_kernel_stat *statbuf	-	-	-	-
7	poll	man/cs/	0x07	struct pollfd *ufds	unsigned int nfds	int timeout	-	-	-
8	lseek	man/cs/	0x08	unsigned int fd	off_t offset	unsigned int whence	-	-	-
9	mmap	man/cs/	0x09	?	?	?	?	?	?
10	mprotect	man/cs/	0x0a	unsigned long start	size_t len	unsigned long prot	-	-	-
11	munmap	man/cs/	0x0b	unsigned long addr	size_t len	-	-	-	-
12	brk	man/cs/	0x0c	unsigned long brk	-	-	-	-	-
13	rt_sigaction	man/cs/	0x0d	int	const struct sigaction *	struct sigaction *	size_t	-	-
14	rt_sigprocmask	man/cs/	0x0e	int how	sigset_t *set	sigset_t *oset	size_t sigsetsize	-	-
15	rt_sigreturn	man/cs/	0x0f	?	?	?	?	?	?
16	ioctl	man/cs/	0x10	unsigned int fd	unsigned int cmd	unsigned long arg	-	-	-

https://chromium.googlesource.com/chromiumos/docs/+/master/constants/syscalls.md#x86-32_bit

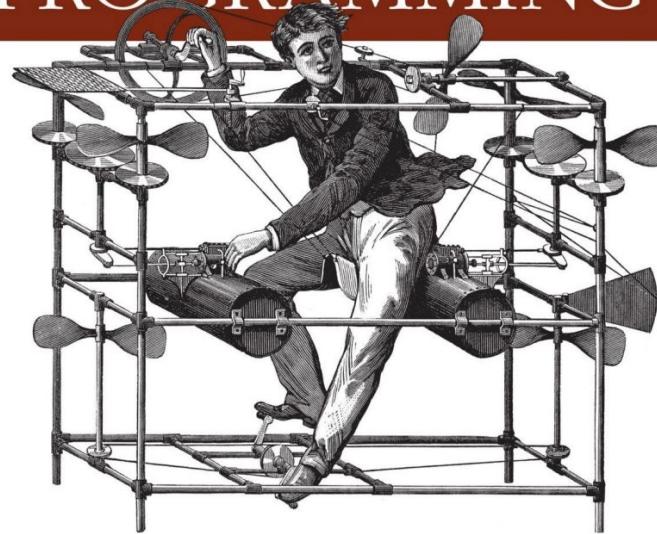
Making a System Call in x86_64 (64-bit) Assembly

NR	syscall name	references	%rax	arg0 (%rdi)	arg1 (%rsi)	arg2 (%rdx)	arg3 (%r10)	arg4 (%r8)	arg5 (%r9)
59	execve	man/ cs/	0x3b	const char *filename	const char *const *argv	const char *const *envp	-	-	-

```
push rax  
xor rdx, rdx  
xor rsi, rsi  
mov rbx,'/bin//sh'  
push rbx  
push rsp  
pop rdi  
mov al, 59  
syscall
```

SYSTEM AND LIBRARY CALLS EVERY PROGRAMMER NEEDS TO KNOW

LINUX SYSTEM PROGRAMMING



O'REILLY®

ROBERT LOVE

Background Knowledge: System Calls

Channels of Communication for Linux Process

Every process in Linux has three initial, standard channels of communication:

- Standard Input (stdin, fd=0) is the channel through which the process takes input.
For example, your shell uses Standard Input to read the commands that you input.
- Standard Output (stdout, fd=1) is the channel through which processes output normal data, such as the flag when it is printed to you in previous challenges or the output of utilities such as ls.
- Standard Error (stderr, fd=2) is the channel through which processes output error details. For example, if you mistype a command, the shell will output, over standard error, that this command does not exist.

Examples

Redirecting output > or 1>

```
echo hi > asdf echo hi 1> asdf
```

Appending output >>

```
echo hi >> asdf
```

Redirecting errors 2>

```
/challenge/run 2> errors.log
```

Redirecting input <

```
rev < messagefile
```

Channels of Communication for Linux Process

Process can also take input from command line arguments

ls -al c

at /flag

cat 1.txt 2.txt 3.txt

Pipe

The | (pipe) operator. Standard output from the command to the left of the pipe will be connected to (piped into) the standard input of the command to the right of the pipe.

```
echo hello-world | wc -c
```

Background Knowledge: Environment and Shell Variables

Environment and Shell Variables

Environment and Shell variables are a set of dynamic **named values**, stored within the system that are used by applications launched in shells.

KEY=value

KEY="Some other value"

KEY=value1:value2

The names of the variables are case-sensitive (UPPER CASE). Multiple values must be separated by the colon : character. There is no space around the equals = symbol.

Environment and Shell Variables

Environment variables are variables that are available **system-wide** and are **inherited** by all spawned child processes and shells.

Shell variables are variables that apply only to the **current shell instance**. Each shell such as zsh and bash, has its own set of internal shell variables.

Common Environment Variables

USER - The current logged in user.

HOME - The home directory of the current user.

EDITOR - The default file editor to be used. This is the editor that will be used when you type edit in your terminal.

SHELL - The path of the current user's shell, such as bash or zsh.

LOGNAME - The name of the current user.

PATH - A list of directories to be searched when executing commands.

LANG - The current locales settings.

TERM - The current terminal emulation.

MAIL - Location of where the current user's mail is stored.

Commands

env – The command allows you to run another program in a custom environment without modifying the current one. When used without an argument it will print a list of the current environment variables.

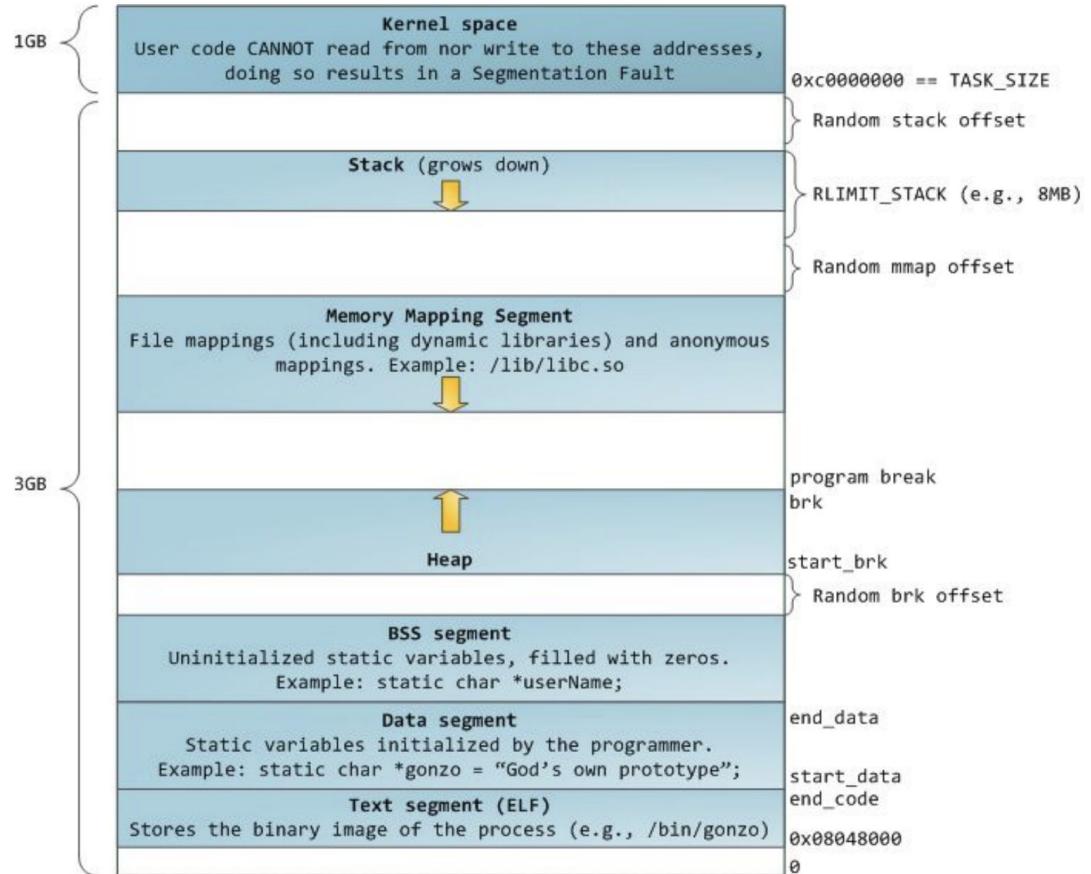
printenv – The command prints all or the specified environment variables.

set – The command sets or unsets shell variables. When used without an argument it will print a list of all variables including environment and shell variables, and shell functions.

unset – The command deletes shell and environment variables.

export – The command sets environment variables

The environment variables live towards the top of the stack, together with command line arguments

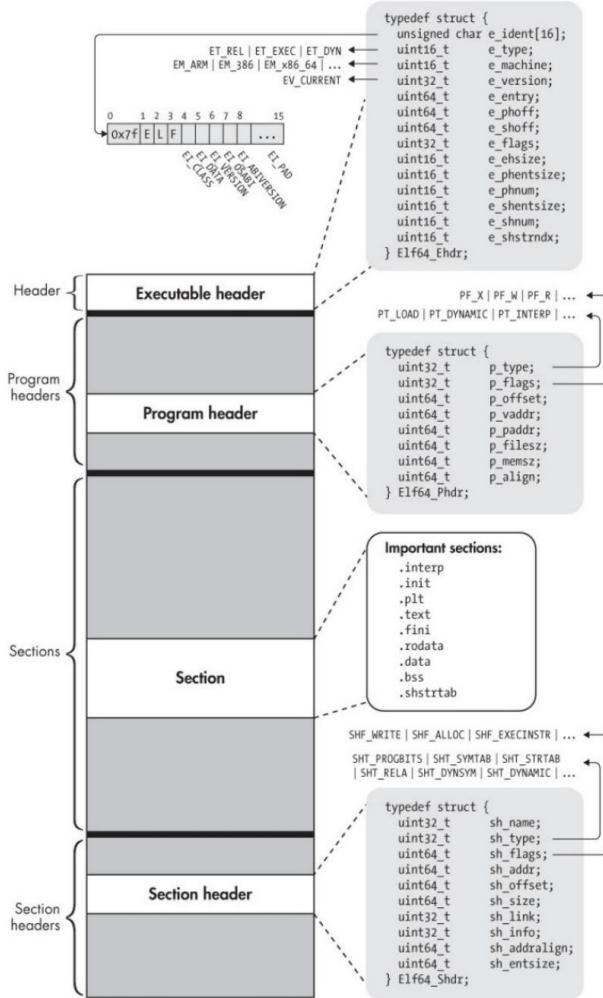


Background Knowledge: Executable and Linkable Format (ELF)

Commands

The **Executable and Linkable Format (ELF)** is a common standard file format for executable files, object code, shared libraries, and core dumps. Filename extension none, .axf, .bin, .elf, .o, .prx, .puff, .ko, .mod and .so

Contains the program and its data. Describes how the program should be loaded (program/segment headers). Contains metadata describing program components (section headers).



- Executable (a.out), object files (.o), shared libraries (.a), even core dumps.
- Four types of components: an **executable header**, a series of (optional) **program headers**, a number of **sections**, and a series of (optional) **section headers**, one per section.

Executable Header

```
typedef struct {
    unsigned char e_ident[16];      /* Magic number and other info      */ /0x7F ELF ..
    uint16_t      e_type;          /* Object file type Executable, obj, dynamic lib */
    uint16_t      e_machine;       /* Architecture x86-64, Arm        */
    uint32_t      e_version;       /* Object file version             */
    uint64_t      e_entry;         /* Entry point virtual address   */
    uint64_t      e_phoff;         /* Program header table file offset */
    uint64_t      e_shoff;         /* Section header table file offset */
    uint32_t      e_flags;         /* Processor-specific flags       */
    uint16_t      e_ehsize;        /* ELF header size in bytes      */
    uint16_t      e_phentsize;     /* Program header table entry size */
    uint16_t      e_phnum;         /* Program header table entry count */
    uint16_t      e_shentsize;     /* Section header table entry size */
    uint16_t      e_shnum;         /* Section header table entry count */
    uint16_t      e_shstrndx;      /* Section header string table index*/
} Elf64_Ehdr;
```

ELF headers

```
arman@aserver:~$ readelf -h /bin/ls
ELF Header:
  Magic: 7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class: ELF64
  Data: 2's complement, little endian
  Version: 1 (current)
  OS/ABI: UNIX - System V
  ABI Version: 0
  Type: DYN (Position-Independent Executable file)
  Machine: Advanced Micro Devices X86-64
  Version: 0x1
  Entry point address: 0x6d30
  Start of program headers: 64 (bytes into file)
  Start of section headers: 140328 (bytes into file)
  Flags: 0x0
  Size of this header: 64 (bytes)
  Size of program headers: 56 (bytes)
  Number of program headers: 13
  Size of section headers: 64 (bytes)
  Number of section headers: 31
  Section header string table index: 30
```

Sections

The code and data in an ELF binary are logically divided into contiguous non-overlapping chunks called sections. The structure of each section varies depending on the contents.

The division into sections is intended to provide a convenient organization for use by the *linker*.

Section Header Format

```
typedef struct {
    uint32_t sh_name;          /* Section name (string tbl index) */
    uint32_t sh_type;          /* Section type */
    uint64_t sh_flags;         /* Section flags */
    uint64_t sh_addr;          /* Section virtual addr at execution */
    uint64_t sh_offset;         /* Section file offset */
    uint64_t sh_size;          /* Section size in bytes */
    uint32_t sh_link;          /* Link to another section */
    uint32_t sh_info;          /* Additional section information */
    uint64_t sh_addralign;     /* Section alignment */
    uint64_t sh_entsize;        /* Entry size if section holds table */
} Elf64_Shdr;
```

SHF_WRITE | SHF_ALLOC | SHF_EXECINSTR | ... ←
SHT_PROGBITS | SHT_SYMTAB | SHT_STRTAB
| SHT_REL | SHT_DYNSYM | SHT_DYNAMIC | ... ←

```
typedef struct {
    uint32_t sh_name;
    uint32_t sh_type;
    uint64_t sh_flags;
    uint64_t sh_addr;
    uint64_t sh_offset;
    uint64_t sh_size;
    uint32_t sh_link;
    uint32_t sh_info;
    uint64_t sh_addralign;
    uint64_t sh_entsize;
} Elf64_Shdr;
```

Each section is described by its section header

redelf -S a.out

sh_flags

SHF_WRITE: the section is writable at runtime.

SHF_ALLOC: the contents of the section are to be loaded into virtual memory when executing the binary.

SHF_EXECINSTR: the section contains executable instructions.

SHF_WRITE | SHF_ALLOC | SHF_EXECINSTR | ...

SHT_PROGBITS | SHT_SYMTAB | SHT_STRTAB
| SHT_REL | SHT_DYNSYM | SHT_DYNAMIC | ...

```
typedef struct {
    uint32_t      sh_name;
    uint32_t      sh_type;
    uint64_t      sh_flags;
    uint64_t      sh_addr;
    uint64_t      sh_offset;
    uint64_t      sh_size;
    uint32_t      sh_link;
    uint32_t      sh_info;
    uint64_t      sh_addralign;
    uint64_t      sh_entsize;
} Elf64_Shdr;
```

```
arman@aserver:~/STC/software-security-course-binaries/misc$ readelf -S add_32
```

```
There are 31 section headers, starting at offset 0x385c:
```

Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	00000000	000000	000000	00		0	0	0
[1]	.interp	PROGBITS	000001b4	0001b4	000013	00	A	0	0	1
[2]	.note.gnu.bu[...]	NOTE	000001c8	0001c8	000024	00	A	0	0	4
[3]	.note.gnu.pr[...]	NOTE	000001ec	0001ec	00001c	00	A	0	0	4
[4]	.note.ABI-tag	NOTE	00000208	000208	000020	00	A	0	0	4
[5]	.gnu.hash	GNU_HASH	00000228	000228	000020	04	A	6	0	4
[6]	.dynsym	DYNSYM	00000248	000248	0000a0	10	A	7	1	4
[7]	.dynstr	STRTAB	000002e8	0002e8	0000bb	00	A	0	0	1
[8]	.gnu.version	VERSYM	000003a4	0003a4	000014	02	A	6	0	2
[9]	.gnu.version_r	VERNEED	000003b8	0003b8	000040	00	A	7	1	4
[10]	.rel.dyn	REL	000003f8	0003f8	000040	08	A	6	0	4
[11]	.rel.plt	REL	00000438	000438	000020	08	AI	6	24	4
[12]	.init	PROGBITS	00001000	001000	000024	00	AX	0	0	4
[13]	.plt	PROGBITS	00001030	001030	000050	04	AX	0	0	16
[14]	.plt.got	PROGBITS	00001080	001080	000010	10	AX	0	0	16
[15]	.plt.sec	PROGBITS	00001090	001090	000040	10	AX	0	0	16
[16]	.text	PROGBITS	000010d0	0010d0	000259	00	AX	0	0	16
[17]	.fini	PROGBITS	0000132c	00132c	000018	00	AX	0	0	4
[18]	.rodata	PROGBITS	00002000	002000	000044	00	A	0	0	4
[19]	.eh_frame_hdr	PROGBITS	00002044	002044	000054	00	A	0	0	4
[20]	.eh_frame	PROGBITS	00002098	002098	00014c	00	A	0	0	4
[21]	.init_array	INIT_ARRAY	00003ed0	002ed0	000004	04	WA	0	0	4
[22]	.fini_array	FINI_ARRAY	00003ed4	002ed4	000004	04	WA	0	0	4
[23]	.dynamic	DYNAMIC	00003ed8	002ed8	0000f8	08	WA	7	0	4
[24]	.got	PROGBITS	00003fd0	002fd0	000030	04	WA	0	0	4
[25]	.data	PROGBITS	00004000	003000	000008	00	WA	0	0	4
[26]	.bss	NOBITS	00004008	003008	000004	00	WA	0	0	1
[27]	.comment	PROGBITS	00000000	003008	00002a	01	MS	0	0	1
[28]	.symtab	SYMTAB	00000000	003034	000490	10		29	47	4
[29]	.strtab	STRTAB	00000000	0034c4	00027d	00		0	0	1
[30]	.shstrtab	STRTAB	00000000	003741	000118	00		0	0	1

Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
L (link order), O (extra OS processing required), G (group), T (TLS),
C (compressed), x (unknown), o (OS specific), E (exclude),
D (mbind), p (processor specific)

redelf -S a.out

Sections

.init: executable code that performs initialization tasks and needs to run before any other code in the binary is executed.

.fini: code that runs after the main program completes.

.text: where the main code of the program resides.

Sections

.rodata section, which stands for “read-only data,” is dedicated to storing constant values. Because it stores constant values, .rodata is not writable.

The default values of initialized variables are stored in the .data section, which is marked as writable since the values of variables may change at runtime.

the .bss section reserves space for uninitialized variables. The name historically stands for “block started by symbol,” referring to the reserving of blocks of memory for (symbolic) variables.

Lazy Binding (.plt, .got, .got.plt Sections)

Binding at Load Time: When a binary is loaded into a process for execution, the dynamic linker resolves references to functions located in shared libraries. The addresses of shared functions were not known at compile time.

In reality - Lazy Binding: many of the relocations are typically not done right away when the binary is loaded but are deferred until the first reference to the unresolved location is actually made.

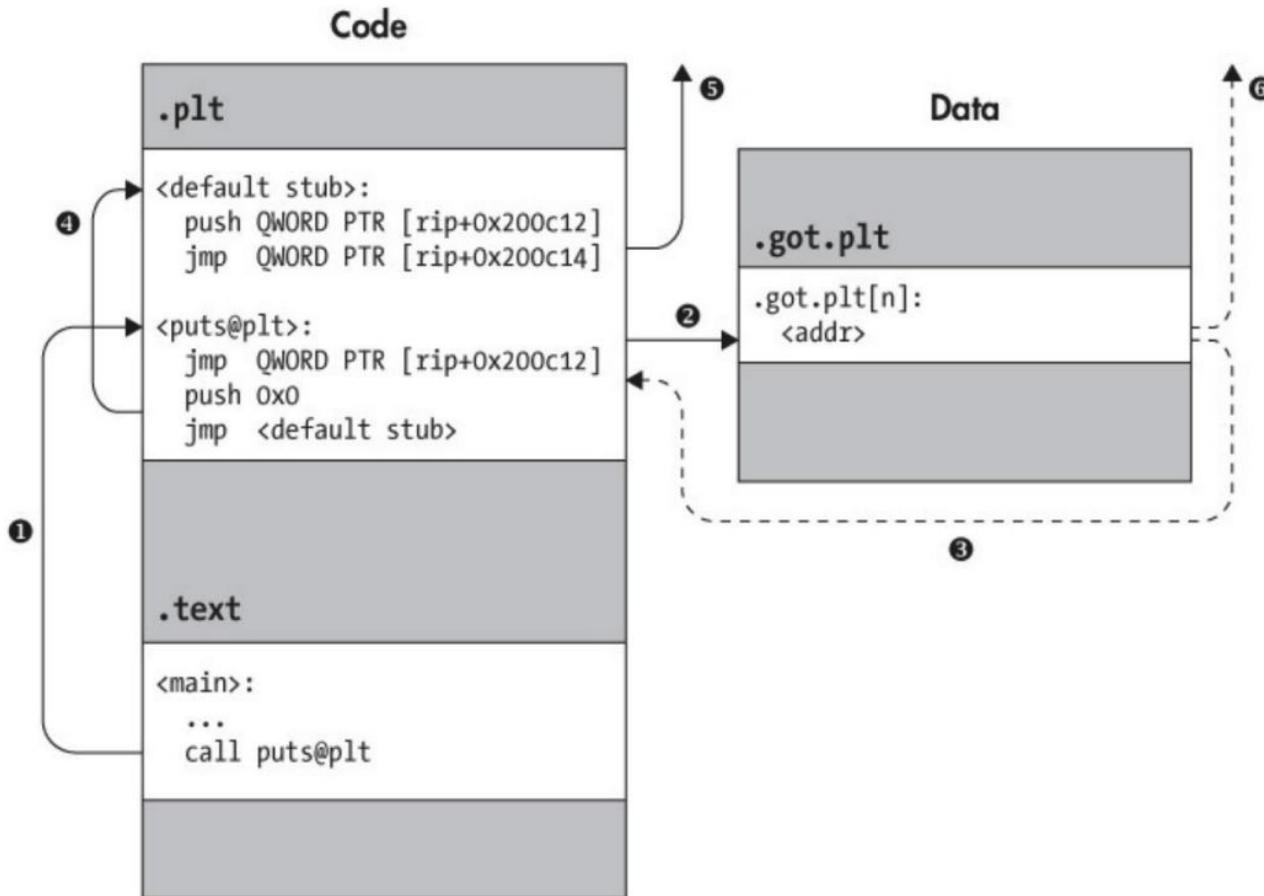
Lazy Binding (.plt, .got, .got.plt Sections)

Lazy binding in Linux ELF binaries is implemented with the help of two special sections, called the Procedure Linkage Table (.plt) and the Global Offset Table (.got).

.plt is a code section that contains executable code. The PLT consists entirely of stubs of a well-defined format, dedicated to directing calls from the .text section to the appropriate library location.

.got.plt is a data section.

Dynamically Resolving a Library Function Using the PLT



Example: Debug misc/lazyb

Section View (Section Header)

VS.

Segment View (Program Header)

The program header table provides a segment view of the binary, as opposed to the section view provided by the section header table.

The section view of an ELF binary is meant for static linking purposes.

The segment view is used by the operating system and dynamic linker when loading an ELF into a process for execution to locate the relevant code and data and decide what to load into virtual memory.

Segments are simply a bunch of sections bundled together.

Program Header Format

```
typedef struct {
    uint32_t sh_name;          /* Section name (string tbl index) */
    uint32_t sh_type;          /* Section type */
    uint64_t sh_flags;         /* Section flags */
    uint64_t sh_addr;          /* Section virtual addr at execution */
    uint64_t sh_offset;         /* Section file offset */
    uint64_t sh_size;          /* Section size in bytes */
    uint32_t sh_link;          /* Link to another section */
    uint32_t sh_info;          /* Additional section information */
    uint64_t sh_addralign;     /* Section alignment */
    uint64_t sh_entsize;        /* Entry size if section holds table */
} Elf64_Shdr;
```

SHF_WRITE | SHF_ALLOC | SHF_EXECINSTR | ... ←
SHT_PROGBITS | SHT_SYMTAB | SHT_STRTAB
| SHT_REL | SHT_DYNSYM | SHT_DYNAMIC | ... ←

```
typedef struct {
    uint32_t sh_name;
    uint32_t sh_type;
    uint64_t sh_flags;
    uint64_t sh_addr;
    uint64_t sh_offset;
    uint64_t sh_size;
    uint32_t sh_link;
    uint32_t sh_info;
    uint64_t sh_addralign;
    uint64_t sh_entsize;
} Elf64_Shdr;
```

Each section is described by its section header

redelf -S a.out

```
arman@aserver:~/STC/software-security-course-binaries/misc$ readelf -l add_32
```

Elf file type is DYN (Position-Independent Executable file)

Entry point 0x1160

There are 12 program headers, starting at offset 52

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
PHDR	0x000034	0x00000034	0x00000034	0x00180	0x00180	R 0x4	
INTERP	0x0001b4	0x000001b4	0x000001b4	0x00013	0x00013	R 0x1	
[Requesting program interpreter: /lib/ld-linux.so.2]							
LOAD	0x000000	0x00000000	0x00000000	0x00458	0x00458	R 0x1000	
LOAD	0x001000	0x00001000	0x00001000	0x00344	0x00344	R E 0x1000	
LOAD	0x002000	0x00002000	0x00002000	0x001e4	0x001e4	R 0x1000	
LOAD	0x002ed0	0x00003ed0	0x00003ed0	0x00138	0x0013c	RW 0x1000	
DYNAMIC	0x002ed8	0x00003ed8	0x00003ed8	0x000f8	0x000f8	RW 0x4	
NOTE	0x0001c8	0x000001c8	0x000001c8	0x00060	0x00060	R 0x4	
GNU_PROPERTY	0x0001ec	0x000001ec	0x000001ec	0x0001c	0x0001c	R 0x4	
GNU_EH_FRAME	0x002044	0x00002044	0x00002044	0x00054	0x00054	R 0x4	
GNU_STACK	0x000000	0x00000000	0x00000000	0x00000	0x00000	RW 0x10	
GNU_RELRO	0x002ed0	0x00003ed0	0x00003ed0	0x00130	0x00130	R 0x1	

Section to Segment mapping:

Segment Sections...

00	
01	.interp
02	.interp .note.gnu.build-id .note.gnu.property .note.ABI-tag .gnu.hash .dynsym .dynstr .gnu.version .gnu.version_r .rel.dyn .rel.pl
03	.init .plt .plt.got .plt.sec .text .fini
04	.rodata .eh_frame_hdr .eh_frame
05	.init_array .fini_array .dynamic .got .data .bss
06	.dynamic
07	.note.gnu.build-id .note.gnu.property .note.ABI-tag
08	.note.gnu.property
09	.eh_frame_hdr
10	
11	.init_array .fini_array .dynamic .got

Background Knowledge: Manual Binary Analysis Tools

Tools for this class

file

readelf

strings

nm

objdump

GDB

[*optional*] *IDA Pro*

[*optional*] *ghidra*

[*optional*] *Binary Ninja*

GDB Cheat Sheet

Start gdb using:

```
gdb <binary>
```

Pass initial commands for gdb through a file

```
gdb <binary> -x <initfile>
```

To start the program and breakpoint at main()

```
start <argv>
```

To start the program and breakpoint at _start

```
start <argv>
```

To run the program without breakpoint

```
r <argv>
```

Use another program's output as stdin in GDB:

```
r <<< $(python2 -c "print '\x12\x34'*5")
```

GDB Cheat Sheet

Set breakpoint at address:

`b *0x80000000`

Set breakpoint at beginning of a function:

`b main`

....

`b <filename:line number>`

`b <line number>`

Disassemble 10 instructions from an address:

`x/10i 0x80000000`

Exam 15 dword (w) from an address; show hex (x):

`x/15wx 0x80000000`

Exam 3 qword (g) from an address; show hex (x):

`x/3gx 0x80000000`

GDB Cheat Sheet

To show breakpoints

`info b`

To remove breakpoints

`clear <function name>`

`clear *<instruction address>`

`clear <filename:line number>`

`clear <line number>`

GDB Cheat Sheet

Use “examine” or “x” command

`x/32xw <memory location>` to see memory contents at memory location, showing 32 hexadecimal words

`x/5s <memory location>` to show 5 strings (null terminated) at a particular memory location

`x/10i <memory location>` to show 10 instructions at particular memory location

See registers

`info reg`

Step an instruction

`si`

GDB Script

Use “examine” or “x” command

`x/32xw <memory location>` to see memory contents at memory location, showing 32 hexadecimal words

`x/5s <memory location>` to show 5 strings (null terminated) at a particular memory location

`x/10i <memory location>` to show 10 instructions at particular memory location

See registers

`info reg`

Step an instruction

`si`

Shell Cheat Sheet

Run a program and use another program's output as a parameter

```
program $(python2 -c "print '\x12\x34'*5")
```

Reading

1. <https://iq.thc.org/how-does-linux-start-a-process>